



US 20170185674A1

(19) United States

(12) Patent Application Publication (10) Pub. No.: US 2017/0185674 A1  
(43) Pub. Date: Jun. 29, 2017

## (54) ONTOLOGY MAPPING METHOD AND APPARATUS

(71) Applicant: Semantic Technologies Pty Ltd, North Sydney, New South Wales (AU)

(72) Inventors: Albert Donald TONKIN, Seaforth, New South Wales (AU); Robert Naughton CONNELL, East Lindfield, New South Wales (AU); Michael Anthony CAVANAGH, Neutral Bay, New South Wales (AU); Kamalraj JAIRAM, Newington, New South Wales (AU); Qingyan WANG, Beecroft, New South Wales (AU); Dung Xuan Thi LE, Homebush, New South Wales (AU)

(73) Assignee: Semantic Technologies Pty Ltd, North Sydney, New South Wales (AU)

(21) Appl. No.: 15/301,357

(22) PCT Filed: Apr. 2, 2015

(86) PCT No.: PCT/AU2015/000195

§ 371 (c)(1),  
(2) Date: Sep. 30, 2016

## Related U.S. Application Data

(60) Provisional application No. 61/974,180, filed on Apr. 2, 2014.

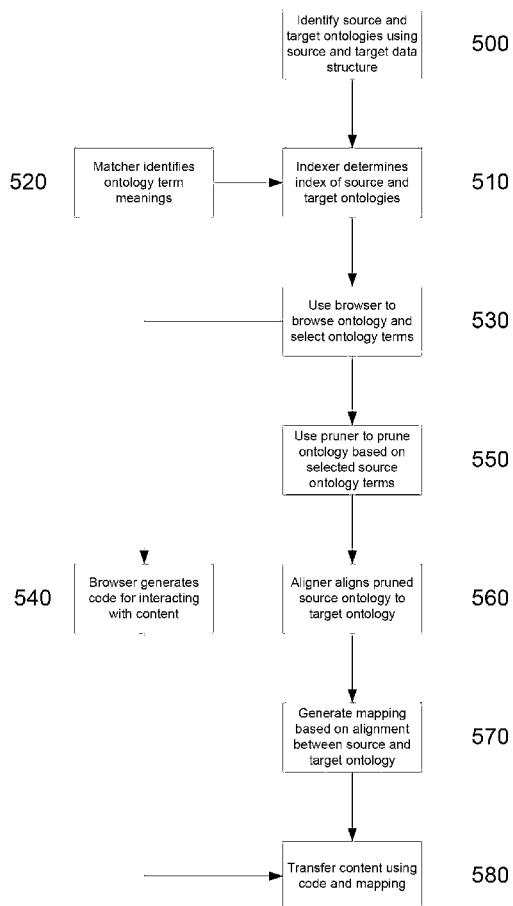
## Publication Classification

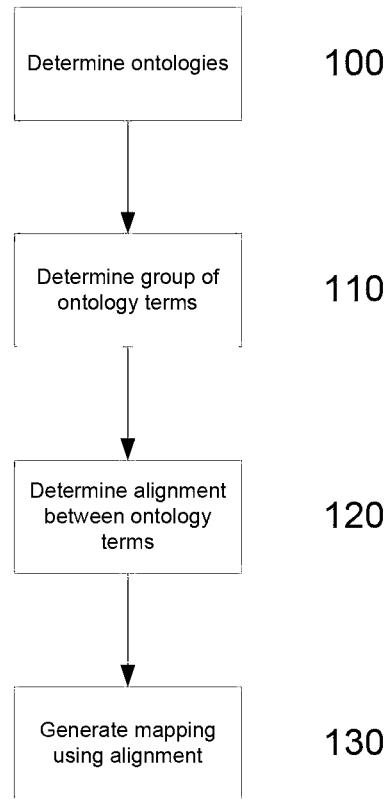
(51) Int. Cl. G06F 17/30 (2006.01)

(52) U.S. Cl. CPC .. G06F 17/30734 (2013.01); G06F 17/30631 (2013.01)

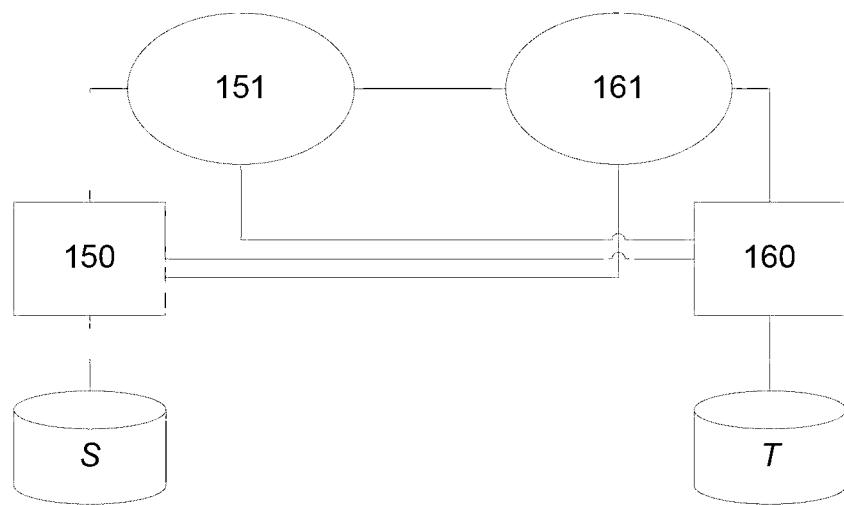
## (57) ABSTRACT

Apparatus for use in generating a mapping using ontologies, the apparatus including at least one electronic processing device that determines ontologies having respective ontology terms, determines a group of ontology terms from at least one of the ontologies, determines alignment between ontology terms in the ontologies for at least some of the group of ontology terms, the alignment being determined at least partially in accordance with an ontology term meaning of the ontology terms and generates a mapping in accordance with the alignment.





**Fig. 1A**



**Fig. 1B**

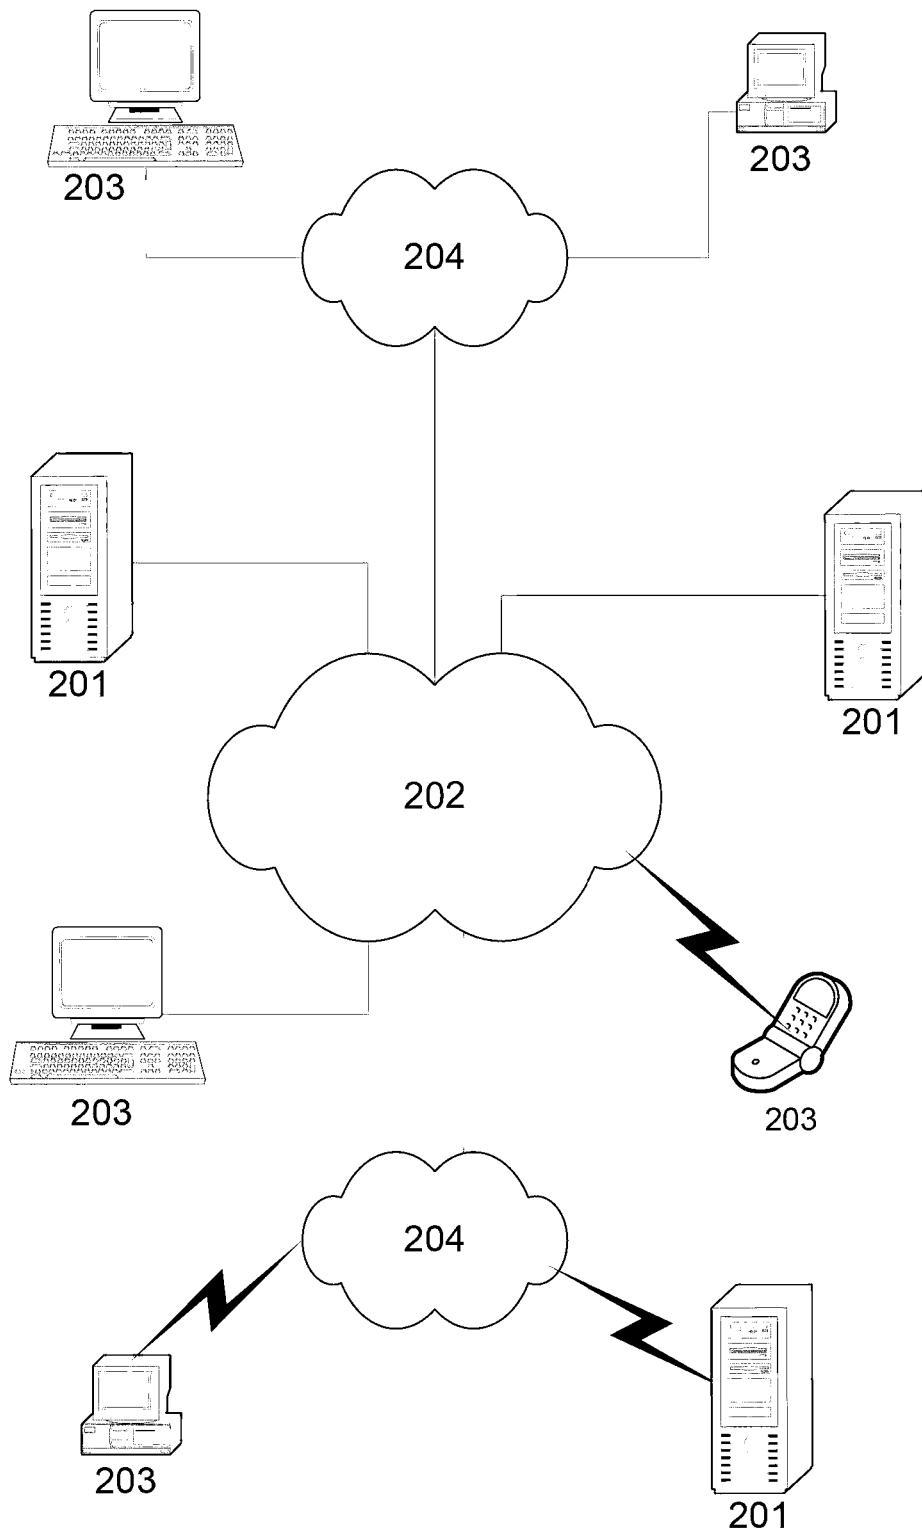
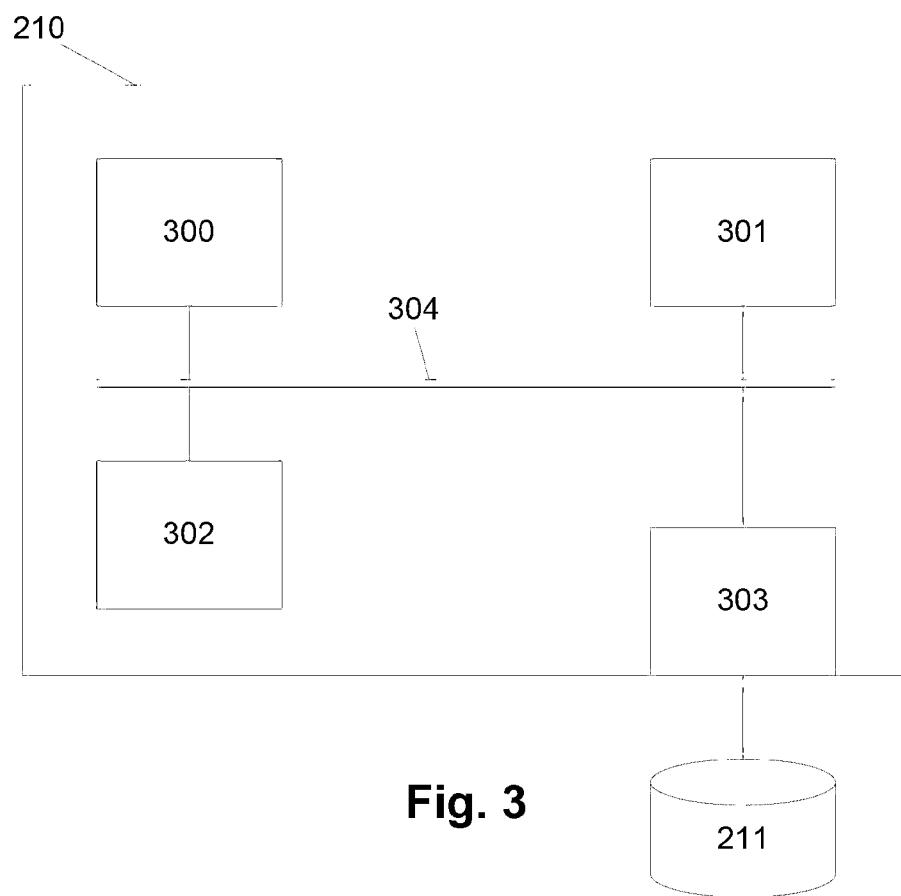
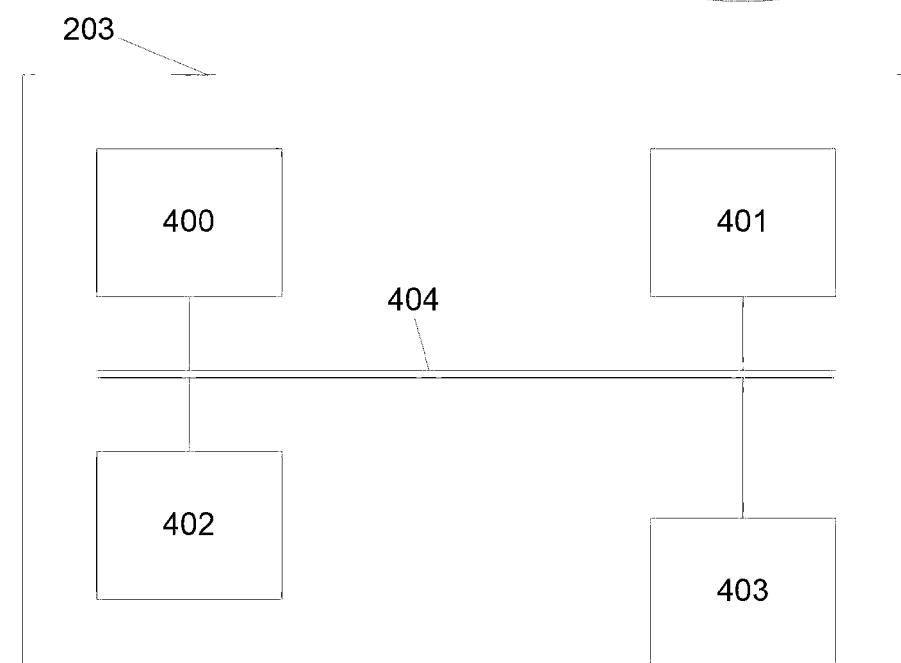


Fig. 2



**Fig. 3**



**Fig. 4**

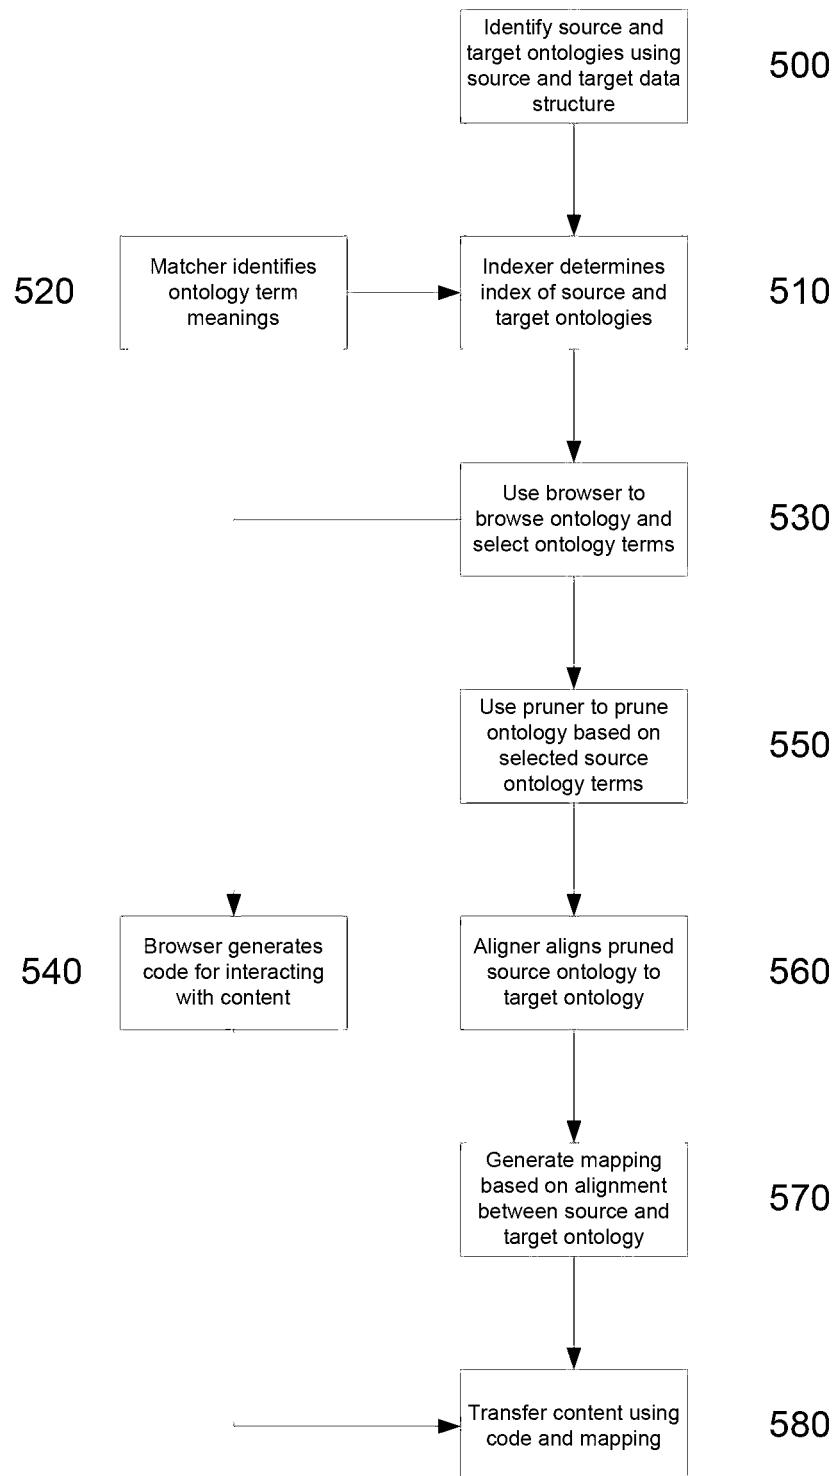
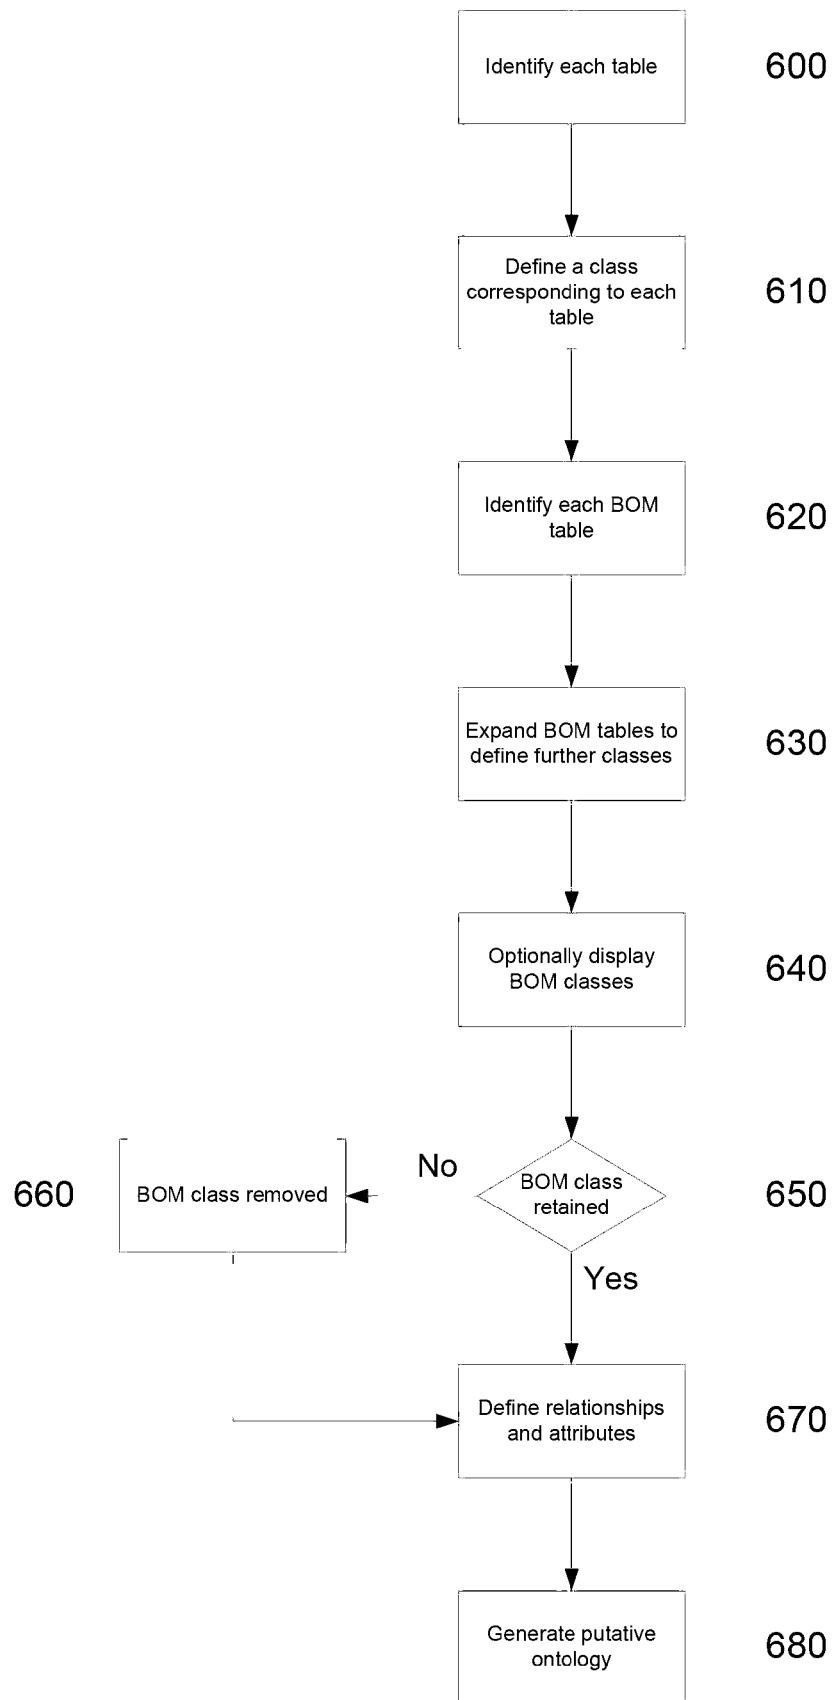


Fig. 5

**Fig. 6**

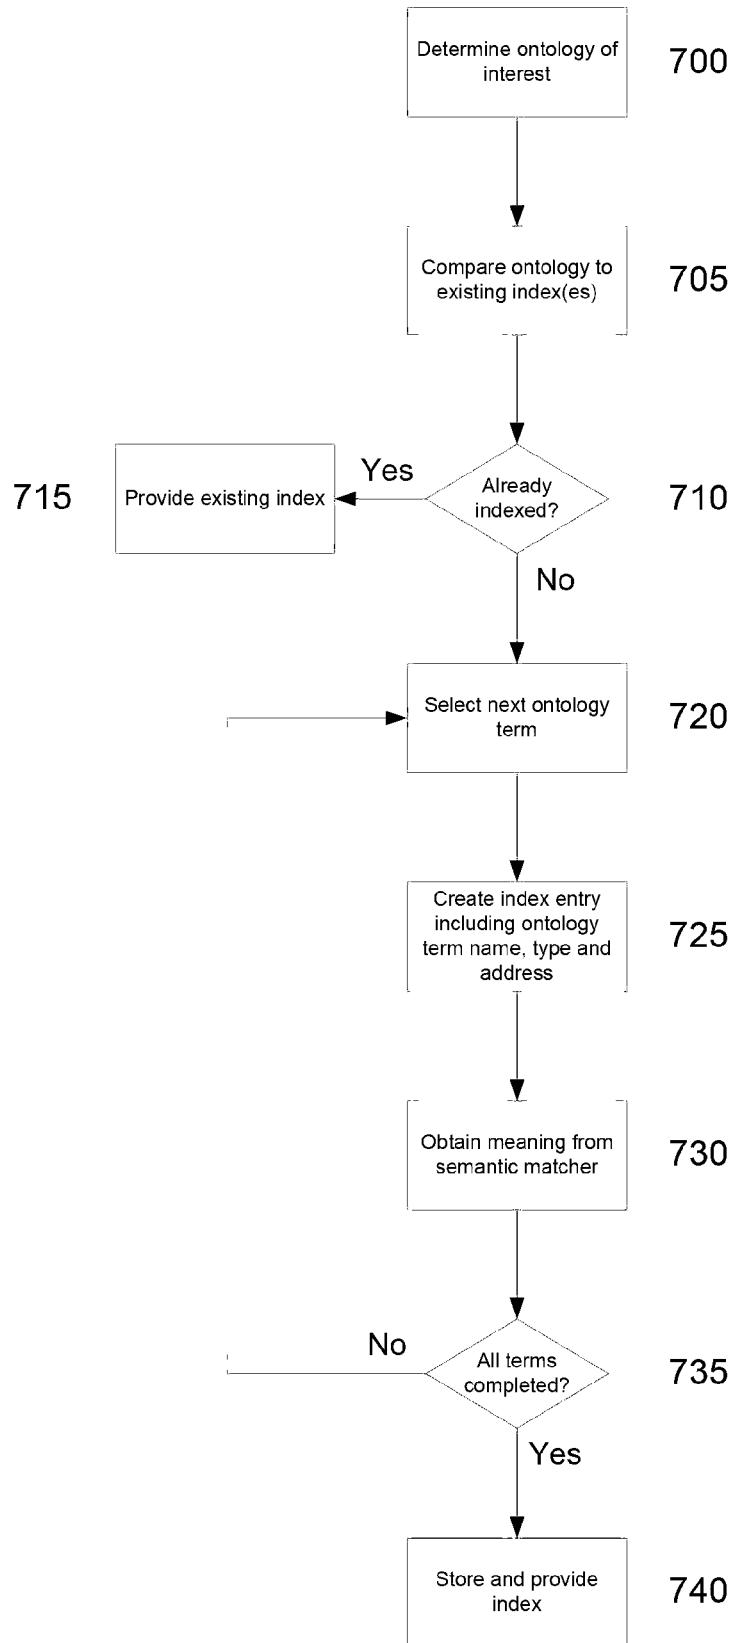


Fig. 7

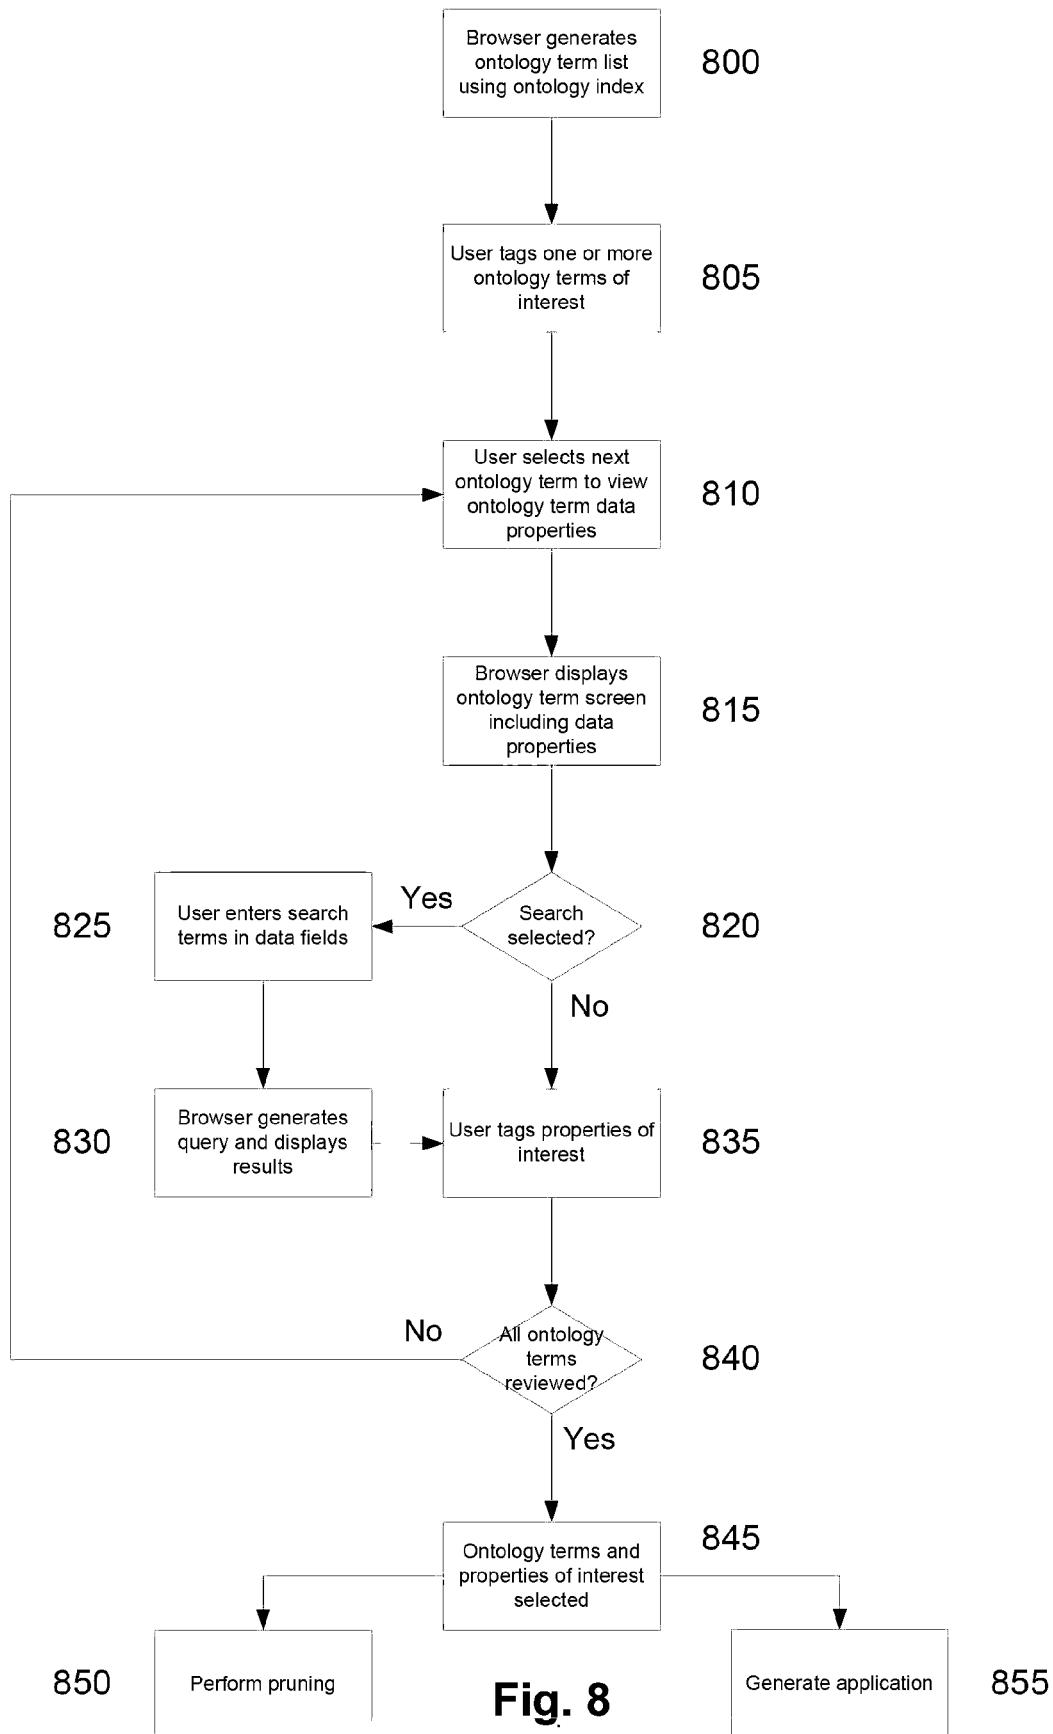


Fig. 8

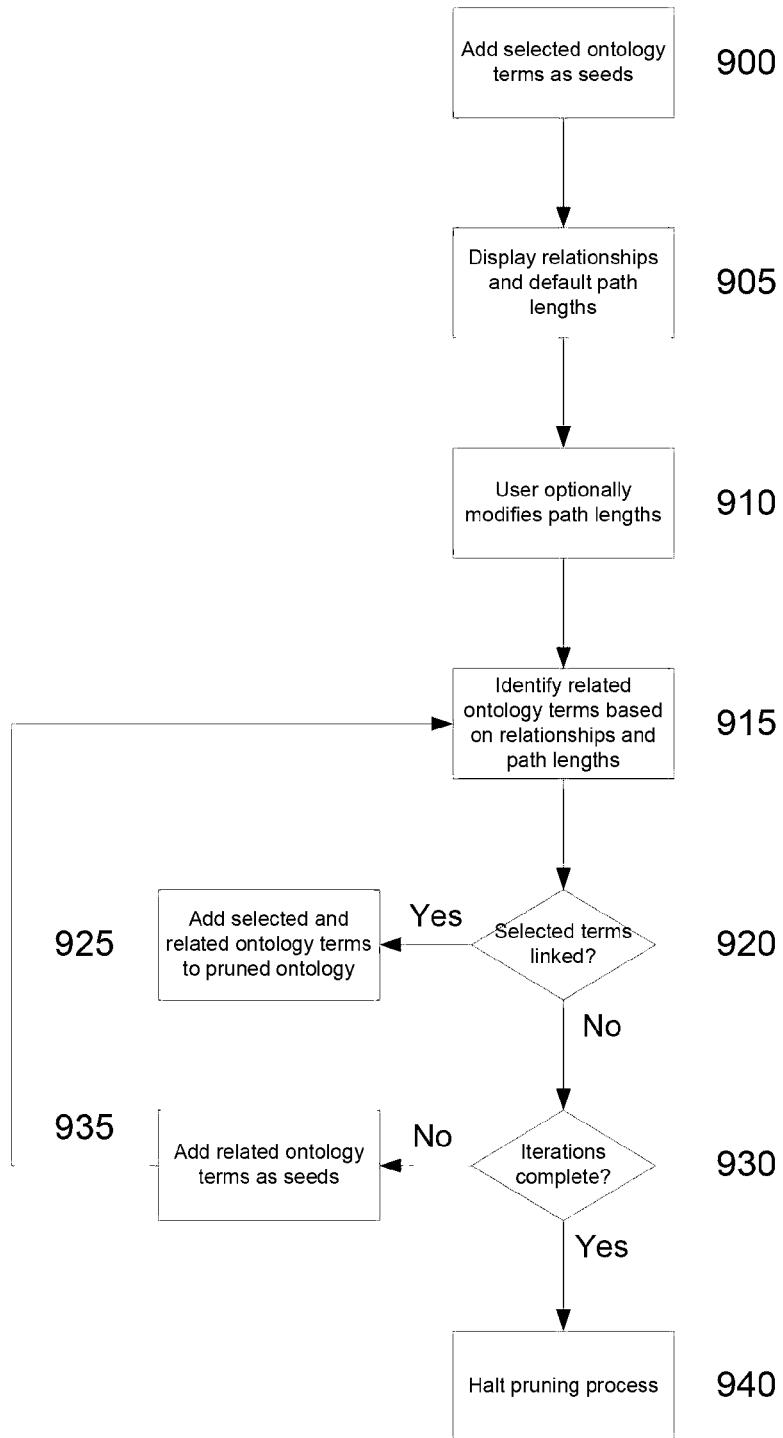


Fig. 9

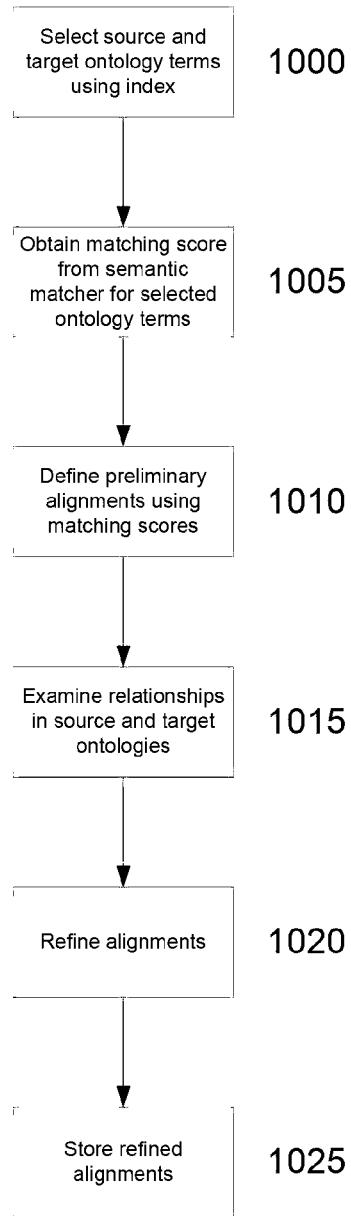


Fig. 10

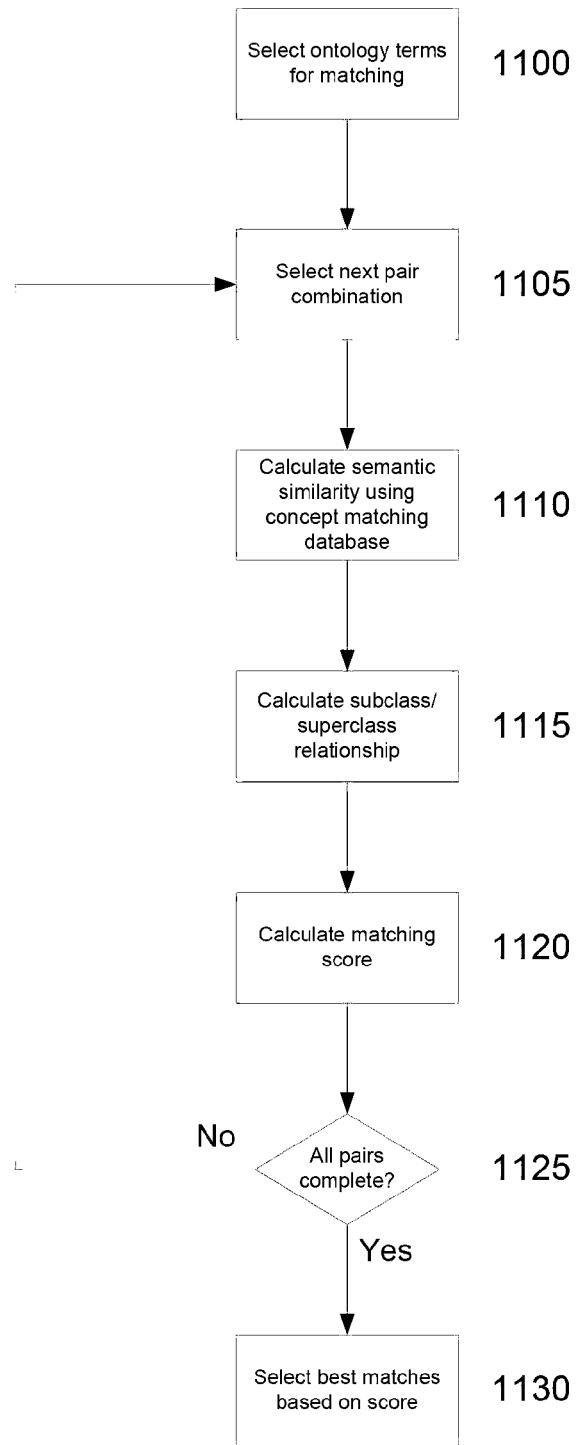
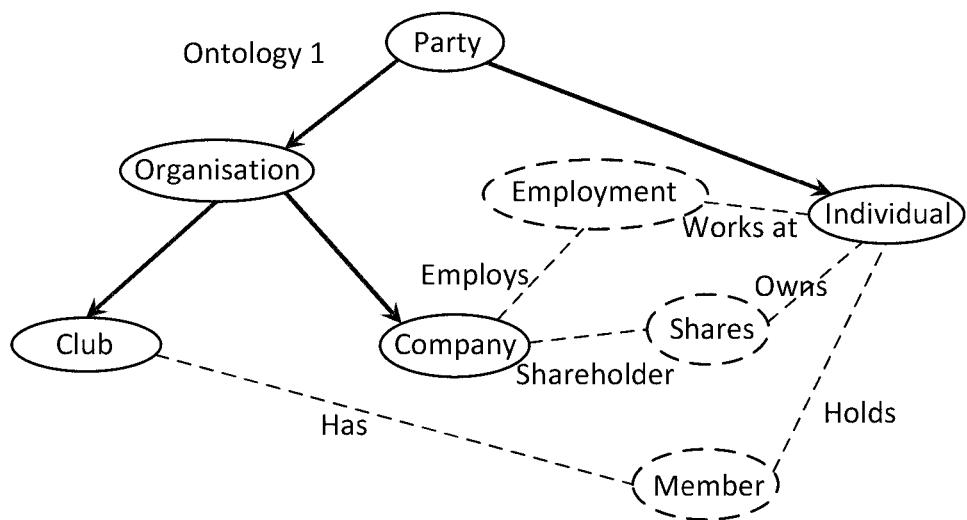
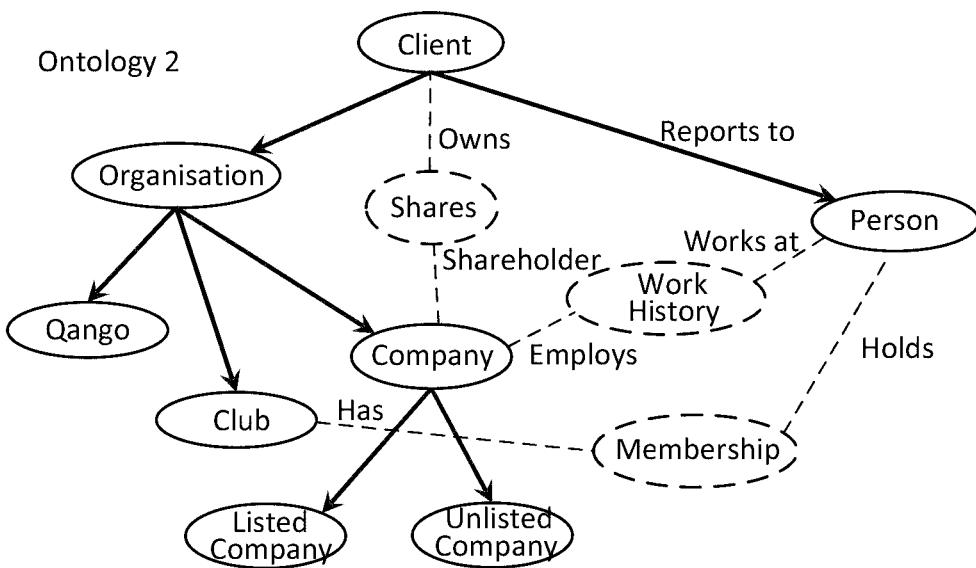


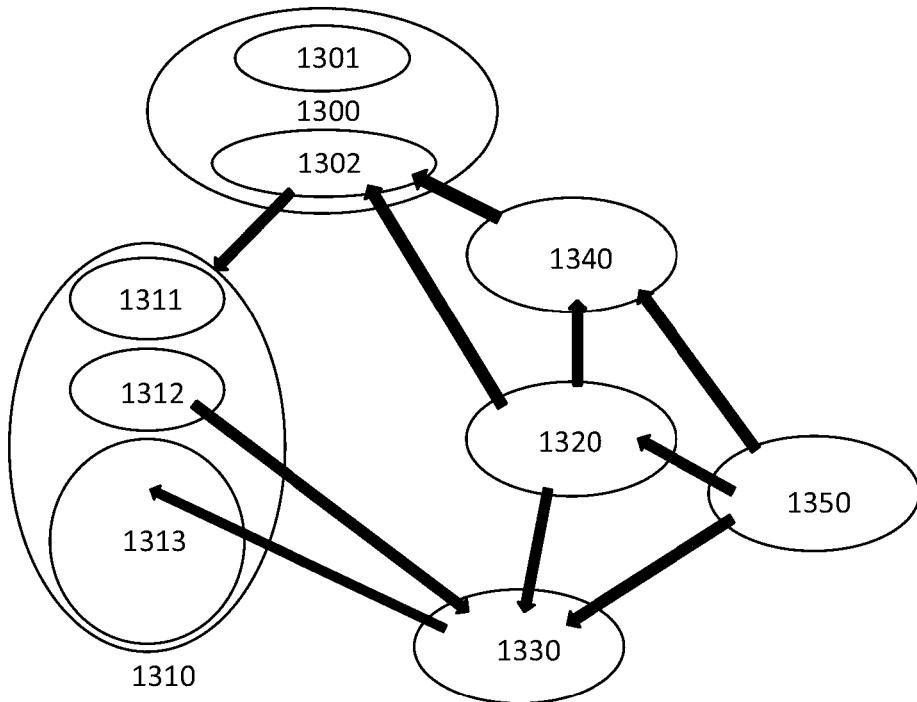
Fig. 11



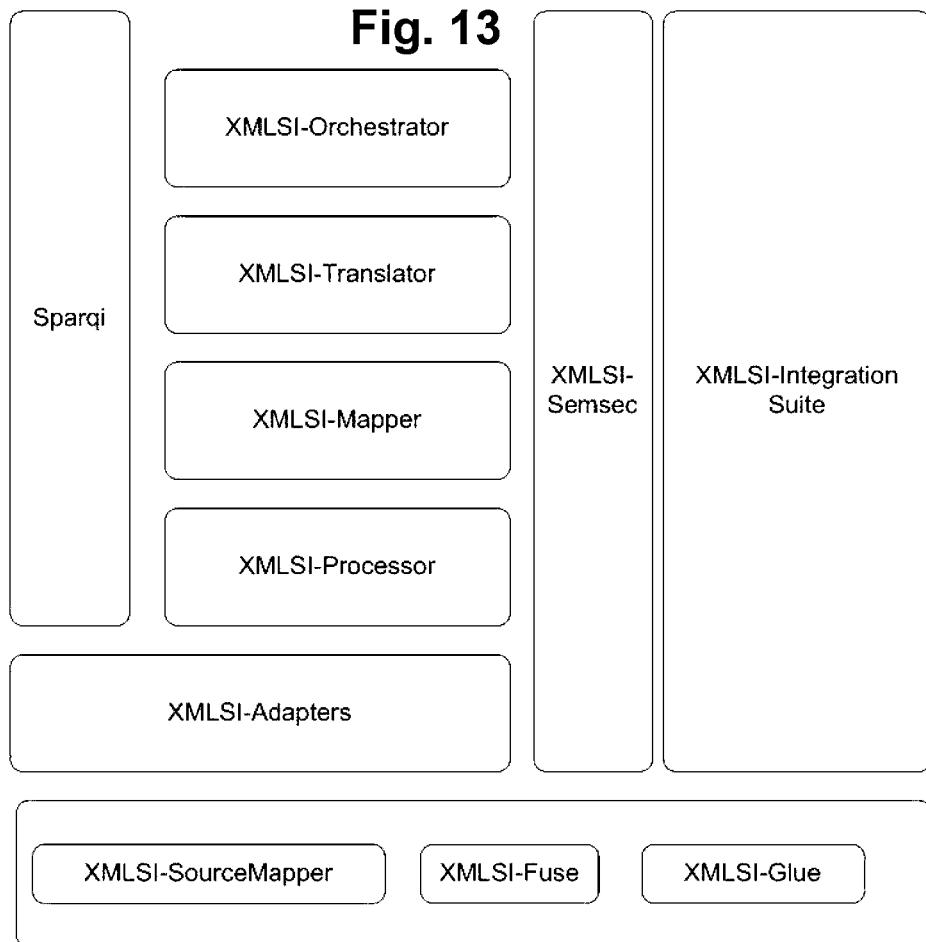
**Fig. 12A**



**Fig. 12B**



**Fig. 13**



**Fig. 14A**

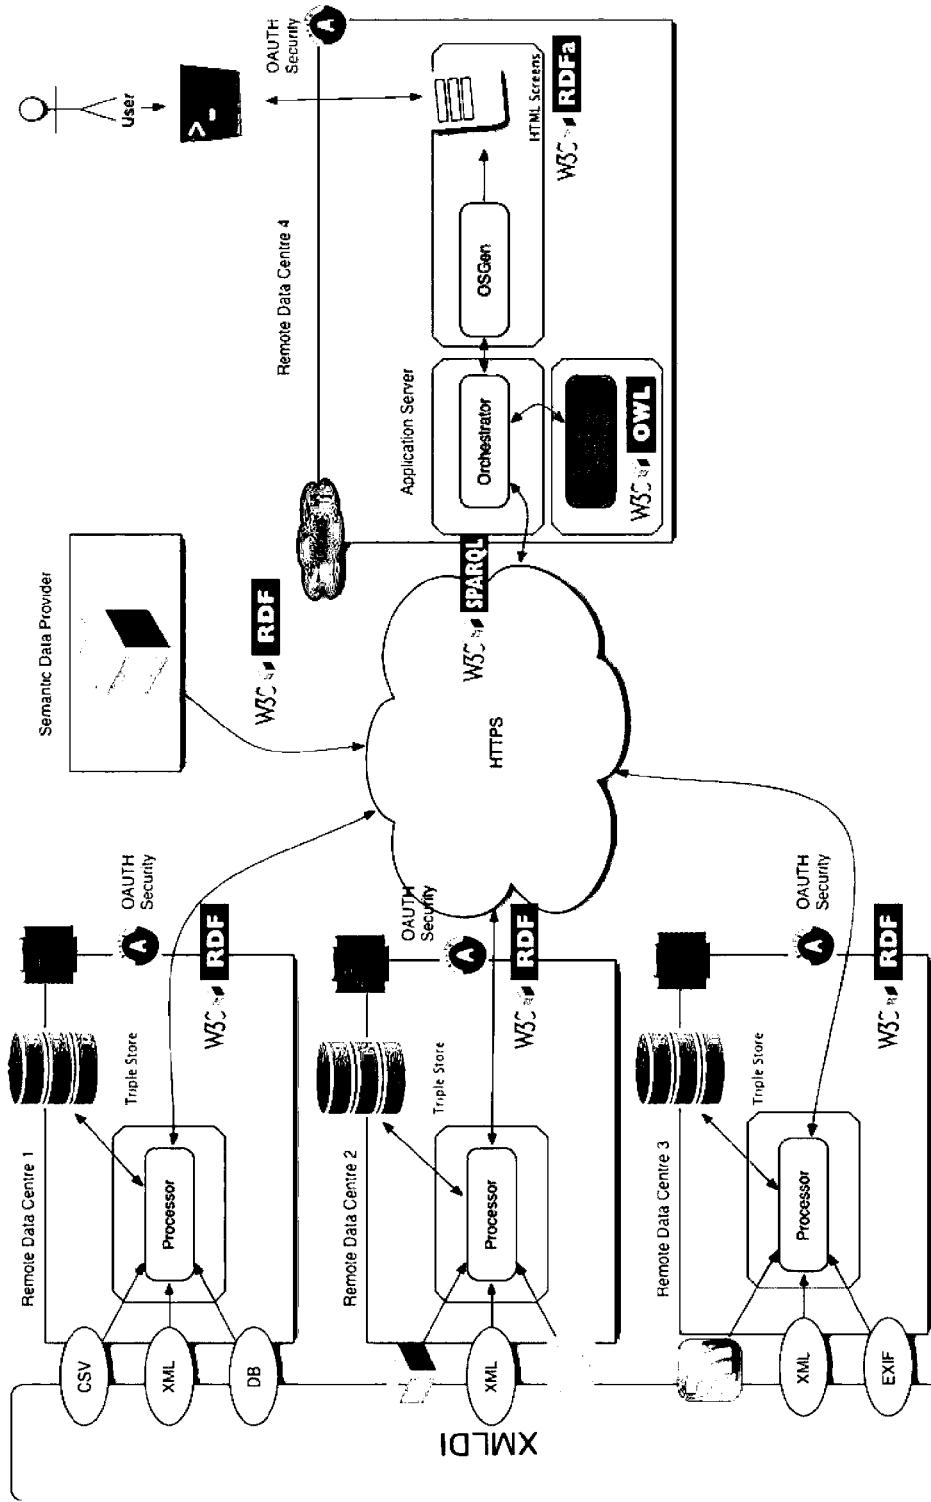
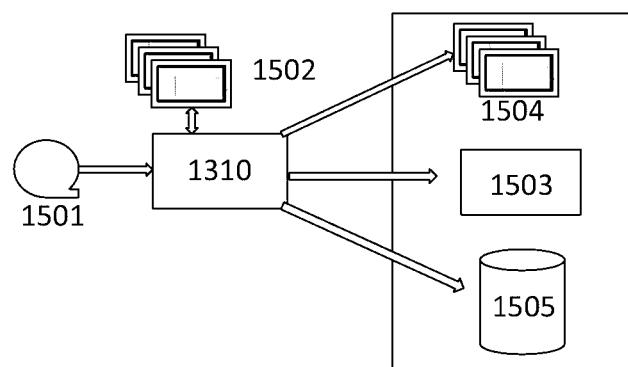
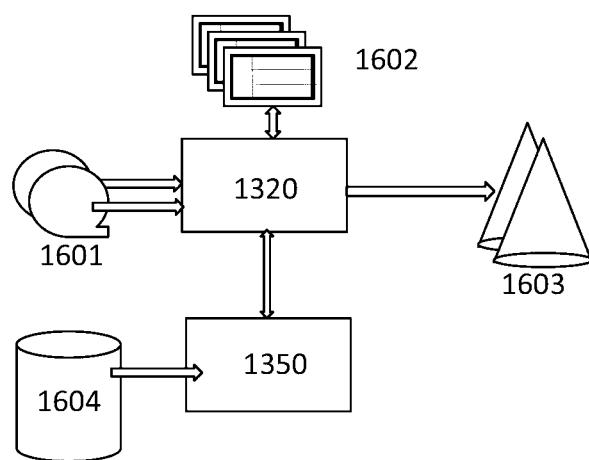


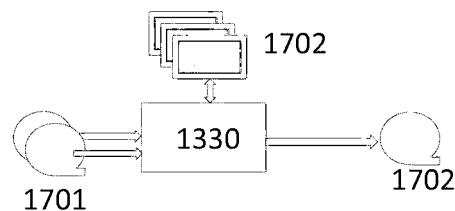
Fig. 14B



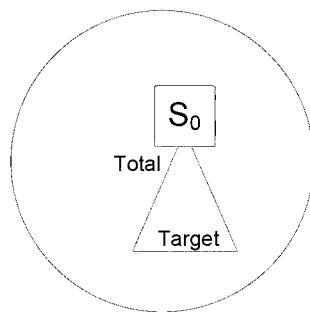
**Fig. 15**



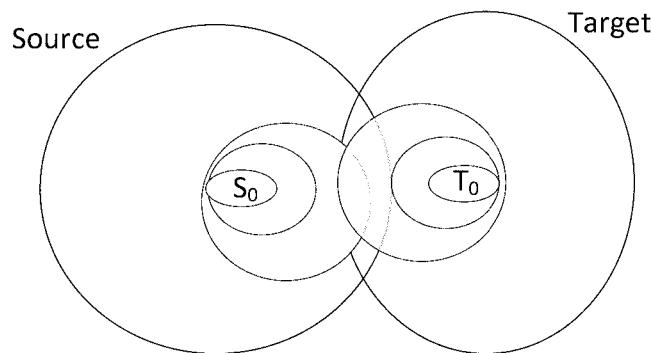
**Fig. 16**



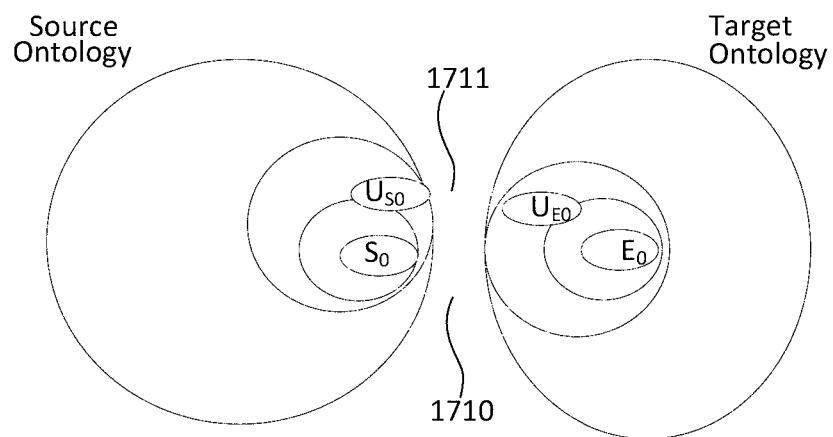
**Fig. 17A**



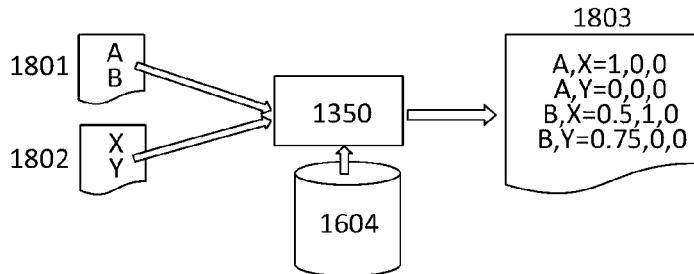
**Fig. 17B**



**Fig. 17C**

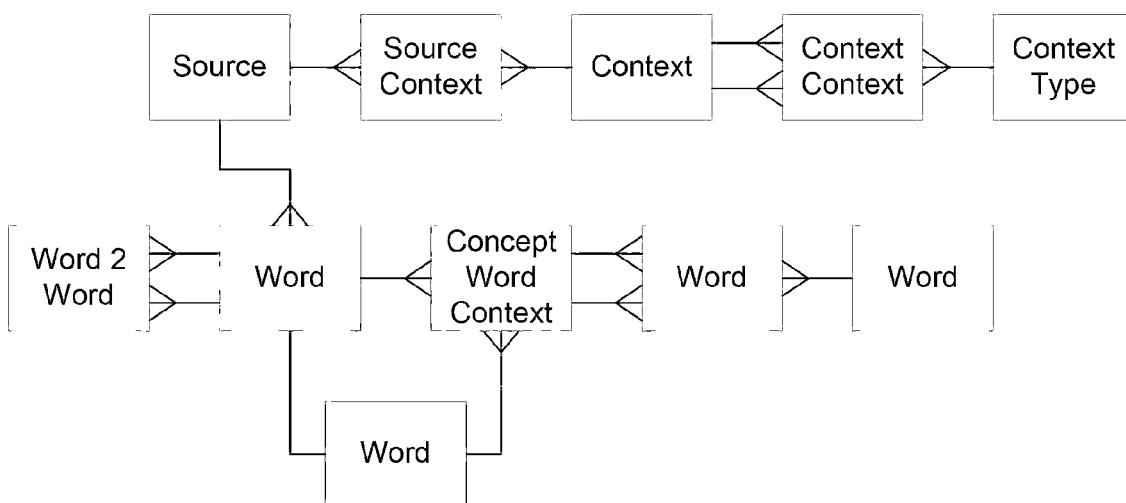
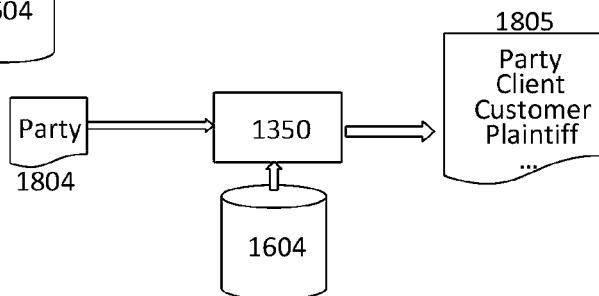


**Fig. 17D**

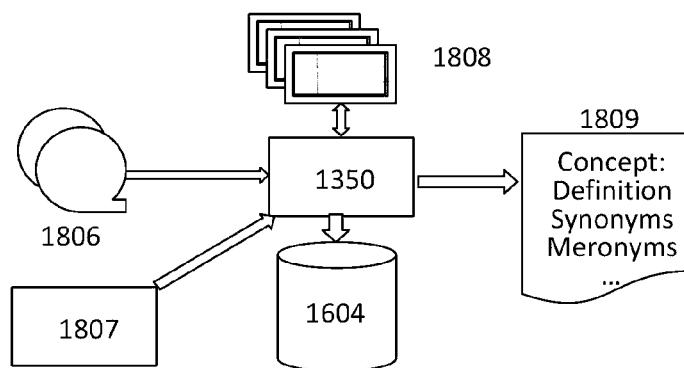


**Fig. 18A**

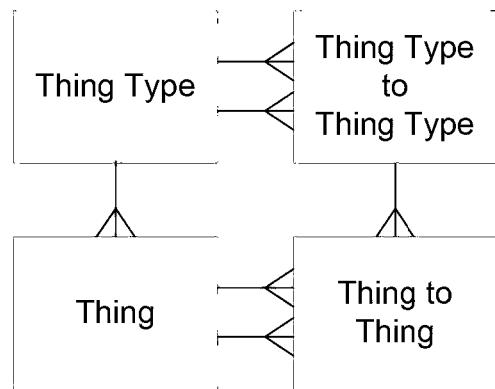
**Fig. 18B**



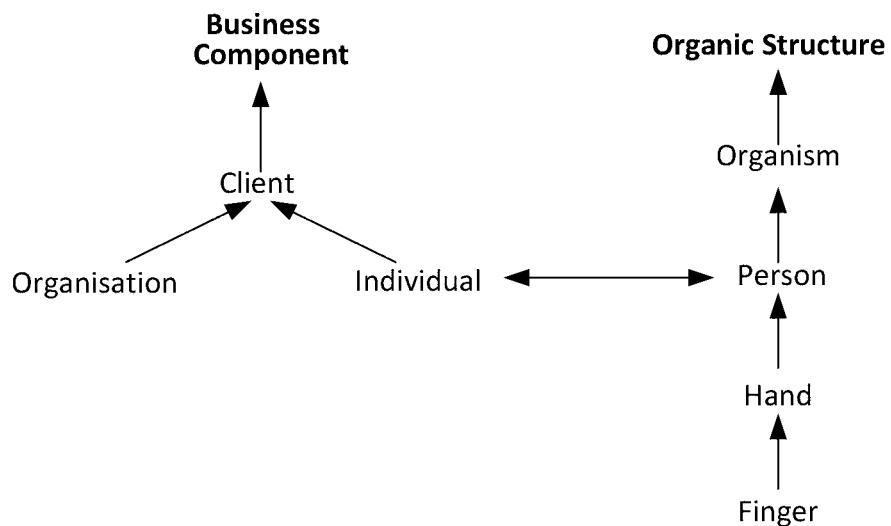
**Fig. 18C**



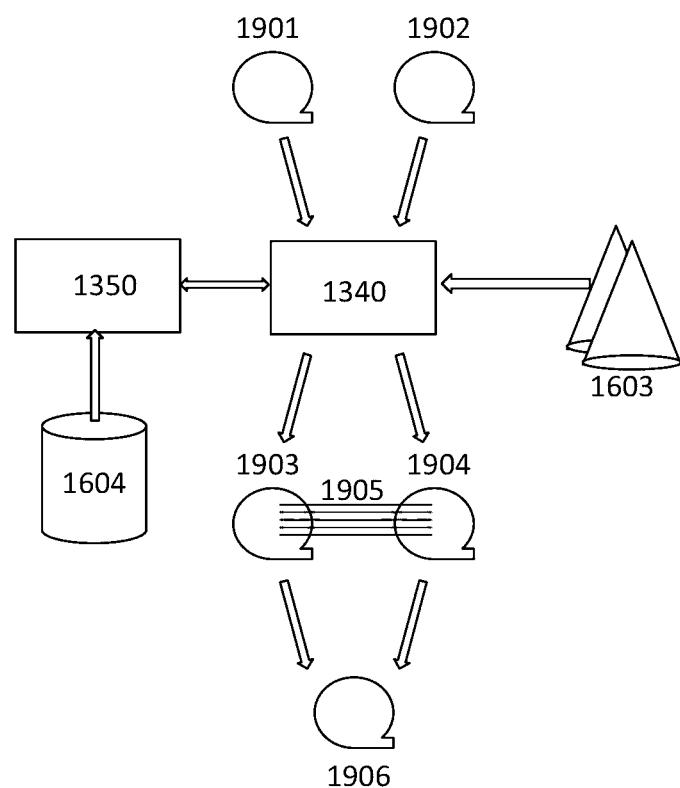
**Fig. 18D**



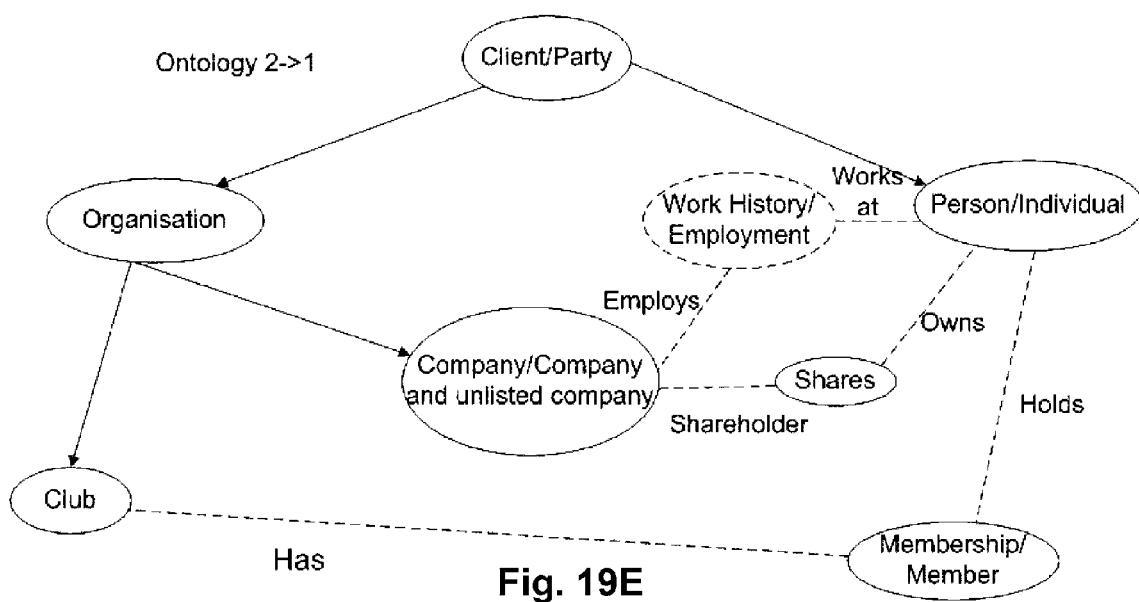
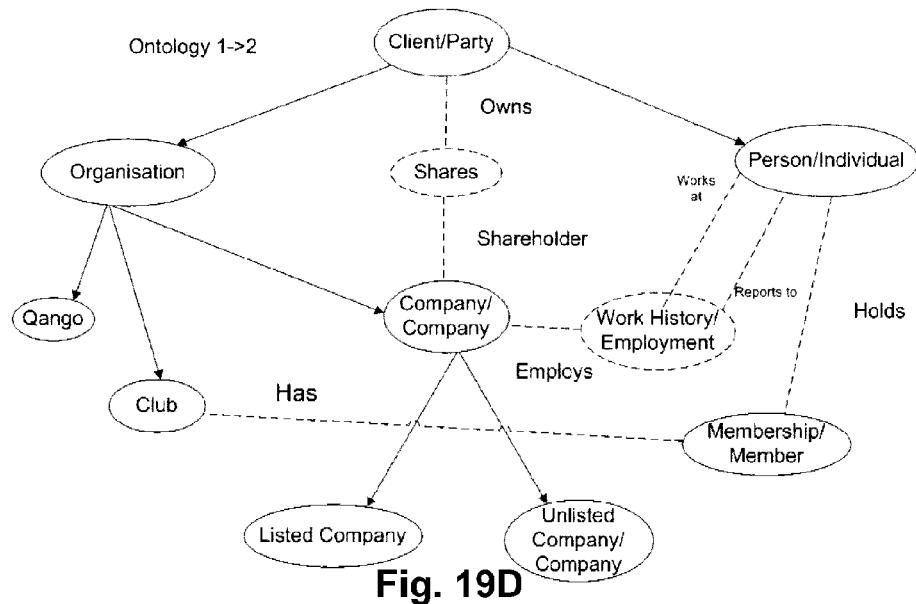
**Fig. 19A**



**Fig. 19B**



**Fig. 19C**



## ONTOLOGY MAPPING METHOD AND APPARATUS

### BACKGROUND OF THE INVENTION

[0001] The present invention relates to a method and apparatus for use in generating a mapping using ontologies.

### DESCRIPTION OF THE PRIOR ART

[0002] The reference in this specification to any prior publication (or information derived from it), or to any matter which is known, is not, and should not be taken as an acknowledgment or admission or any form of suggestion that the prior publication (or information derived from it) or known matter forms part of the common general knowledge in the field of endeavour to which this specification relates.

[0003] There are many thousands of public and private ontologies describing every aspect of the scientific, engineering and business worlds. This invention describes a set of processes which allow the knowledge and metadata in these advanced ontologies to be applied to the management of any data regardless of the data structure or lack of structure.

[0004] The explosive growth of knowledge and data is beyond the ability of traditional information management mechanisms to manage or even describe. Semantic Web technologies such as ontologies and new languages such as OWL (Web Ontology Language) and RDF (Resource Description Framework) enable the description of linked concepts such as health, medicine or engineering to be described in previously impossible detail and in a manner which is both human and machine understandable. These ontologies are typically created by teams of subject matter experts (ontologists) and are frequently publicly available.

[0005] Currently the manipulation of ontologies defining linked concepts is largely confined to academics and professional ontologists. These ontologies may contain many thousands of linked concepts. Removing even one concept, axiom or data property could render many of the relationships invalid. Determining the effects manually is a long and tedious process.

[0006] Among the business needs being addressed by this project is to enable a person with little or no ontological expertise to access all the details of an ontology in a simple comprehensible fashion. Furthermore that person will be able to select and inspect the data described by the ontology using a simplified query construction mechanism. They will be able to add records to the data with all the constraints and inferences which exist in the original ontology still being enforced. Finally they will be able to generate and deploy code and screens as stand-alone applications suitable for use by front-office personnel.

[0007] The particular business or research need may require concepts from multiple ontologies, possibly from disparate disciplines. In such cases an alignment between the concepts in the two ontologies may be required. These aligned ontologies may then be pruned to give a target ontology specific to the particular needs of the business or organisation.

[0008] This invention not only provides the alignment and pruning methods required, but also provides technical tools to perform semantic matching of concepts and broad naive indexes utilising synonyms and homonyms specific to the business context.

[0009] With this toolset a user can access disparate data sources, map them to context specific ontologies, align the disparate ontologies and then prune the aligned ontologies to contain only the relevant information. They can then extract the data into a format of their choosing and access them via the generated application which understands the axioms and inferences of the underlying ontologies and implements them through program code against the target database.

[0010] For example, a doctor could access patient records from multiple practice management systems, hospital and health systems plus diagnostic and imaging systems without being concerned about the various record layouts being accessed. A medical researcher could join such data to laboratory research data such as genome analysis, drug development and testing and so forth.

[0011] Traffic engineers can join signal data to barometric data, CCTV feeds, twitter and event calendars. Telecommunications companies can accept 99% of new Call Data Record structures without needing to write specific code for each new phone feature created by each manufacturer for each model in their range.

[0012] Finally, in the business sphere the time to market of new IT software is massively reduced by using an appropriately tailored ontology to describe the business problem and then to generate the code to execute the business rules described. These business rules can access all aspects of Big Data as well as the conventional data.

[0013] U.S. Pat. No. 7,464,099 provides a method of transferring content from a file and a database. In this case, the file includes content instances, each content instance being associated with a respective field, and each field having a respective type. The transfer is achieved by determining the type of each field, and then storing each content instance in a store in accordance with the determined field type of the associated field. Each content instance can then be transferred to the database in accordance with the determined field type. A similar procedure is provided for creating XML files based on content within the database.

### SUMMARY OF THE PRESENT INVENTION

[0014] In a First aspect, the present invention provides an apparatus for use in generating a mapping using ontologies, the apparatus including at least one electronic processing device that:

[0015] determines ontologies having respective ontology terms;

[0016] determines a group of ontology terms from at least one of the ontologies;

[0017] determines alignment between ontology terms in the ontologies for at least some of the group of ontology terms, the alignment being determined at least partially in accordance with an ontology term meaning of the ontology terms; and

[0018] generates a mapping in accordance with the alignment;

[0019] wherein the ontology term meaning is determined by comparing an ontology term with a number of potential ontology terms and generating a matching score for the results of each comparison, the ontology term meaning determining the matching score.

[0020] Preferably, the ontology term meaning is converted to a set of RDF triples and that set of RDF triples is compared with a set of other RDF triples for potential ontology terms and generating the matching score for the

comparison, the matching score determining the nature of a relationship between the ontology term meanings, thereby enabling semantic matching between the terms.

[0021] Preferably, the mapping may be at least one of:

[0022] a merged ontology; and  
[0023] an alignment index.

[0024] Preferably, the ontologies may include:

[0025] a putative ontology; and  
[0026] a formalized ontology.

[0027] Preferably, the electronic processing device may generate a mapping to map a putative ontology to a formalized ontology.

[0028] Preferably, the mapping may be used in transferring content from a source data store having a source data structure including a number of source data fields to a target data store having a target data structure including a number of target data fields, wherein the ontologies are associated with the source data structure and the target data structure and wherein the electronic processing device transfers content between the source data fields and target data fields using the mapping.

[0029] Preferably, the ontologies may include a source ontology associated with the source data structure and a target ontology associated with the target data structure.

[0030] Preferably, the source and target ontologies may be putative ontologies and the source and target ontologies are mapped to one or more formalized ontologies.

[0031] Preferably, the electronic processing device may determine an ontology by:

[0032] generating a putative ontology; and  
[0033] selecting one of a number of existing ontologies.  
[0034] Preferably, the electronic processing device may select an ontology using at least one of:  
[0035] metadata associated with a data structure; and  
[0036] data fields of a data structure.

[0037] Preferably, the electronic processing device may select one of a number of existing ontologies by:

[0038] comparing data fields to ontology terms of a number of existing ontologies; and  
[0039] selecting one of the number of existing ontologies in accordance with the results of the comparison.

[0040] Preferably, the electronic processing device may generate a putative ontology from a database schema by:  
[0041] identifying tables in the schema;

[0042] creating an ontology term corresponding to each table;  
[0043] identifying at least one bill of material table; and  
[0044] creating an ontology term corresponding to each entry in the bill of materials table.

[0045] Preferably, the electronic processing device may generate a putative ontology from a database schema by:

[0046] displaying an indication of the ontology term corresponding to each entry in the bill of materials table; and  
[0047] adding the ontology term to the putative ontology in response to user input commands.

[0048] Preferably, the electronic processing device may generate relationships between ontology terms using a table structure defined by the database schema.

[0049] Preferably, the putative ontology may include:

[0050] classes corresponding to ontology terms;

[0051] data properties for at least some of the classes; and  
[0052] object properties defining relationships between classes.

[0053] Preferably, the electronic processing device:

[0054] May determine an index for at least one ontology, the index including an indication of the ontology terms of the at least one ontology; and uses the index to determine:  
[0055] the group of ontology terms; and  
[0056] alignment between ontology terms.

[0057] Preferably, for each ontology term, the index may include an indication of:  
[0058] an ontology term meaning; and  
[0059] an ontology term type.

[0060] Preferably, the electronic processing device may generate an index by, for each ontology term:  
[0061] identifying an ontology term name;  
[0062] identifying an ontology term type;  
[0063] identifying an ontology meaning of each ontology term using a semantic matching process; and

[0064] creating an index entry including an indication of the ontology term name, the ontology term type and the ontology term meaning.  
[0065] Preferably, the electronic processing device generates a single index for multiple ontologies and wherein the index entry includes an indication of the ontology associated with the ontology term.

[0066] Preferably, the electronic processing device may:  
[0067] display an indication of one or more ontology terms in an ontology;

[0068] identify at least one ontology term in response to user input commands; and  
[0069] display details of at least one of:  
[0070] at least one identified ontology term; and  
[0071] data properties associated with the at least one identified ontology term.

[0072] Preferably, the electronic processing device may determine the indication of the one or more ontology terms using an index.

[0073] Preferably, the electronic processing device may:  
[0074] determine user selected ontology terms and corresponding user selected data properties in response to user input commands; and

[0075] generate executable code in accordance with the user selected ontology terms and corresponding user selected data properties, the executable code, when executed on a computer system, causes the computer system to display a user interface for allowing a user to interact with content stored in a data store having a data structure, the content being stored in data fields corresponding to the user selected ontology terms.

[0076] Preferably, the executable code may cause the computer system to generate queries for interacting with data stored in a source or target data structure in accordance with a source or target ontology.

[0077] Preferably, the executable code may cause the computer system to generate queries in accordance with at least one of data properties and relationships between the user selected ontology terms of an ontology.

[0078] Preferably, the executable code may cause the computer system to:

[0079] display an indication of one or more ontology terms;

[0080] determine selection of at least one ontology term in response to user input commands;

[0081] and query data stored in a data field associated with the selected ontology term.

[0082] Preferably, the electronic processing device may:  
[0083] compare an ontology term to a number of potential ontology term meanings; and  
[0084] select one of the potential ontology term meanings as the ontology term meaning in accordance with the results of the comparison.  
[0085] Preferably, the electronic processing device may determine an ontology term meaning by:  
[0086] generating a matching score for the results of each comparison; and  
[0087] determining an ontology term meaning in accordance with matching scores.  
[0088] Preferably, the electronic processing device may determine if a potential ontology term meaning is at least one of:  
[0089] a synonym;  
[0090] an antonym;  
[0091] a meronym;  
[0092] a superclass; and  
[0093] a subclass.  
[0094] Preferably, the electronic processing device may store an indication of the ontology term meaning in an index, the indication including at least one of:  
[0095] an identifier indicative of a defined meaning; and  
[0096] a list of equivalent meanings.  
[0097] Preferably, the electronic processing device may determine an alignment between ontology terms by:  
[0098] comparing ontology term meanings of a number of ontology terms;  
[0099] generating a matching score for the results of each comparison; and  
[0100] determining an alignment in accordance with matching scores.  
[0101] Preferably, the electronic processing device may further determine the alignment based on at least one of:  
[0102] relationships between ontology terms; and  
[0103] data properties of the ontology terms.  
[0104] Preferably, the electronic processing device may:  
[0105] determine relationships between source ontology terms in a source ontology;  
[0106] determine relationships between the target ontology terms in a target ontology;  
[0107] compare the relationships; and  
[0108] determine the alignment in accordance with the results of the comparison.  
[0109] Preferably, the electronic processing device may determine alignment between one or more source ontology terms from a source ontology and one or more target ontology terms from a target ontology.  
[0110] Preferably, the electronic processing device may determine the group of ontology terms by:  
[0111] determining selected ontology terms; and  
[0112] determining the group of ontology terms at least partially in accordance with the selected ontology terms and relationships between the selected ontology terms.  
[0113] Preferably, the electronic processing device may determine related ontology terms by:  
[0114] for each selected ontology term, identifying ontology terms within a defined relationship path length for each of a number of different types of relationship; and  
[0115] adding the identified ontology terms to the group of ontology terms.  
[0116] Preferably, the electronic processing device may repeatedly add identified ontology terms until all the selected ontology terms are connected by relationships.

[0117] Preferably, the electronic processing device may use a different relationship path length for different types of relationship.  
[0118] Preferably, the electronic processing device may determine a relationship path length in accordance with user input commands.  
[0119] Preferably, the apparatus may include:  
[0120] an indexer module that generates an index indicative of ontology terms in an ontology;  
[0121] a browser module that enables browsing of ontology terms in an ontology and generates code embodying at least part of the ontology thereby allowing a user to interact with data stored in a data structure in accordance with the ontology;  
[0122] an aligner module that determines alignment between ontology terms different ontologies;  
[0123] a pruner module that determines a group of ontology terms within at least one ontology at least in part using relationships between the ontology terms; and  
[0124] a semantic matcher module that identifies ontology term meanings.  
[0125] In a second aspect, the present invention provides a method for use in generating a mapping using ontologies, the method including, in at least one electronic processing device:  
[0126] determining ontologies having respective ontology terms;  
[0127] determining a group of ontology terms from at least one of the ontologies;  
[0128] determining alignment between ontology terms in the ontologies for at least some of the group of ontology terms, the alignment being determined at least partially in accordance with an ontology term meaning of the ontology terms; and  
[0129] generating a mapping in accordance with the alignment;  
[0130] wherein the ontology term meaning is determined by comparing an ontology term with a number of potential ontology terms and generating a matching score for the results of each comparison, the ontology term meaning determining the matching score.

#### BRIEF DESCRIPTION OF THE DRAWINGS

[0131] An example of the present invention will now be described with reference to the accompanying drawings, in which:  
[0132] FIG. 1A is a flow chart of an example of a method for use in generating a mapping using ontologies;  
[0133] FIG. 1B is a schematic diagram of examples of different mappings based on example ontology structures;  
[0134] FIG. 2 is a schematic diagram of an example of a distributed computer architecture;  
[0135] FIG. 3 is a schematic diagram of an example of a base station processing system;  
[0136] FIG. 4 is a schematic diagram of an example of a computer system;  
[0137] FIG. 5 is a flow chart of an example of a method for use in generating a mapping for transferring content between source and target data structures;  
[0138] FIG. 6 is a flow chart of an example of a method of generating a putative ontology;  
[0139] FIG. 7 is a flow chart of an example of a method of determining an index;

- [0140] FIG. 8 is a flow chart of an example of a method of browsing an ontology;
- [0141] FIG. 9 is a flow chart of an example of a method for pruning an ontology;
- [0142] FIG. 10 is a flow chart of a second example of a method for aligning ontologies;
- [0143] FIG. 11 is a flow chart of an example of a semantic matching method;
- [0144] FIGS. 12A and 12B are schematic diagrams of example ontologies;
- [0145] FIG. 13 is a schematic diagram of the modules used for interacting with ontologies;
- [0146] FIG. 14A is a schematic diagram of an example of the software stack of the ETL (Extraction Transformation Load) module of FIG. 13;
- [0147] FIG. 14B is a schematic diagram of an architecture used for implementing the ETL module if FIG. 13;
- [0148] FIG. 15 is a schematic diagram of an example of the functionality of the browser module of FIG. 13;
- [0149] FIG. 16 is a schematic diagram of an example of the functionality of the indexer module of FIG. 13;
- [0150] FIG. 17A is a schematic diagram of an example of the functionality of the pruner module of FIG. 13;
- [0151] FIGS. 17B to 17D are schematic diagrams of examples of a pruning process;
- [0152] FIG. 18A is a schematic diagram of a first example of the functionality of the semantic matcher module of FIG. 13;
- [0153] FIG. 18B is a schematic diagram of a second example of the functionality of the semantic matcher module of FIG. 13;
- [0154] FIG. 18C is a schematic diagram of an example of relationships between tables;
- [0155] FIG. 18D is a schematic diagram of a third example of the functionality of the semantic matcher module of FIG. 13;
- [0156] FIG. 19A is a schematic diagram of an example of a “thing database”;
- [0157] FIG. 19B is a schematic diagram of an example of a framework for unifying disparate sources;
- [0158] FIG. 19C is a schematic diagram of an example of the functionality of the aligner module of FIG. 13; and
- [0159] FIGS. 19D and 19E are schematic diagrams of examples of merged ontologies.

#### DETAILED DESCRIPTION OF THE PREFERRED EMBODIMENTS

- [0160] An example of a method of generating a mapping using ontologies will now be described with reference to FIG. 1A.
- [0161] For the purpose of this example, it is assumed that the process is performed at least in part using an electronic processing device, such as a microprocessor of a computer system, as will be described in more detail below.
- [0162] For at least some of the example, it is also assumed that content is stored as one or more content instances in content fields of a data store acting as a content repository, such as database or file. Thus, the content fields could be database fields of a database, with a content instance corresponding to a database record, including values stored across one or more database fields. Alternatively, content fields could be fields defined within a file, such as an XML file, which may be used for transporting data, for example, when data is to be extracted from and/or transferred to a

database, as will become apparent from the description below. As another alternative, content fields could be fields defined within a file, such as an RDF triple store, which may be used for transporting data, for example, when data is to be extracted from and/or transferred to a RDF triple store database, as will also become apparent from the description below. It is assumed that the content is stored in accordance with a data structure, such as a database schema, XML document definition, ontology or schema, or the like.

[0163] For the purpose of illustration throughout the following description, the term “source” is used to refer to a data store, such as a database or file from which data is being extracted, whilst the term “target” is used to refer to a data store, such as a database or file into which data is being stored. These terms are for the purpose of illustration only, for example to distinguish between possible sources and targets, and are not intended to be limiting.

[0164] The term “content instance” refers to an individual piece of content that is being extracted from a source and/or transferred to a target and is also not intended to be limiting. For example the term content instance could refer to a database record having values stored in a number of different database fields, or a set of related database records, or could alternatively refer to a single value stored within a single field.

[0165] The term “ontology” represents knowledge as a set of concepts within a domain, using a shared vocabulary to denote the types, properties and interrelationships of those concepts. Ontologies typically include a number of components such as individuals, classes, objects, attributes or the like and the term “ontology terms” is generally used to refer to these components and optionally specific ones of these concepts.

[0166] The term “meaning” is intended to refer to the semantic interpretation of a particular ontology term, content field name, or the like. The term meaning therefore encompasses the intended meaning of the ontology term or content field, for example to account for issues such as homonyms, synonyms, meronyms, or the like, as will be described in more detail below.

[0167] In this example, at step 100 the electronic processing device determines ontologies having respective ontology terms. This process can be performed in any appropriate manner, and can include having the electronic processing device select one or more of a number of existing ontologies, stored for example in one or more ontology databases, or could alternatively be achieved by generating putative ontologies. In one example, the ontologies selected correspond to a source ontology associated with a source data structure and a target ontology associated with the target data structure, although this is not essential.

[0168] Selection of ontologies could be achieved on the basis of the source or target data structure, and could include comparing fields within the data structures to ontology terms until a suitable match is found. This process could also involve examining domains of a number of ontologies and selecting domains, and hence ontologies that are of relevance to a particular subject matter field of the content, the industry to which the content relates, or the like. Selection of the ontologies can be automated, for example by providing the electronic processing device with an indication of the subject matter field of the relevant content; manual, for example by having the electronic processing device display

details of available ontologies allowing the user to select these, or through a combination of manual and automated processes.

[0169] Alternatively, this could involve generating a ‘putative’ ontology from the source or target data structure, for example using meta data, a database schema, or the like. This process can involve deriving some ontological axioms from relational referential integrity constraints, but most axioms would need to be manually added or ignored. This putative ontology may then be aligned with an existing rich ontology to add metadata.

[0170] At step 110, the electronic processing device determines a group of ontology terms from at least one of the ontologies. The group of ontology terms are typically a sub-group of the respective ontology and represent a group of related ontology terms. The group typically includes selected ontology terms, for example corresponding to ontology terms of interest. Thus, for example, when used in mapping source and target data structures, the group of ontology terms can include selected ontology terms corresponding to the source or target data fields that are to be mapped to the target or source data fields, for example to allow content therein to be extracted and transferred to the target data store. The group further typically includes ontology terms required to maintain relationships between the selected ontology terms.

[0171] The manner in which the group is determined will vary depending on the preferred implementation, and could include selecting one of a number of previous determined groups, for example on the basis of the content to be extracted. Alternatively however, this will involve generating the group by identifying the selected ontology terms and then examine related ontology terms progressively, until a path connecting each of the selected ontology terms is found. This process is sometimes referred to as pruning as it effectively involves pruning an ontology to a reduced number of ontology terms that can be more easily managed than the entire ontology, but which still maintains the structure and axioms of the ontology.

[0172] At step 120 the electronic processing device determines alignment between ontology terms for at least some of the group of ontology terms, with the alignment being determined at least partially in accordance with an ontology term meaning of the ontology terms. Thus, the electronic processing device examines at least the meaning of the ontology terms in one ontology and attempts to identify equivalent ontology terms in the other ontology. Thus at the simplest level, this could include identifying source and target ontology terms that have the same meaning, but more typically will examine relationships between the source ontology terms, and also between the target ontology terms, to identify if the source and target ontology terms also have similar relationships with other source and target ontology terms.

[0173] It should be noted that the alignment need not be a one-to-one mapping, and a single ontology term in one ontology may align with multiple ontology terms in the other ontology. For example, a source ontology may define an individual’s name as a single concept, whereas a target ontology may define first and last names as separate concepts, in which case content will need to be merged as it is transferred from the source to the target.

[0174] At step 130, the electronic processing device operates to generate a mapping in accordance with the alignment.

The mapping can represent a mapping between the ontologies themselves and/or can represent a mapping between source data fields and target data fields in respective source and target data structures, based on alignment between the target ontology terms and source ontology terms. Thus, for example, the electronic processing device can determine for the source data fields, corresponding source ontology terms, and then using the alignment, aligned target ontology terms, and hence the target data fields. The mapping is then typically stored in the form of a mapping file, index, table, merged ontology, or the like, allowing it to be subsequently used in transferring data.

[0175] Accordingly, the above process allows two ontologies to be aligned, with this in turn being used to create a mapping. The mapping can be used to equate ontology terms in different ontologies, allowing this in turn to be used to map source data fields to target data fields using respective ontologies. This mapping can be used in transferring content between a source and target, with the use of the ontologies helping to ensure that the structure of the content is maintained even when content is being transferred between data stores having different data structures. The use of ontologies can also help ensure that the integrity of content being transferred is being maintained or improved as it is being transferred, for example to ensure that the content is in the required normal form, and even providing normalisation if required. Thus, using ontologies allows additional relationship restrictions to be imposed as content is being transferred from the source to the target data store, so that even if data in the source data store is not stored in third normal form, it can be converted to third normal as it is transferred.

[0176] When the process is being used to create a mapping between source and target data structures, source and target ontologies can be previously defined ontologies, such as official created ontologies such as the Galen ontology (generally referred to as a formalised ontology), or can alternatively be putative ontologies, generated based on the data structure itself. This allows a number of different mapping relationships to be established as will now be described with reference to FIG. 1B.

[0177] In this example, source and target data stores having respective source and target data structures S, T are shown. Each data store can have a corresponding putative ontology 150, 160, with corresponding formalised ontologies 151, 161 also existing. This allows a number of different mappings to be established.

[0178] For example, mappings can be created directly between the putative ontologies 150, 160, allowing content to be transferred directly between the source and target S, T.

[0179] Additionally, and/or alternatively, a mapping could be created between a putative ontology 150, 160 and a corresponding formalised ontology 151, 161 using the above described techniques. This process allows a putative ontology 150, 160 to be created based on the source or target data structure S, T, and then substantially automatically mapped to a formalised ontology 151, 161 to thereby provide meaning, definitions and potentially relationships, to ontology terms in the putative ontology 150, 160. This process can also be used in normalising the data stored in the respective data structures.

[0180] In the event that the ontologies 151, 161 are identical, mapping the putative ontologies 150, 160 to the formalised ontologies 151, 161 allows for content to be transferred between the source and target S, T. Alternatively,

the formalised ontologies could also be aligned to allow for mappings therebetween to again allow transfer of content between the source and target S, T.

[0181] Thus, it will be appreciated that the above described technique can be used to create mappings between ontologies, with these in turn being used to establish a mapping directly between a source and target data structure. This facilitates transfer of content between disparate source and targets, and allows this to be performed in accordance with formalised data structures, thereby ensuring that content is provided in a normalised form. This therefore not only allows for content transfer between disparate data structures, but can also be used to overcome issues associated with data stored in a denormalised form.

[0182] A number of further features will now be described. In this regard, the electronic processing device is typically adapted to perform a number of different functions to facilitate the above described process, including generating an index of ontologies, allowing users to browse and interact with ontologies, align ontologies, prune ontologies, and interpret the meaning of ontology terms, as will now be further described.

[0183] As mentioned above, the mapping can be of any form. When the mapping is between ontologies, this can be in the form of a merged ontology or alternatively an alignment index, as will be described in more detail below.

[0184] The ontologies that are involved can include putative ontologies and/or formalized ontologies depending on the preferred implementation. In one particular example, the electronic processing device generates a mapping to map a putative ontology to a formalized ontology.

[0185] As also mentioned, the mapping can be used in transferring content from a source data store having a source data structure including a number of source data fields to a target data store having a target data structure including a number of target data fields, wherein the ontologies are associated with the source data structure and the target data structure and wherein the electronic processing device transfers content between the source data fields and target data fields using the mapping. Thus, the ontologies can include a source ontology associated with the source data structure and a target ontology associated with the target data structure, with the source and target ontologies being putative ontologies mapped to one or more formalized ontologies.

[0186] Accordingly, in the above described process, the electronic processing device can determine an ontology by generating a putative ontology or selecting one of a number of existing putative or formalised ontologies.

[0187] The electronic processing device typically determines a source or target ontology using either metadata associated with the source or target data structure or source and target data fields of the source or target data structure. This process could include generating a putative ontology or selecting one of a number of existing ontologies, for example from ontologies stored in a store such as an ontology database. In this latter case, the electronic processing device can compare the data structure data fields to ontology terms of a number of existing ontologies and select one of the number of existing ontologies in accordance with the results of the comparison.

[0188] Alternatively, when generating an ontology, for example from a database schema, the electronic processing device typically identifies tables in the schema, creates an ontology term corresponding to each table, identifies at least

one bill of material table and creates an ontology term corresponding to each entry in the bill of materials table. Thus, this process operates to examine the content of any denormalised database tables and expands the contents of this table to identify additional ontology terms. As part of this process, the electronic processing device can display an indication of the ontology term corresponding to each entry in the bill of materials table and add the ontology term to the putative ontology in response to user input commands. This allows the user to override the creation of ontology terms if required.

[0189] When generating a putative ontology, the electronic processing device can further generate relationships between ontology terms using a table structure defined by the database schema. This process allows the electronic processing device to generate a putative ontology including classes corresponding to ontology terms, data properties for at least some of the classes and object properties defining relationships between classes.

[0190] In one example, the electronic processing device determines an index for at least one ontology, the index including an indication of the ontology terms of the at least one ontology and uses the index to determine the group of ontology terms and alignment between ontology terms. Whilst the use of an index is not essential, it substantially reduces the amount of data that needs to be handled compared to use of an entire ontology, thereby making browsing, grouping (pruning) and alignment processes more manageable.

[0191] The index can be of any appropriate form but typically includes, for each ontology term, a name of the ontology term and an indication of an ontology term meaning and an ontology term type. The index may also include additional information such as an address for the ontology term in the respective ontology, which could take the form of a URI (Uniform Resource Identifier) or the like.

[0192] The electronic processing device generates the index by, for each ontology term, identifying an ontology term name, identifying an ontology term type, identifying an ontology meaning of each ontology term using a semantic matching process and creating an index entry including an indication of the ontology term name, the ontology term type and the ontology term meaning.

[0193] In one example, the electronic processing device generates a single index for multiple ontologies and wherein the index entry includes an indication of the ontology associated with the ontology term, although this is not essential and alternatively a separate index can be used for each ontology that is indexed.

[0194] The electronic processing device is typically further adapted to display an indication of one or more ontology terms in the ontology, identify at least one ontology term in response to user input commands and display details of at least one of at least one identified ontology term and data properties associated with the at least one identified ontology term. Thus, this allows the user to browse ontologies and view details of the ontology terms defined therein. This can be useful to allow users to assess whether the correct ontologies are being used, understand the scope of the ontologies, explore relationships between different ontology terms, and assess how a source or target data structure relate to the ontologies.

[0195] The electronic processing device typically determines the indication of the one or more ontology terms using

an ontology index. Thus, it will be appreciated that the index provides a rapid mechanism for the electronic processing device to display a list of ontology terms, and then subsequently explore data properties associated with selected ontology terms.

[0196] The electronic processing device is typically adapted to determine user selected ontology terms and corresponding user selected data properties in response to user input commands. This can be used not only to display details of the selected ontology terms and data properties, but also allows the electronic processing device to generate executable code.

[0197] The executable code is based on the user selected ontology terms and data properties, and when executed on a computer system, causes the computer system to display a user interface for allowing a user to interact with content stored in a data store having a data structure corresponding to the user selected ontology terms. Thus, this provides a mechanism for allowing the electronic processing device to automatically generate code that can be used to display an interface allowing users to interact with and subsequently export content from or import content to respective source and/or target data structures.

[0198] Thus, it will be appreciated that this allows a user to browse through the ontology terms within an ontology and then select ontology terms that correspond to data fields in a data structure. This in turn allows code to be generated which can act as an interface allowing the user to then interact with content stored within the data structure.

[0199] In one example, the executable code causes the computer system to generate queries for interacting with data stored in the source or target data structure in accordance with the source or target ontology, for example by querying and retrieving the content, for example using SPARQL (SPARQL Protocol and RDF Query Language). This provides a mechanism for rapidly deploying computer software that can act as an interface to a database. Furthermore, by virtue of this being generated based on the ontologies, it can integrate relationships defined within the ontology into the structure of the code and hence resulting queries. Thus, this allows the computer system to generate queries in accordance with data properties or relationships between the ontology terms selected by the user.

[0200] To achieve this, the computer system typically displays an indication of one or more ontology terms, for example from a source or target ontology, determines selection of at least one ontology term in response to user input commands, and queries data stored in a corresponding data field.

[0201] In order to determine meanings for ontology terms, and compare the similarity of the meaning of different ontology terms, the electronic processing device is typically adapted to perform semantic matching. This typically involves comparing an ontology term to a number of potential ontology term meanings, for example stored in a reference database, such as concept matching database, and selecting one of the potential ontology term meanings as the ontology term meaning in accordance with the results of the comparison.

[0202] To achieve this, the electronic processing device determines an ontology term meaning by generating a matching score for the results of each comparison and determining an ontology term meaning in accordance with matching scores. In determining the matching score, the

electronic processing device determines if a potential ontology term meaning is at least one of a synonym, an antonym, a meronym, a superclass and a subclass of the ontology term under consideration.

[0203] Once an ontology term meaning has been created, an indication of this is typically stored as part of an index, the indication including at least one of an identifier indicative of a defined meaning and a list of equivalent meanings. Thus, a single unique identifier can be assigned to each unique concept within a number of different ontologies, allowing common meanings to be identified by common identifiers across different ontologies. The identifier could be a unique alphanumeric code, lemma, root word, or the like, depending on the preferred implementation. This makes the process of identifying equivalent, and hence aligned, ontology terms far easier.

[0204] Similarly, the electronic processing device typically determines an alignment between ontology terms by comparing ontology term meanings of a number of ontology terms, generating a matching score for the results of each comparison and determining an alignment in accordance with matching scores. More typically the electronic processing device will also determine the alignment based on relationships between ontology terms or data properties of the ontology terms. Thus, this takes into account not only the absolute meaning of terms, but also the relative structures of the ontologies to determine if an alignment exists. Thus, the electronic processing device can determine relationships between the source ontology terms, determine relationships between target ontology terms, compare the relationships and determine alignment between source and target ontology terms in accordance with the results of the comparison.

[0205] The electronic processing device typically determines alignment between one or more source ontology terms from the source ontology and one or more target ontology terms from the target ontology, so this allows one-to-one, many-to-one and one-to-many relationships to be encompassed.

[0206] As mentioned above, the electronic processing device can determine a group of ontology terms, for example by pruning an ontology so that ontology terms that are irrelevant or unused for the current application can be removed and only those relevant to the current situation remain. In one example, this is achieved by determining selected ontology terms and then determining the group of ontology terms at least partially in accordance with the selected ontology terms and relationships between the ontology terms.

[0207] Thus, for each selected ontology term, the electronic processing device identifies ontology terms within a defined relationship path length for each of a number of different types of relationship and adds identified ontology terms to the group of source ontology terms. This process can be performed iteratively so that the electronic processing device repeatedly adds identified ontology terms until a desired end point is reached, such as when all the selected ontology terms are connected by relationships. This then establishes a group of related ontology terms that allow relationships within the ontology to be preserved and used in other processes, such as aligning ontologies, creating mappings, or the like.

[0208] As part of the above process, a number of different relationships can be examined in an attempt to establish paths between all of the selected ontology terms, with the

electronic processing device optionally using a different relationship path length for different types of relationship. Thus, for example, some types of relationships would be deemed more important, and hence would typically be given a longer path length when identifying relation ontology terms. In one example, the path length(s) could be defined by a user, allowing manual intervention to be used in controlling the creation of the group.

[0209] In one example, in order to allow the above described process to be performed, a number of different tools can be used to assist in generating the mappings and managing the ontologies. In one example the tools are provided as part of a software suite forming an integrated package of ontology and data management tools. In one example, the tools include an indexer module that generates an index indicative of ontology terms in an ontology, a browser module that enables browsing of ontology terms in an ontology and generates code embodying at least part of the ontology thereby allowing a user to interact with data stored in a data structure in accordance with the ontology, an aligner module that determines alignment between ontology terms different ontologies, a pruner module that determines a group of ontology terms within at least one ontology at least in part using relationships between the ontology terms and a semantic matcher module that identifies ontology term meanings. However, the use of respective modules is not essential and other arrangements can be used.

[0210] In one example, the processes can be performed at least in part using a processing system, such as a suitably programmed computer system. This can be performed on a standalone computer, with the microprocessor executing applications software allowing the above described method to be performed. Alternatively, the process can be performed by one or more processing systems operating as part of a distributed architecture, an example of which will now be described with reference to FIG. 2.

[0211] In this example, two base stations 201 are coupled via a communications network, such as the Internet 202, and/or a number of local area networks (LANs) 204, to a number of computer systems 203. It will be appreciated that the configuration of the networks 202, 204 are for the purpose of example only, and in practice the base station 201, computer systems 203 can communicate via any appropriate mechanism, such as via wired or wireless connections, including, but not limited to mobile networks, private networks, such as an 802.11 networks, the Internet, LANs, WANs, or the like, as well as via direct or point-to-point connections, such as Bluetooth, or the like.

[0212] In one example, each base station 201 includes a processing system 210 coupled to a database 211. The base station 201 is adapted to be used in managing ontologies, for example to perform pruning or alignment, as well as generating mappings for example for use in transferring content between source and target data stores. The computer systems 203 can be adapted to communicate with the base stations 201 to allow processes such as the generation of mappings to be controlled, although this is not essential, and the process can be controlled directly via the base stations 201.

[0213] Whilst each base station 201 is shown as a single entity, it will be appreciated that the base station 201 can be distributed over a number of geographically separate locations, for example by using processing systems 210 and/or databases 211 that are provided as part of a cloud based environment. In this regard, multiple base stations 201 can

be provided each of which is associated with a respective data stores or ontology, although alternatively data stores could be associated with the computer systems 203.

[0214] However, the above described arrangement is not essential and other suitable configurations could be used. For example, the processes could be performed on a stand-alone computer system.

[0215] An example of a suitable processing system 210 is shown in FIG. 3. In this example, the processing system 210 includes at least one microprocessor 300, a memory 301, an input/output device 302, such as a keyboard and/or display, and an external interface 303, interconnected via a bus 304 as shown. In this example the external interface 303 can be utilised for connecting the processing system 210 to peripheral devices, such as the communications networks 202, 204, databases 211, other storage devices, or the like. Although a single external interface 303 is shown, this is for the purpose of example only, and in practice multiple interfaces using various methods (e.g. Ethernet, serial, USB, wireless or the like) may be provided.

[0216] In use, the microprocessor 300 executes instructions in the form of applications software stored in the memory 301 to allow index generation, mapping and content transfer to/from the database 211 to be performed, as well as to communicate with the computer systems 203. The applications software may include one or more software modules, and may be executed in a suitable execution environment, such as an operating system environment, or the like.

[0217] Accordingly, it will be appreciated that the processing system 210 may be formed from any suitable processing system, such as a suitably programmed computer system, PC, database server executing DBMS, web server, network server, or the like. In one particular example, the processing system 210 is a standard processing system such as a 32-bit or 64-bit Intel Architecture based processing system, which executes software applications stored on non-volatile (e.g. hard disk) storage, although this is not essential. However, it will also be understood that the processing system could be any electronic processing device such as a microprocessor, microchip processor, logic gate configuration, firmware optionally associated with implementing logic such as an FPGA (Field Programmable Gate Array), or any other electronic device, system or arrangement.

[0218] As shown in FIG. 4, in one example, the computer system 203 includes at least one microprocessor 400, a memory 401, an input/output device 402, such as a keyboard and/or display, and an external interface 403, interconnected via a bus 404 as shown. In this example, the external interface 403 can be utilised for connecting the computer system 203 to peripheral devices, such as the communications networks 202, 204, databases 211, other storage devices, or the like. Although a single external interface 403 is shown, this is for the purpose of example only, and in practice multiple interfaces using various methods (e.g. Ethernet, serial, USB, wireless or the like) may be provided.

[0219] In use, the microprocessor 400 executes instructions in the form of applications software stored in the memory 401 to allow communication with the base station 201, for example to allow an operator to provide control inputs.

[0220] Accordingly, it will be appreciated that the computer systems 203 may be formed from any suitable pro-

cessing system, such as a suitably programmed PC, Internet terminal, lap-top, hand-held PC, smart phone, PDA, web server, or the like. Thus, in one example, the processing system 100 is a standard processing system such as a 32-bit or 64-bit Intel Architecture based processing system, which executes software applications stored on non-volatile (e.g. hard disk) storage, although this is not essential. However, it will also be understood that the computer systems 203 can be any electronic processing device such as a microprocessor, microchip processor, logic gate configuration, firmware optionally associated with implementing logic such as an FPGA (Field Programmable Gate Array), or any other electronic device, system or arrangement.

[0221] Examples of the operation of the system to generate mappings, allow browsing, indexing of and interaction with ontologies, including aligning and pruning ontologies will now be described in further detail.

[0222] For the purpose of these examples, it is assumed that the processing system 210 of the base station 201 hosts applications software for performing the processes, with actions performed by the processing system 210 being performed by the processor 300 in accordance with instructions stored as applications software in the memory 301 and/or input commands received from a user via the I/O device 302, or commands received from the computer system 203. In this regard, for the purpose of the following examples, the processing system 210 executes applications software having a number of modules including an indexer module, a browser module, an aligner module, a pruner module, a semantic matcher module and an ETL module. However, the use of respective modules is not essential and other arrangements can be used.

[0223] It will also be assumed that the user interacts with applications software executed by the processing system 210 via a GUI, or the like, presented either on the input/output device 302 or the computer system 203. Actions performed by the computer system 203 are performed by the processor 400 in accordance with instructions stored as applications software in the memory 401 and/or input commands received from a user via the I/O device 402. The base station 201 is typically a server which communicates with the computer system 203 via the particular network infrastructure available, and may for example be in the form of an enterprise server that interacts with a database 211 for users of one or more computer systems 203.

[0224] However, it will be appreciated that the above described configurations are for the purpose of example only and are not intended to be limiting, so in practice any database management system can be used. It will also be appreciated that the partitioning of functionality between the computer system 203, and the base station 201 may vary, depending on the particular implementation.

[0225] An overview of the process for determining a mapping and using this to transfer content from a source to a target will now be described with reference to FIG. 5. For the purpose of this example it will be assumed that the processing system 210 implements a number of different modules for providing different functionalities.

[0226] In this example, at step 500, the processing system 210 identifies source and target ontologies using the source and target data structures. This can be achieved in any manner, but typically involves creating a putative ontology based on the source and target data structures for source and target data stores. For example, the names of the different

source and target data fields could be equated to ontology terms, with relationships between the ontology terms being identified from the relationships in the source and target data structures. A specific example of the process of generating putative ontologies will be described in more detail with reference to FIG. 6.

[0227] At step 510, the indexer module determines an index of source and target ontologies. The index is typically in the form of a list including an entry indicative of each ontology term, an associated ontology term type if this is known, and also optionally an ontology term meaning. In this regard, the ontology term meanings are typically determined by the semantic matcher module at step 520 that compares the ontology term to a concept matching database, and uses the results of the comparison to identify a meaning for each ontology term in the index.

[0228] At step 530, the browser module is used to browse an ontology and select source or target ontology terms. This allows a user to select those ontology terms that are of interest, typically corresponding to content to be extracted from the source data store or imported into the target data store.

[0229] The selected ontology terms can then be used at step 540 to allow the browser module to generate code for interacting with content stored in a data store in accordance with the respective data structure. In particular, this can include code for allowing a computer system to generate a user interface which the user can use to review data fields of the data structure, select content to be extracted/imported and then generate the necessary queries to perform the extraction/importation, as will be described in more detail below.

[0230] Alternatively, at step 550, the selected ontology terms are used by the pruner module to prune either the source and/or target ontology. In particular, this allows the user to select only those parts of the ontology that are of interest, with the processing system 210 then selecting additional ontology terms required to maintain relationships between the selected ontology terms as will be described in more detail below.

[0231] Once one or more of the ontologies have been pruned, at step 560, the processing system 210 uses the aligner module to align the source and target ontologies. This identifies a correspondence between one or more of the source ontology terms and one or more of the target ontology terms, thereby allowing a mapping between the source and target data structures to be determined at step 570, which in turn can be used together with code generated by the browser module to transfer content from the source data store to the target data store.

[0232] An example of the process for generating a putative ontology from a data structure, such as a database schema or the like, will now be described with reference to FIG. 6.

[0233] Whilst this example is specific to generating a putative ontology for a relational database, it will be appreciated that similar concepts can be applied to other data structures, and that this example is for the purpose of illustration only and is not intended to be limiting.

[0234] In this example, at step 600, the processing system 210 determines each table in the database, typically by extracting this information from metadata defining the database schema. At step 610, the processing system 210 defines a class corresponding to each table in the database. In this regard, the term class refers to a specific ontology term

corresponding to a concept within the ontology, as will be described in more detail below.

[0235] At step 620, the processing system 210 identifies any database tables having a BOM (Bill Of Materials) structure or a Type structure. In this regard, a BOM table has two “one to many” relationships and is used to list of all parts constituting an item, object or article. The Type structure has one “many to one” relationship and has only one relevant attribute or column which is used to limit the range of values in the related table. Such tables are often used to de-normalise data and can therefore contain many concepts or classes that should each represent a respective ontology term. Accordingly, at step 630, the processing system expands each Type table and each BOM table to define further classes corresponding to each unique entry in the table.

[0236] At step 640, the processing system 210 optionally displays each identified class from within the Type or BOM table, allowing a user to confirm whether the class should be retained at step 650. If it is indicated that the Type or BOM class should not be retained, it is removed at step 660.

[0237] Once the relevant BOM classes have been selected, the processing system 210 defines relationships and attributes (also referred to as data objects and data properties) based on the database schema. Thus, the table structure can be used to identify relationships between the identified classes, whilst data fields in the tables are used to identify attributes of the classes. The relationships and attributes are in turn used to define object properties and data properties in the ontology, thereby allowing the putative ontology to be generated and saved, for example in an ontology database at step 680.

[0238] Thus, this allows a putative ontology to be created in a substantially automated fashion solely from an analysis of the data structure of a data store, such as a database, structured file, or the like. Following this, in the event that it is required to define meanings for the different classes within the putative ontology, the putative ontology can be aligned with a formalised ontology, as will be described in more detail below.

[0239] An example of the process for generating an index will now be described with reference to FIG. 7.

[0240] In this example, at step 700 the indexer module determines an ontology of interest. This may be determined based on user input commands, for example supplied via the browser module, or could be received from another module requiring an index. For example, an ETL module that has generated a putative ontology may require this be indexed and provide an indication of the ontology to the indexer module, or alternatively, a pruner module may request an index allowing pruning to be performed on an ontology.

[0241] At step 705, the indexer module compares the ontology to one or more existing indexes, typically stored in an index database, and determines if an index already exists. This can be achieved by comparing metadata associated with the ontology, such as an ontology name and/or address, with corresponding information associated with the indexes, or alternatively by comparing one or more ontology terms to ontology terms in existing indexes.

[0242] If it is determined that an index exists at step 710, then the index is provided at step 715, for example by providing the index to the module that requested the index. Otherwise, the index must be generated, in which case the indexer module selects a next ontology term at step 720, and

then creates an index entry including an indication of the ontology term name, an ontology term type and an ontology term address, typically indicative of a URI (Uniform Resource Identifier) or similar, at step 725. At step 730, the indexer module obtains a semantic meaning for the ontology term from a semantic matcher module, as will be described in more detail below, and adds this to the index entry.

[0243] At step 735, the indexer module determines if all ontology terms have been completed and if not the process returns to step 720, allowing a next ontology term to be selected. Otherwise, at step 740, the index is stored and optionally provided to another module.

[0244] An example of a process for browsing of an ontology will now be described with reference to FIG. 8.

[0245] In this example, at step 800, the browser module generates an ontology term list for a selected ontology, using an ontology index. Accordingly, as part of this process, the browser module can request the ontology index from the indexer module, for example based on the identity of a selected ontology. The ontology term list can then be displayed to a user via an appropriate GUI (graphical user interface).

[0246] A step 805, the user tags one or more ontology terms of interest, before selecting a next ontology term to view at step 810 allowing the browser module to display a ontology term screen including data properties for the selected ontology term at step 815. In this regard, the data properties correspond to attributes of the ontology term, which are defined as part of the ontology.

[0247] At step 820, the browser module determines if a search option has been selected by the user, in which case the user enters search terms in the data fields of the data properties at step 825. The browser module then generates and performs a query of data associated with the respective ontology term data properties, returning and displaying results to the user at step 830. Thus, this process allows the user to review the content that would be associated with respective data properties in the corresponding source or target data store, thereby allowing the user to ascertain whether the ontology term and associated data properties are of interest.

[0248] Once the search has been performed, or in the event that no search is performed, the user tags one or more data properties of interest at step 835. Thus, this process allows the user to review the ontology terms and associated data properties and then select ontology terms and data properties of interest by tagging them.

[0249] At step 840, the ontology terms are reviewed to determine if all ontology terms and data properties of interest to the user have been selected. If not, the process returns to step 810 allowing further ontology terms to be reviewed.

[0250] Otherwise, at step 845 the browser module selects the tagged ontology terms and associated data properties, allowing these to be used in other processes, such as to perform pruning at step 850 or to generate an application at step 855. In this regard, generation of an application involves uses scripts or the like to generate executable code, that when executed on a computer system allows the computer system to display a user interface for interacting with content in fields in the source or target corresponding to the selected ontology terms or data properties, as will be described in more detail below.

[0251] Thus, the above described process can be used to allow a user to browse ontology terms and associated data properties to identify which of those are of interest in respect of the content they wish to export from a source or import into a target.

[0252] An example of the process for pruning an ontology will now be described with reference to FIG. 9.

[0253] In this example, at step 900, the selected ontology terms are added as seeds for the pruning process. Following this, an iterative process is performed to repeatedly explore ontology terms related to the seed ontology terms until a path is identified that interconnects the seed ontology terms. To achieve this, at step 905, different types of relationships and associated default path lengths are displayed. In this regard, ontology terms can be related by different types of relationships, such as parent, child, sibling, or the like. As certain types of relationship may be more important than others, different relationship types may have different lengths. Additionally, the length of path that is explored for each type of relationship can be varied thereby ensuring that a larger number of ontology terms connected to the seed ontology terms via the more important relationships are included. Accordingly, at step 910, the user can adjust the path lengths for the different relationships, thereby allowing the pruning process to be tailored by the user, for example to control the extent and/or direction of pruning.

[0254] At step 915, ontology terms related to the selected ontology terms are determined, by identifying those ontology terms related by relationships of the specified path length. At step 920, the pruner module determines if the selected seed terms are linked. In other words there is a series of interconnected ontology terms that links the seed ontology terms, and if so, the pruning process can end with the selected and related ontology terms identified being used to define the pruned ontology at step 925, which can be stored as a pruned ontology or pruned index.

[0255] Otherwise, at step 930 it is determined if the iterations are complete, and if not the related ontology terms are added the selected ontology terms and the process returns to step 915, allowing further related ontology terms to be identified. Thus, the number of ontology terms related to the seed ontology terms is gradually increased until the seed ontology terms are connected by a path of relationships.

[0256] Thus, the above described process is repeated either until the ontology is successfully pruned, at which time the seed ontology terms are interconnected via a path of related ontology terms, or until a predetermined number of iterations are completed and no path is identified, in which case the process is halted at step 940. In this latter case, this typically suggests that the ontology terms are from different ontologies, in which case the pruning process is performed in conjunction with an alignment process, allowing the pruning process to span multiple ontologies as will be described in more detail below. Alternatively, this indicates that the ontology terms cannot be easily linked.

[0257] An example of the process for aligning source and target ontologies will now be described with reference to FIG. 10.

[0258] In this example, at step 1000 source and/or target ontology terms are selected using the index. This may involve having the user select ontology terms using the browser module, or more typically select two pruned ontologies corresponding to pruned versions of source and target ontologies that contain source and/or ontology terms of

interest. At step 1005, the matcher module is used to determine a matching score for different combinations of pairs of source and target ontology terms. These scores are used to define preliminary alignments solely based on how similar the meanings of the source and target ontologies are at step 1010.

[0259] At step 1015, the aligner module examines relationships (object properties) and attributes (data properties) of the source and target ontology terms to determine whether the preliminary alignments are correct. Thus, for example, this will examine if preliminary aligned source and target ontology terms have similar number of attributes, and also if these have similar relationships with other source or target ontology terms. This can be used to identify inexact matches, for example if each of the terms first name and last name may be preliminary matched to name, with the examination of the relationships being used to demonstrate this should be a many to one relationship.

[0260] At step 1020, this can be used to refine the alignments, allowing these to be stored to represent the alignment between the source and target ontologies at step 1025. This can be in the form of a merged ontology, or alternatively an alignment index.

[0261] An example of a semantic matching process will now be described with reference to FIG. 11.

[0262] In this example, at step 1100, the matcher module receives ontology terms for matching. This could be based on user selection via the browser module, but more typically is by receiving terms from the indexer module or the aligner module. At step 1105, a next pair combination is selected, either by comparing a single ontology term to a plurality of respective terms in a matching database, or by selecting a next pair of received source and target ontology terms.

[0263] At step 1110, the semantic matcher module calculates a semantic similarity using a concept matching database. The score can be determined in any one of a number of manners, but typically involves applying a predetermined formula that calculates a score based on whether the meanings are in any way related, such as whether they are antonyms, synonyms, or the like. In one particular example, this involves matching ontology terms with definitions, for example using a dictionary, such as WordNet, or the like. In this regard, WordNet is a large lexical database of English. Nouns, verbs, adjectives and adverbs are grouped into sets of cognitive synonyms (synsets), each expressing a distinct concept and is described in Fellbaum, Christiane (2005). WordNet and wordnets. In Brown, Keith et al. (eds.), Encyclopedia of Language and Linguistics, Second Edition, Oxford: Elsevier, 665-670.

[0264] Once a definition has been identified, this is expressed in terms of RDF triples, which are then stored in a database. The RDF triples for two different meanings can then be queried to determine a similarity between the triples, which is used to determine a similarity score indicative of the similarity of the meaning of the two ontology terms.

[0265] Following this, at step 1115, the semantic matcher module determines whether the terms are related by subclass and superclass arrangements. This information is then combined with the similarity score to calculate a matching score at step 1120. At step 1125, it is determined if all pairs are completed and if not the process returns to step 1105 allowing a next pair of source and target ontologies to be selected and a matching score is calculated. Once all potential pairs of ontology terms or ontology terms and matching

concepts in the database have been checked, the semantic matcher module can select the best match and then provide an indication of this at step 1130.

[0266] Accordingly, it will be appreciated that the above described processes allow users to interact with ontologies, select ontology terms of interest and use this to generate software for interacting with content stored in a data store, such as a database or XML file, in accordance with a respective ontology. The users can further investigate the ontology and then prune this using a pruner module, allowing a minimal ontology to be determined which allows the user to interact with content of interest. The pruned ontology can then be aligned with another pruned ontology, so that this can be used to define a mapping therebetween, which can in turn be used to transfer data between data stores having a source and target data structure.

[0267] A more specific example will now be described. For the purpose of this example, an ontology is defined as follows:

[0268] A set of related Concepts, also called Classes or Objects, some of which are related to each other using sub/super class relationships also called ‘inheritance’ relationships. Examples are ‘Organisation’, ‘Company’, ‘Club’ which display inheritance and ‘Land Mass’, ‘Gender’, ‘Person’ which do not display inheritance.

[0269] A set of Object Properties, which provide an additional mechanism for relating Classes. For example ‘is Located at/in’ has Gender’. These relationships allow inferencing of concepts, relationships and properties

[0270] A set of Data Properties associated with each Class. For example the class “Person” may have Data Properties of Name, Title, Date-of-Birth, and Gender.

[0271] A set of axioms providing a formulaic relationship between any of the preceding properties. For example, “if a Person has a Title of ‘Mrs’ then the gender must be female” or “if two objects have the same unique identifier then they are the same object”. These axioms allow further inferencing of concepts, relationships and properties.

[0272] An ontology can be described in a number of languages such as RDFS, XML, DAML, OIL, N3 and OWL. These languages may have different dialects such as OWL-Lite or OWL-DL. From a functionality perspective they differ in their ability to manage and describe complex relationships and axioms.

[0273] An ontology may contain hundreds of thousands of concepts. A user may be interested in a subset of these concepts. This subset may be from:

[0274] a single ontology;

[0275] multiple overlapping ontologies; or multiple disparate ontologies.

[0276] Some concepts in a target ontology may not be pre-defined, and may not exist in any of the source ontologies. In such a case the user may need to manually add the missing concepts. The required subset may have both or either starting and ending concepts.

[0277] For the purpose of illustration two extremely simple example ontologies are shown in FIGS. 12A and 12B. It will be appreciated that these are utilised to illustrate the processes of indexing, pruning semantic matching and alignment and are not intended to be limiting.

[0278] In these examples, there are two types of relationships, those which are hierarchically connected and those which are not. In these examples, hierarchically connected classes are represented by solid ellipses, which are hierar-

chically connected by solid lines pointing from the superclass to the subclass. Each subclass inherits all the properties of its superclass. The non-hierarchically connected set of classes, shown as broken ellipses, are connected to any class by a named Object Property line shown here as a dashed line. Each class has a set of data properties some of which are shown in Table 1 for illustration.

TABLE 1

Class Name 1	Data Property 1	Class Name 2	Data Property 2
Party 1.0	Name	Client 2.0	Name
Individual 1.1	First Name	Person 2.1	Given Name
	Last Name		Family Name
	Date of Birth		Date of Birth
	Gender		Gender
Organisation 1.2	Date of Incorporation or founding	Organisation 2.2	Date of Incorporation or founding
Club 1.3	Type	Club 2.3	Type
Company 1.4	Type	Company 2.4	Type
	Registered in/on	Listed Company 2.5	Stock Exchange Registered in
		Unlisted Company 2.6	
		Qango 2.7	Ministry
Member 1.5	Type	Membership 2.8	Type
	Joined Date		Joined Date
	Exit Date		Exit Date
Employment 1.6	Role	Work History 2.9	Role
	Start Date		Start Date
	End Date		End Date
	Reports to		Reports to
Shares 1.7	Purchase Date	Shares 2.10	Purchase Date
	Number		Number
	Type		Type

[0279] It will be appreciated that the ontologies show similar concepts, but that there are some differences:

[0280] Some concepts have different names. Can we say that a ‘Party’ is identical to a ‘Client’, a ‘Person’ is identical to an ‘Individual’, ‘Member’ to ‘Membership’ and ‘Employment’ to ‘Work History’?

[0281] In each case except ‘Employment’, the classes each have identical Data Properties so we can assume that they are nearly identical. Mathematically the Sameness( $C_{1i}, C_{2i}$ )  $\sim 1.0$  where  $C_{1i}$  is a concept from the first ontology and  $C_{2i}$  is a concept from the second ontology.

[0282] Some Concepts have different Data Properties. In the case of ‘Employment’ and ‘Work History’ they have some identical Data Properties and one, ‘Reports To’, which applies only to ‘Work History’. In fact ‘Work History’ violates 4<sup>th</sup> Normal Form as it is ambiguous as to whether the ‘Start Date’ and ‘End Date’ refer to the ‘Role’ or the ‘Reports To’ Data Property.

[0283] Some Concepts have different Object Properties. ‘Work History’ has two Object Properties with ‘Person’, whereas ‘Employment’ has only one. In Ontology 1 ‘Shares’ relates ‘Company’ to ‘Individual’ whereas in Ontology 2 it relates ‘Company’ to ‘Client’.

[0284] Some Concepts do not exist in one Ontology. ‘Listed Company’ exists in Ontology 2 but not in Ontology 1.

[0285] For the purpose of these examples, the system performs the functionality shown in FIG. 13, with these being implemented by respective modules. In this regard the modules include:

[0286] ETL (Extraction-Transformation-Loading) module **1300**. This extracts, transforms and loads content within structured data sources. This includes two sub-components, including:

[0287] Processor **1301** that extracts source data either via a specified ontology, or, in the absence of an ontology, via a putative ontology which the Processor creates to describe the data. The Processor can be deployed either in the Cloud or on the same machine as the data or on a machine which can access the data via messaging, ODBC, https, SOAP or any equivalent protocol. Multiple copies of the Processor can be deployed in order to obtain data from multiple sources.

[0288] Orchestrator **1302** that collects data from the various Processors and maps the source ontologies to the target ontology. Queries are written using the target ontology and are translated into equivalent source ontology queries, allowing data to be returned using the target ontology.

[0289] Ontology Browser module **1310** including a browser **1311**, editor **1312** and generator **1313**. This generates screens and the associated software and data to manage them, which enables a user to browse and edit an ontology and the data described by the ontology. These screens appear in two stages. The first stage is during the generation process. In this stage the screens are dynamically created and display additional information to enable the user to select which features are to be generated. In the second stage the screens are hard coded and only display the information specified for generation.

[0290] Ontology Indexer module **1320**. The Indexer module creates a set of linked indexes on one or more ontologies, of all the class names, data property names and, object property names. Additionally the index includes semantically equivalent terms (synonyms and homonyms for example) which come from the source ontologies plus from a semantic equivalence function.

[0291] Ontology Pruner module **1330**. The Pruner module takes an ontology and allows a user to specify which classes, data properties, object properties and axioms they wish to retain. Using those retained the Pruner module checks to see that the relational and axiomatic integrity defined in the ontology is maintained.

[0292] Ontology Aligner module **1340**. The Aligner module takes two or more ontologies and uses a number of techniques to align the concepts in the various ontologies, either with each other or with a specified target ontology. The techniques utilise the indexes created by the indexer module to find concepts which are semantically similar. Each data property and concept is compared using the semantic matcher module. It refines the matching based upon the ontology structure and the data properties.

[0293] Semantic Matcher module **1350**. The matcher module compares two terms or two lists of terms to determine whether they have a mathematically defined degree of semantic equivalence within a specified context, for example medicine or engineering, or, in another instance, given a single term, will provide a list of synonyms, homonyms, etcetera based upon a specified context.

[0294] Typically an ontology does not have any data instances except as examples, however an ontology can be matched to existing data in one of two ways:

[0295] The ontology is constructed from the existing data. For example a relational database could be automatically converted to a ‘putative’ ontology by relational Entities (tables) being defined as ontological Classes, relational Relationships as ontological Object Properties, and relational Attributes (columns) as ontological Data Properties. Some ontological axioms could be derived from relational referential integrity constraints, but most axioms would need to be manually added or ignored. This putative ontology may then be aligned with an existing rich ontology to add metadata.

[0296] Matching the ontology to the data. There are a number of tools for doing this (e.g. S-Match).

[0297] Regardless of the data format a putative ontology can be automatically generated from the source data using methods appropriate to the source data structure and metadata (if it exists). This putative ontology may be manually updated using the ontology editor, or used as generated. In either case the putative ontology is then aligned using the aligner module with a subject area ontology (invoked by the ETL module processor) and with the target ontology (invoked by the ETL module orchestrator).

[0298] The target ontology may be pruned using the pruner module, to ensure that it contains only the desired concepts plus those concepts, axioms, properties, inferences and provenance details which are required to ensure the integrity of the desired concepts.

[0299] All these tools make use of the services provided by the semantic matcher module to check if two semantic concepts match, and the indexer module to look for matching concepts and conceptual structures in the various source and target ontologies.

[0300] Examples of the respective modules will now be described in further detail.

#### ETL Module

[0301] The ELT module performs the functions of data extraction, transformation and loading common to all ETL tools without the use of a metadata repository. It does this by using metadata associated with the source data to determine data structure and then by mapping this metadata to an ontology. It also assigns meaning to the data and hence is able to achieve a high level of automation in mapping and transforming the data.

[0302] Eliminating the need for a metadata repository means that the flexibility of the processes is not constrained by the human interface required to maintain it. New data formats and technologies can be automatically accommodated.

[0303] At a high level there are two major processes performed. The code to perform these processes is called the processor and the orchestrator. Numerous copies of the processor may be deployed to read data at any defined location. The processor can be co-located on the same device as the data or it can be located in the cloud and access the data using a remote access protocol. The processor extracts metadata from the source and creates a putative ontology from that metadata. It then performs some elementary data transformations and passes the data and the ontology to the orchestrator.

**[0304]** The orchestrator receives input from the various processors and aligns their ontologies. It then applies a mapping from the aligned source ontologies to the user defined target ontology. The user can now see all the data from the various source ontologies. Data can be extracted either by specifying a specific query against the target ontology or by using the ontology browser module to create the query, as will be described in more detail below.

**[0305]** An example ETL module software stack including the various software components which are required to achieve this outcome are shown in FIG. 14A, whilst FIG. 14B shows an example deployment in which a number of processors are coupled to a single orchestrator via a network arrangement.

**[0306]** The processor is responsible for reading data from disparate data source, exposing the data as RDF and creating a putative ontology to describe the data. The high level functions are as follows:

**[0307]** Register disparate data sources by adding meta-data and mapping files.

**[0308]** Convert unstructured data into RDF.

**[0309]** Load RDF into triple-store.

**[0310]** Convert the mapping files into putative ontologies.

**[0311]** Expose a SPAQRL endpoint for each source.

**[0312]** The orchestrator is responsible for reading target ontologies and mapping files and orchestrating the transformation of request and response. The high level functions are as follows:

**[0313]** Register target ontologies.

**[0314]** Read mapping files and index them.

**[0315]** Transform SPARQL queries from target to mapped source vocabularies.

**[0316]** Transform Response from source to target vocabularies.

**[0317]** Store transformation rules.

**[0318]** Expose a SPARQL endpoint for the target.

### Ontology Browser Module

**[0319]** The ontology browser module operates to automatically create a set of screens to enable a user to browse an ontology, query data defined by an ontology and add instance data to data defined by an ontology. The screens thus generated can then be used independently of the ontology and the creating tools, as a complete stand-alone application.

**[0320]** In this regard, currently the use of ontologies to define linked concepts and to access data is largely confined to academics and professional ontologists. The reason for this is that there is no simple mechanism for allowing users to browse ontologies, and then use this in guiding their interaction with data stored in structured data stores. Accordingly, by providing a tool that enables a person with little or no ontological expertise to access all the details of an ontology in a simple comprehensible fashion, this allows the user to select and inspect the data described by the ontology using a simplified query construction mechanism. They will be able to add records to the data with all the constraints and inferences which exist in the original ontology still being enforced. Finally they will be able to deploy the generated screens as stand-alone applications suitable for use by front-office personnel.

**[0321]** When inspecting the data, the user can display it in a number of formats. The underlying data can be stored as

RDF Triples, for example. These can be displayed as relational tables, spreadsheets, name-value pairs or any user defined format.

**[0322]** The ontology browser module can exist in two major forms, either as a stand-alone tool, or second as a plug-in to existing ontology tools (such as Protégé). In either form it can generate an application specific to the ontology selected.

**[0323]** The generated application can be used without the ontology as a full function code-set for accessing, updating, deleting and adding records with all the data rules defined in the original ontology being enforced.

**[0324]** Thus, the ontology browser module provides a set of processes which can be implemented in a computer program which generates screens and the associated software and data to manage them which enables a user to browse and edit an ontology and the data described by the ontology. These screens appear in two stages. The first stage is during the generation process. In this stage the screens are dynamically created and display additional information to enable the user to select which features are to be generated. In the second stage the screens are hard coded and only display the information specified for generation.

**[0325]** A brief description of the screens is set out in Table 2 below.

TABLE 2

#	Screen Title	User Story	Notes
1	Landing Page	The user will first access the 'landing page' which will have the capability of listing the available ontologies. The user will select an ontology. Having selected the ontology the tool will generate the screens to manage that ontology.	This is the entry point for this application. It could be done in an object oriented fashion. This screen is not generated in the deployable code.
2	Class List	Having selected an ontology the user will be presented with a list of classes in the chosen ontology. The User will select one class as the entry point to the ontology.	Each list item could include a label to help in identification. This screen is generated in the deployable code. It would be the entry point for deployable code.
3	Class Display	All the data property fields for the class are displayed in the main frame, along with four additional frames which are: The parent/super classes of the selected class-a clickable link utilising super class relationships. The child/sub classes of the selected class-a clickable link utilising subclass relationships. The related classes of the selected class-a clickable link using Object Property details. The axioms which impact that class, This frame is only displayed during the generation process. In the generated screens axioms are not displayed except as an error message if invalid data is entered. The field names on the screen are displayed with an adjacent data entry field which is blank when browsing an ontology.	A screen or a set of related screens is generated in the deployable code, as a screen or set of screens specific to each class. Class specific screens can be generated using a number of templating tools to ensure that a particular look and feel is generated.

TABLE 2-continued

#	Screen #	Title	User Story	Notes
4	Query		<p>Editing mechanisms are provided to select the classes and properties for the screens which are to be generated. Fields can be marked 'non-searchable' to control resource usage.</p> <p>A query is performed by adding data into a data property field in one or more class screens. Additional constraints can be defined by the user. Once the query has been defined the user selects a 'Search' option and the records meeting the search criteria are returned.</p>	

[0326] These screens are available without generation in a generic format such that a single screen is used for each type of screen. The screen layout is dynamically determined by the ontology content.

[0327] Generic screens are not user friendly and cannot be customised. Therefore the process allows the user to generate a complete set of screens whose look and feel can be parametrically predetermined using facilities such as cascading style sheets, Templates, icons and user supplied parameters.

[0328] An example of the arrangement of the browser module is shown in FIG. 15.

[0329] In this regard, the browser module 1310 takes a target ontology 1501 from the orchestrator 1302, or any ontology defined by the user. The Browser module 1310 displays the set of screens 1502 which allowing the user to browse the ontology and to specify which components of the ontology to generate into a standalone application.

[0330] The browser module 1302 generates a standalone application 1503 including a set of computer screens 1504 to manage the data using the structure and rules specified in the target ontology. The application can be generated in a number of modes, such as purely an ontology or data browser module, or as a full function data add, update and delete application. In this case the user now has a complete application 1503 to manage the data described by the ontology.

[0331] Ontologies using OWL or RDF files have enough information to generate web pages and create a corresponding database 1505 to store the information. The RDF or OWL file may have been created by an ontologist based upon their detailed business knowledge.

[0332] Thus the browser module 1310 creates an application 1503 for end users to query or enter transaction data. The OWL or RDFS file is fed into the browser module 1310 along with application customisation files, database connection details and any other metadata required to create the application.

[0333] The browser module 1310 can create web pages, for example using HTML5, JSP, JSF or any similar technology. For each class in the ontology browser module 1310 creates a web page and each property associated with that class is created as a field within the page. The application 1503 bridges between the generated webpages and the database 1505. It performs the processes to persist the data from the web pages to the database 1505, to extract data from the database 1505, to query data in the database 1505 and to display data on the web page. The browser module

1310 then creates database scripts for creating and loading a database of the type specified in the user supplied metadata. This could be a relational database (RDBMS), a Triple Store, NOSQL, NewSQL, Graph Database or any other recognised database.

[0334] Operation of the browser module will now be described in more detail. In this regard, in order to browse an ontology a user must be able to find ontology terms:

[0335] concepts;

[0336] data properties;

[0337] object properties; and

[0338] inferences.

[0339] This requires two mechanisms, namely:

[0340] a method for indexing the above ontology terms from an ontology, in order to search for any such ontology term by name, as described with respect to the indexer module below; and

[0341] a method for displaying all the related data and object properties once a particular property has been chosen.

[0342] To achieve this, the user initially selects the ontology to be browsed in the 'Landing screen' described in Table 2. The ontology can be selected from a file or a Web address. Once the ontology has been selected a class list is generated using an index of the ontology. This list displays the name and description of each class. For larger lists a list search function is provided enabling the user to search by class name or part of a class description. It is also possible to search on a data property. In either case the search would return a list of classes which contain that data property.

[0343] The user then selects the class of interest, which causes a 'Class screen' to be displayed including four components, in the form of frames or tagged sub-screens, as follows:

[0344] The Data Property Component. The name of each data property is displayed in a list format with a description box beside the field. Clicking on an information icon beside the field will display all the field attributes and any axioms related to that field. Optionally (clickable), data properties of a parent/super or related class or classes may also be shown

[0345] The parent/super Class Component. This displays the name and description of the parent/super class of the displayed class, with a clickable link to it. Clicking on this link will cause the browser module to display a screen displaying the Parent of the current class.

[0346] The child/sub Class Component. This displays the name and description of the subclasses of the displayed class, with a clickable link utilising subclass relationships. Clicking on one of these links will cause the browser module to display a Child/sub class or subclass of the current class.

[0347] The Object Property Component. This displays the related Classes of the selected class, each with a clickable link using the object property. Clicking on one of these links will cause the browser module to display a class related to the current class.

[0348] By selecting a 'Search' option on a class screen a query is issued to return all the data instances for that class. This is displayed as a list with one row for each instance of the class. By clicking on a particular row, that row is displayed as a formatted screen similar to the ontology class screen. In one example, the data returned maybe restricted

by executing a query which would filter the results. The construction and use of such a query will now be described in more detail.

[0349] In this regard, filtering the data returned to the user is achieved by capturing from the user, the user's exact requirements of the data to be returned, in the form of a filter and then generating a query based on that filter. The filter is constructed by entering values or expressions into the data property fields on a class screen. For example, using the sample ontologies described above, to find out how many shares John Doe owns, the following steps would be required:

[0350] Select the 'Individual' class from the class list screen.

[0351] In the Data Property fields enter 'John' into Given-Name and Doe' into Last-Name.

[0352] From the Object Property frame of the 'Individual' class screen, select the 'Shares' class.

[0353] Select the Search option.

[0354] By selecting the 'Search' option on a Shares Class screen a query is issued to return all the data properties for that class but only those owned by John Doe. The filter has been transformed by the generated application 1503 into a SPARQL or functionally equivalent query which can be executed against the data stored in the database 1505.

[0355] To allow the browser module 1310 to generate the application 1503, the following process is performed:

[0356] Optionally configure metadata for the application to be generated including items such as:

[0357] Company name, logo etc.

[0358] Name of the application to be generated.

[0359] Name and type of database to be created.

[0360] Location of the database.

[0361] Naming and coding specification and standards for the application to be generated. This includes style sheets, Templates, Java scripts and other display specifications.

[0362] Icons to be associated with classes and actions.

[0363] Location and contact details of help desk.

[0364] Verbosity of error and log messages.

[0365] On the 'Landing Screen' select the ontology from which to generate, resulting in a 'Class List' screen being displayed by the browser module 1310.

[0366] On the Class List screen tag each class to be generated with 'g'.

[0367] Select each class to be generated, causing the browser module 1310 to display the 'Class Display' screen.

[0368] On the Class Display screen all fields are initially tagged with a 'g'. Review each data property field, each super/subclass link and each object property link to be generated, removing the tag if it is not required.

[0369] By default all fields are searchable (i.e. can be added to a filter). Adding an 'ns' tag to a data property field will mean that that field will be non-searchable in the generated application.

[0370] There are additional field tag positions on each of the super/subclass link fields and the object property link fields. By setting an 'I' tag in these fields it will generate data fields from the linked class into the generated screen. These fields will be displayed as non-updateable fields.

[0371] If any fields from linked classes are to be displayed, select the linked class and tag the appropriate fields with an 'I'.

[0372] Return to the Class Display screen and remove the tag from each axiom description if it is not to be enforced. It is important to remove fields before axioms as otherwise there may be a loss of integrity in the generated application.

[0373] Repeat Steps 3-9 until all the required classes have been selected for generation.

[0374] Return to the Class List screen and selects a 'Generate Application' option.

[0375] The application will be generated by the browser module 1310 and saved into the location specified in the application metadata (Step 1). The database creation and load scripts will be created. Run these scripts to ready the application for use.

[0376] Accordingly, the above described browser module 1310 allows a user to browse and interact with ontologies, and then by selecting specific classes and data properties, generate an application 1503 that can be used to interact with data stored in a data store 1505 in accordance with the selected classes and data properties.

#### Ontology Indexer Module

[0377] The indexer module automatically creates a set of indexes of the terms used in a collection of one or more ontologies to assist a user to browse an ontology and to expedite the querying of data defined by an ontology. These indexes are used by the other modules to assist in the alignment, pruning and browsing of ontologies.

[0378] The indexer module indexes one or more ontologies by creating a set of linked indexes of all the class names, data property names and object property names and relationships. The index includes semantically equivalent terms which come from the source ontologies plus from a semantic equivalence function.

[0379] An example of the functionality of the indexer will now be described with reference to FIG. 16.

[0380] In this example, the indexer module 1320 receives an ontology 1601 from the orchestrator 1302, or any ontologies defined by the user, via a set of screens 1602, or by the processor 1301 and creates indexes 1603 of all the class names, data property names and, object property names. It will be appreciated that the screens may be generated by the browser module 1310 as previously described.

[0381] As each ontology term is indexed, synonyms of that item, obtained from the semantic matcher module 1350, using a concept matching database 1604, are also indexed. For Object Properties, the concepts linked by the object property are cross referenced in an index.

[0382] A sample of the Concept-Data Property-Object Property (CDO) index based on the example ontologies above is shown in Table 3. It should be noted that this is a display form of the index for the purpose of illustration but that in practice the index may be stored in a more complex index structure as will be described in more detail below.

TABLE 3

CDO	Type	Address
Client	Concept	Ont 2.0
Club	Concept	Ont 1.3

TABLE 3-continued

CDO	Type	Address
Club	Concept	Ont 2.3
Company	Concept	Ont 1.4
Company	Concept	Ont 2.4
Employment	Concept	Ont 1.6
Individual	Concept	Ont 1.1
Listed Company	Concept	Ont 2.5
Member	Concept	Ont 1.5
Membership	Concept	Ont 2.8
Organisation	Concept	Ont 1.2
Organisation	Concept	Ont 2.2
Party	Concept	Ont 1.0
Person	Concept	Ont 2.1
Qango	Concept	Ont 2.7
Shares	Concept	Ont 1.7
Shares	Concept	Ont 2.10
Unlisted Company	Concept	Ont 2.6
Work History	Concept	Ont 2.9
Date of Birth	Data Property	Ont 1.1
Date of Birth	Data Property	Ont 2.1
Date of Incorporation or founding	Data Property	Ont 1.2
Date of Incorporation or founding	Data Property	Ont 2.2
End Date	Data Property	Ont 1.6
End Date	Data Property	Ont 2.9
Exit Date	Data Property	Ont 1.6
Exit Date	Data Property	Ont 2.9
Family Name	Data Property	Ont 2.1
First Name	Data Property	Ont 1.1
Gender	Data Property	Ont 1.1
Gender	Data Property	Ont 2.1
Given Name	Data Property	Ont 2.1
Joined Date	Data Property	Ont 1.5
Joined Date	Data Property	Ont 2.8
Last Name	Data Property	Ont 1.1
Ministry	Data Property	Ont 2.7
Name	Data Property	Ont 1.0
Name	Data Property	Ont 2.0
Number	Data Property	Ont 1.7
Number	Data Property	Ont 2.10
Purchase Date	Data Property	Ont 1.7
Purchase Date	Data Property	Ont 2.10
Registered in	Data Property	Ont 2.6
Registered in/on	Data Property	Ont 1.4
Reports to	Data Property	Ont 2.9
Role	Data Property	Ont 1.6
Role	Data Property	Ont 2.9
Start Date	Data Property	Ont 1.6
Start Date	Data Property	Ont 2.9
Stock Exchange	Data Property	Ont 2.5
Type	Data Property	Ont 1.3
Type	Data Property	Ont 1.4
Type	Data Property	Ont 1.5
Type	Data Property	Ont 1.7
Type	Data Property	Ont 2.10
Type	Data Property	Ont 2.3
Type	Data Property	Ont 2.4
Type	Data Property	Ont 2.8
Employs	Inv Obj Prop	Ont 1.6
Employs	Inv Obj Prop	Ont 2.9
Has	Inv Obj Prop	Ont 1.5
Has	Inv Obj Prop	Ont 2.8
Holds	Inv Obj Prop	Ont 1.5
Holds	Inv Obj Prop	Ont 2.8
is a	Inv Obj Prop	Ont 1.0
is a	Inv Obj Prop	Ont 1.0
is a	Inv Obj Prop	Ont 1.2
is a	Inv Obj Prop	Ont 1.2
is a	Inv Obj Prop	Ont 2.0
is a	Inv Obj Prop	Ont 2.0
is a	Inv Obj Prop	Ont 2.1
is a	Inv Obj Prop	Ont 2.2
is a	Inv Obj Prop	Ont 2.2
is a	Inv Obj Prop	Ont 2.2
is a	Inv Obj Prop	Ont 2.4

TABLE 3-continued

CDO	Type	Address
is a	Inv Obj Prop	Ont 2.4
Owns	Inv Obj Prop	Ont 1.7
Owns	Inv Obj Prop	Ont 2.10
Reports to	Inv Obj Prop	Ont 2.1
Shareholder	Inv Obj Prop	Ont 1.7
Shareholder	Inv Obj Prop	Ont 2.10
Works at	Inv Obj Prop	Ont 1.6
Works at	Inv Obj Prop	Ont 2.9
Employs	Obj Property	Ont 1.4
Employs	Obj Property	Ont 2.4
Has	Obj Property	Ont 1.3
Has	Obj Property	Ont 2.3
Holds	Obj Property	Ont 1.1
Holds	Obj Property	Ont 2.1
is a	Obj Property	Ont 1.1
is a	Obj Property	Ont 1.2
is a	Obj Property	Ont 1.3
is a	Obj Property	Ont 1.4
is a	Obj Property	Ont 2.1
is a	Obj Property	Ont 2.2
is a	Obj Property	Ont 2.3
is a	Obj Property	Ont 2.4
is a	Obj Property	Ont 2.5
is a	Obj Property	Ont 2.6
is a	Obj Property	Ont 2.7
Owns	Obj Property	Ont 1.1
Owns	Obj Property	Ont 2.0
Reports to	Obj Property	Ont 2.9
Shareholder	Obj Property	Ont 1.4
Shareholder	Obj Property	Ont 2.4
Works at	Obj Property	Ont 1.1
Works at	Obj Property	Ont 2.1

[0383] Even without the inclusion of synonyms this is an extremely useful index. For example, every concept which has the same name in two different ontologies can potentially be aligned. The Aligner module would take each such pair and compare first their Object Properties and then their Data Properties.

[0384] For example, the concept ‘Shares’ appears in both ontologies as concepts Ont 1.7 and Ont 2.10. At this level they appear to be similar ( $S_{1,7,2,10}=1.0$  because the names are identical) and from an indexer module point of view that is sufficient.

[0385] Further analysis could be performed by the aligner module described in more detail below. By examining the Object Properties it would find that the Object Properties are different as shown in Table 4 below. Although they match in number and Object Property name, one of the related concepts is different giving  $S_{1,7,2,10}=0.8571$ . By examining the Data Properties we find that they have identical Data Properties giving  $S_{1,7,2,10}=1.0$ .

[0386] The source information on which the aligner module performed the preceding calculations is all available in the indexes created by the indexer.

TABLE 4

	Ontology 1		Ontology 2	
Individual	Owns	Shares	Client	Owns
Company	Shareholder	Shares	Company	Shareholder
				Shares

[0387] Further analysis of the other concepts using the semantic matcher module would show that an “Individual” is a subclass of “Client” hence giving  $S_{1,7,2,10}=0.8>0.95$ .

Ontology 2 is a more generic model than Ontology 1. This similarity range is adequate to establish anchor points between Shares in the two ontologies. The calculations of  $S_{ij}$  are performed by the aligner module.

[0388] The relationship between concepts is extracted in the Concept to Concept (C2C) table shown in display form in Table 5, which shows how Concept C1 relates to Concept C2.

TABLE 5

Object Rel	C1	C2
Employs	Ont 1.4	Ont 1.6
Employs	Ont 2.4	Ont 2.9
has	Ont 1.3	Ont 1.5
has	Ont 2.3	Ont 2.8
Holds	Ont 1.1	Ont 1.5
Holds	Ont 2.1	Ont 2.8
is a	Ont 1.1	Ont 1.0
is a	Ont 1.2	Ont 1.0
is a	Ont 1.3	Ont 1.2
is a	Ont 1.4	Ont 1.2
is a	Ont 2.1	Ont 2.0
is a	Ont 2.2	Ont 2.0
is a	Ont 2.3	Ont 2.2
is a	Ont 2.4	Ont 2.2
is a	Ont 2.5	Ont 2.4
is a	Ont 2.6	Ont 2.4
is a	Ont 2.7	Ont 2.1
is a	Ont 2.7	Ont 2.2
Owns	Ont 1.1	Ont 1.7
Owns	Ont 2.0	Ont 2.10
Reports to	Ont 2.9	Ont 2.1
shareholder	Ont 1.4	Ont 1.7
shareholder	Ont 2.4	Ont 2.10
Works at	Ont 1.1	Ont 1.6
Works at	Ont 2.1	Ont 2.9

[0389] The indexes are constructed in multiple formats, corresponding to sorting the above tables into different sequences. The aligner module can perform many of its tasks by executing SQL queries against the indexes.

[0390] An example of the index structure will now be described in more detail. In this regard, using the semantic matcher module, a root word or lemma is determined for each synonym set. The semantic matcher module requires that the context be set in order to obtain the optimum results. In general, when constructing indexes over a number of ontologies the context of each ontology is known, narrow and related to the other ontologies of interest.

[0391] The final set of indexes is created in a multi-step process summarised below:

[0392] Extract all concepts, Object Properties and Data Properties from the ontology being indexed.

[0393] Load these values into temporary tables (CDO and C2C) with the format described in Tables 3 and 5. These tables are created or recreated empty for each ontology being indexed.

[0394] The ontology is loaded into the semantic matcher module. This will examine every word semantically using any definitions contained in the ontology and comparing them with those definitions already loaded into the semantic matcher module or available from public dictionaries such as WordNet. The context is supplied by the ontology (e.g. Medical/Surgical or Geographical Location)

[0395] The semantic matcher module defines a Concept Id, a unique number corresponding to the lemma or root word for every family of synonyms.

[0396] The synonym table is then loaded with terms matching the terms in the temporary tables described above with the Concept Id.

[0397] All synonyms identified by the semantic matcher module for each term in the ontology being indexed are also loaded into the Synonym table.

[0398] The final CDO index is then created by substituting the appropriate Concept Id for each term in the CDO table.

[0399] The final C2C index is then created by substituting the appropriate Concept Id for each term in the C2C table.

[0400] The temporary (display versions) of the index are deleted.

[0401] The next ontology to be indexed is then loaded by repeating all the preceding steps.

[0402] When all the relevant ontologies have been indexed, a final pass of the synonym table against the semantic matcher module is performed in case any new synonyms have been identified during the loading process.

[0403] The indexes are loaded into an appropriate database structure and tuned for performance. Typically this will involve creating multiple database indexes over the ontology index tables.

[0404] It will be appreciated that there is no direct user interaction with his tool or with the indexes. Instead the indexer module provides a service which is used by other modules, tools or components.

[0405] Some of the services which this index can provide include the enhanced ability to:

[0406] choose the best ontology from a selection of ontologies;

[0407] align or merge multiple ontologies;

[0408] navigate an ontology;

[0409] extract synonyms;

[0410] perform semantic matching.

#### Ontology Pruner Module

[0411] The pruner module is designed to enable a user to take a large ontology or a collection of aligned ontologies and prune them down to the classes of interest for the user's needs, without losing integrity by inadvertently deleting a component which contains data or axioms relevant to their ontology terms of interest.

[0412] For example, issues arise when constructing and utilising a large reference ontology, such as the Foundational Model of Anatomy (FMA). In this regard, the FMA is very large and highly detailed, though also very general in nature (e.g. non-application specific). It is also rigorous in its adherence to proper modelling principles. These criteria together lend the FMA to many possible applications. However, they have also rendered it cumbersome (i.e. overly large or detailed or principled) for use by any specific application.

[0413] As a result, potential users of the FMA had requests of the following basic form, "we really like the FMA, but it is too large or too detailed for our needs, we really only need something based on subsets of the whole FMA". The basis for division varied, application to application, but examples include:

- [0414] Region-based, i.e. the brain or the abdomen.
- [0415] System-based, i.e. the cardiovascular system or the skeletal system.
- [0416] Granularity-based, i.e. only items visible in an x-ray or only cellular and sub-cellular components.
- [0417] Though the desired ontology derivative was generally based on a subset extraction such as those above, it was then often further manipulated to better suit the needs of the application (i.e. classes added, classes removed, properties removed, properties added, etc.).
- [0418] Such requests could be handled in one of three ways:
- [0419] Writing procedural code specific to each new request, which is not a generic solution.
  - [0420] Creating views over the ontology, which needs a language for defining the desired application knowledge base (KB) (not always a proper ontology) as well as an engine that could generate the application KB
- [0424] As another example a user may wish to create a new ontology from components of several existing source ontologies and then add their own additions. The combined ontology would contain many irrelevant concepts which would need to be removed. For example, a parcel delivery company combining a transport ontology with a geo-location ontology to create an ontology which enables delivery routes to be determined and optimised. By combining these ontologies and adding axioms such as aeroplanes start and stop their journeys at airports, ships at ports and trains at stations, it would be possible to construct an information base covering every concept in their business model. However much of each source ontology would not be needed.
- [0425] The pruned ontology definition may be used in place of a view over the complete ontology. This view could be used for a number of purposes such as access control, scope management etc.
- [0426] To achieve this, the pruner module operates in conjunction with the browser module to perform the functions set out in Table 6 below.

TABLE 6

#	Title	Function
1	Prune Single Ontology	Use the updated ontology browser module to enable the tagging of classes by the user in a manner which allows the creation of a coherent, integrated subset of the ontology with all relevant object properties, axioms and inferences from the source ontology.
2	Prune Overlapping Ontologies	Using the semantic matching tool and the updated ontology browser module to enable the tagging of classes in a manner which allows the creation of a coherent, integrated subset of the ontology with all relevant object properties, axioms and inferences from the source ontologies. Include a mechanism to determine if ontologies are in fact disparate.
3	Prune Disparate Ontologies	Include a mechanism to add the necessary detail to join the disparate ontologies. Then iteratively apply the mechanism above to establish the pruned ontology. For example-given a geo ontology and a transport ontology construct a journey ontology which would allow analysis of the appropriate transport mechanisms between two locations.

- from the definition and the source ontology(ies). This has problems with adding and removing properties.
- [0421] Pruning the ontology to deliver a well modelled subset ontology.
- [0422] Thus, there are many needs for a pruned ontology, such as relevance, performance, manageability and testability and these requirements should be met by a tool which enables a person with little or no ontological expertise to safely prune unneeded concepts. Furthermore that person should be able to select and inspect the data described by the ontology by using a simplified query construction mechanism. They then will be able to study the effects of removing components from the ontology before committing to their removal, and then save the pruned ontology as a new ontology.
- [0423] For example, SNOMED-CT is a large medical ontology of medical terms used in clinical documentation. It consists of 300,000+ concepts with about 1,400,000 relationships between them. The concepts are divided into 19 functional areas. A researcher may only be interested in one of these areas, say mental health. Removing the other 18 areas would break many of the relationships between medical health terms and pharmaceutical terms. Obviously they may wish to retain these items. To do so manually would require many months of work with existing tools and would be prone to error.
- [0427] The pruner module interacts with the browser module to allow a user to specify which classes, data properties, object properties and axioms of a selected ontology they wish to retain. Using those retained the pruner module checks to see that the relational and axiomatic integrity defined in the ontology is maintained.
- [0428] In another version the user may specify two essential concepts within a single ontology which must be retained in the pruned ontology. The invention then maps all the conceptual relationships between classes, tagging all classes which are required to analyse the specified concept. Additional classes, object properties and axioms are then included from the source ontology to ensure the integrity of the pruned ontology.
- [0429] In another version the user may specify two essential concepts from disparate ontologies which must be retained in the pruned ontology. The pruner module then attempts to map all the conceptual relationships between classes, tagging all classes which are required to analyse the specified concept. If no connecting paths are identified the software will recognise the potential impossibility of creating a pruned ontology which connects the two starting concepts. The user will be asked to:

- [0430] Abandon the attempt, or  
 [0431] Redefine their goals and start again, or  
 [0432] Enlarge the scope by adding additional classes either manually or from another ontology and start again.
- [0433] Assuming success the user now has a complete ontology which is greatly reduced in size from the combined source ontologies.
- [0434] An example of the arrangement of the pruner module is shown in FIG. 17A.
- [0435] In this example, the pruner module 1330 opens ontologies 1701 defined in OWL and RDFS files, with the user then interacting with the pruner module 1330 via a set of screens 1702 as defined in Table 7 below, to thereby produce a pruned ontology 1703. It will be appreciated that the screens may be generated by the browser module 1310 as previously described.
- | # | Screen Title  | User Story   |
|---|---------------|--|
| 1 | Landing page  | The user will first access the ‘landing page’ which will have the capability of listing the available ontologies. The user will select an ontology. Having selected the ontology the tool will generate the screens to manage that ontology.   |
| 2 | Class List    | Having selected an ontology the user will be presented with a list of classes in the chosen ontology. The User will select one class as the entry point to the ontology  |
| 3 | Class Display | All the data property fields for the class are displayed in the main frame, along with four additional frames which are:<br>The parent/super classes of the selected class-a clickable link utilising super class relationships.<br>The child/sub classes of the selected class-a clickable link utilising subclass relationships.<br>The related classes of the selected class-a clickable link using Object Property details.<br>The axioms which impact that class.<br>The field names on the screen are displayed with an adjacent data entry field which is blank when browsing an ontology.<br>Editing mechanisms are provided to select the classes and properties for the screens which are to be retained in the pruned ontology. |
- [0436] When pruning a single ontology this is a tool assisted manual process, as will now be described with reference to FIG. 17B.
- [0437] In this example, the user selects the concepts that they require and the tool identifies and adds the components required for completeness and integrity. The user selects a class as a starting seed point  $S_0$  in the source ontology and tags it as  $K_0$  for keep.
- [0438] The computer identifies and tags as ' $K_1$ ' all parents of classes marked ' $K_0$ ', all classes and inferences from classes and inferences tagged as  $K_0$ . These tagged variables are called the  $S_1$ -shell. The user reviews the computer tagged items and retags them as  $K_1$  for Keep,  $M_1$  for Maybe and  $D_1$  for Discard. All axioms are loaded for the tagged  $M_i$  and  $K_i$  components. The process is then repeated, incrementing  $i$  each time until the user has tagged all the components for the appropriate ontology.
- [0439] A reasoner is then applied to the resulting ontology to identify potential errors and add inferred values. Any concepts, inferences or axioms thus added are tagged  $K_n$  and the tagged components are exported as the pruned ontology.
- [0440] For multiple overlapping ontologies, the process is as shown in FIG. 17C.
- [0441] In this example, the user selects a class as a starting seed point  $S_0$  in one ontology and another as ending seed point  $E_0$  in either the same or another ontology and tags them both as  $K$  for Keep with ' $K_{0s}$ ' or ' $K_{0e}$ '.
- [0442] The computer identifies and tags as ' $K_{1s}$ ' or ' $K_{1e}$ ' all parents of classes marked ' $K_{0x}$ ', and all subclasses and inferences from classes and inferences tagged as ' $K_{nx}$ ' where  $n=1$ . These tagged variables are called the  $S_1$ -shell and the  $E_1$ -shell. The variables in the  $S$  and  $E$  shells are compared by the semantic matcher module described in more detail below. The matcher module returns a numeric value for the match quality between variables in each shell. If the predetermined match quality is met then a path has been determined between the two shells. This should only occur if the shells overlap. If the start and end point are in the same ontology the match quality must be 1.0 or exact.
- [0443] At any stage, the data properties of a tagged data class may be pruned. This is performed by selecting the class and marking the data fields (data properties) as ‘ID’ for Discard. Any inferences based upon the existence of the discarded field will be ignored.
- [0444] These steps are iterated, incrementing  $n$  by 1 each time until a predetermined number of variables have appropriate match quality or a predefined depth of shell is reached. The shell paths of the matching variables are tagged ' $P_{jx}$ '. If the predefined depth of shell is reached without establishing any paths then the process has failed and the ontologies are considered disparate. The process stops. At this point it is possible to increase the predefined shell depth, and to manually change the tag of any concepts which are considered out of scope from K to D for Discard. The process can be restarted.
- [0445] Once these have been established, the paths  $P_j$  between  $S_0$  and  $E_0$  can be populated and a skeletal pruned ontology can be defined in terms of these paths. All class parents and inferred parents for tagged  $P_j$  path components are also tagged as belonging to the path  $P$ . All axioms are loaded for the tagged  $P_j$  path components thus creating an expanded ontology.
- [0446] A reasoner is applied to the expanded ontology to identify potential errors and add inferred values. Any con-

cepts, inferences or axioms thus added are tagged and exported as part of the pruned ontology.

[0447] For disparate ontologies the process is as shown in FIG. 17D. In this regard, disparate ontologies can arise for two possible reasons:

[0448] The user did not realise that they were disparate until they attempted to align them or to extract a subset ontology from concepts in the two ontologies. This is a potential failed outcome of the previous section; or

[0449] The user knows that they are disparate and is supplying concepts and properties to enable them to join.

[0450] In either case, the user must supply the information to enable the ontologies to be joined. This is effectively the starting point for the process.

[0451] The user selects a class as a starting seed point  $S_0$  in one ontology and another as ending seed point  $E_0$  in the other ontology and tags them both as K for Keep with ' $K_{os}$ ' or ' $K_{oe}$ '. In addition they define a set of user defined paths which connect the ontologies, as shown by the lines 1710.

[0452] These paths have start and end points ' $U_{osi}$ ' and ' $U_{oei}$ ' where T is the path number being defined. These paths form a contiguous set of related concepts, starting with a class in one ontology and ending with a class in another ontology.

[0453] The process described above for overlapping ontologies then applied to each concept pair  $S_0$  and ' $U_{osi}$ ' and  $E_0$  and ' $U_{oei}$ ' to establish paths  $P_{si}$  and  $P_{ei}$  between the starting/end points and the user defined concepts T. Once these have been established, the paths  $P_i$  between  $S_0$  and  $E_0$  can be populated and a skeletal pruned ontology can be defined in terms of these paths. All class parents and inferred parents for tagged  $P_i$  path components are also tagged as belonging to the path  $P_i$ . All axioms are loaded for the tagged  $P_i$  path components. This is called the expanded ontology.

[0454] A reasoner is applied to the expanded ontology to identify potential errors and add inferred values. Any concepts, inferences or axioms thus added are included in the pruned ontology 1711, which can now be exported.

[0455] When a concept is selected by the user as the starting point for pruning it is necessary to determine which additional concepts should be included. There are a number of algorithms base on Object Properties and Data Properties which are applied to make this determination. In this regard, object properties have the following attributes:

[0456] They name a relationship between two concepts.

[0457] The relationship has a direction. This is defined as from a 'Domain' concept to a 'Range' concept. In relational database terminology, the primary key of a Domain becomes a foreign key in a Range.

[0458] Optionally the relationship has a type, including:

[0459] Functional

[0460] Inverse Functional

[0461] Transitive

[0462] Symmetric

[0463] Asymmetric

[0464] Reflexive

[0465] Irreflexive

[0466] Also the super/sub class relationship is equivalent to a special case of an object property. A subclass 'inherits' all the Data Properties and all the Object Properties of its superclass.

[0467] Using the sample ontology described above, if the starting point for pruning were 'Club' then it would be

necessary to include all the super classes of Club, namely Organisation and Party in the pruned ontology. The class Member would not be included as the direction and type of that relationship precludes its automatic inclusion. For the same reason the subclasses of Organisation and Party would not be automatically included and neither would any subclasses of club be included had there been any.

[0468] However if Member had been included then the direction and type of the Object Properties 'Has' and 'Holds' would ensure that Club and Individual and all their superclasses were automatically included.

[0469] The Data Property 'Type' in any concept raises a red flag as it implies the existence of an unmodelled concept, viz. 'Type of Club' in Club, 'Type of Member' in Member and so forth. For example the 'Type of Club' concept could contain a list of all the valid values such as Sailing, Chess, Gymnastics etcetera. The Type\_of\_Club concept would have an Object Property called 'Has Type' with Range of Club. This concept would be automatically included in the pruned ontology.

[0470] All automatic inclusions and exclusions can be modified either across all concepts, or on a concept by concept basis. The user specifies 'Include', 'Exclude' or 'Ask' for each type of Object Property.

[0471] The decisions to include a particular concept are made by a specialised Semantic Reasoner using the ontology rules, in particular the Object Properties as input to an inference engine. First order predicate logic is initially used to get explicit inclusions and exclusions. Further inferences as in the example of a 'Type' Data Property must be determined using forwards and backwards inference chaining. To obtain the best result Novamente's probabilistic logic network techniques can be applied to each localised problem area.

[0472] An example of operation of the pruner module will now be described in more detail. In this example, in order to prune an ontology it is necessary to identify the concepts, data properties, object properties and inferences that are included in the ontology. In one example, this is achieved using the indexer module to index the ontology items, and then using the browser module to display the ontology terms for selection as previously described.

[0473] In particular, the user selects the ontology to be pruned in the browser module 'Landing screen'. In this regard, the ontology can be selected from any source, such as a file, Web address, or the like. Once the ontology has been selected the Class List is generated using the index of the ontology. This list displays the name and description of each class. For larger lists a list search function is provided enabling the user to search by class name or part of a class description. It is also possible to search on a data property. In either case the search would return a list of classes which contain that data property. The user then selects a class as the starting point and tags it  $S_0$ .

[0474] Optionally the user then selects an end point  $E_0$ . If the user does not select an endpoint then they will need to manually control the pruning operation as described above. The user may also return to the Landing Screen and select another ontology for the end point or could alternatively add a set of bridging concepts and relationships if they are aware that the chosen ontologies are disparate. If the user does not specify bridging concepts then the process will proceed on

the basis of the overlapping ontologies process described above, otherwise it will proceed as per the disparate ontologies process.

[0475] To control the pruning process, a number of metadata parameters can be set, including:

- [0476] Location to store the pruned ontology.
- [0477] Shell depth for examination.
- [0478] Match quality for accepting sameness.
- [0479] Whether to pause the process at the completion of each shell to allow manual editing.
- [0480] Maximum run time.

[0481] Verbosity of error and log messages.

[0482] An example of the manual pruning process will now be described in more detail.

[0483] In this example, the user only specifies a starting point from which to start the pruning process. They can perform manual pruning in one of two manners, which can be used interchangeably at any time.

[0484] From the Class List screen, typically displayed by the browser module 1310, they can tag classes to be retained with a X'. At any time they can select a 'Validate' option which will automatically tag any related classes and axioms and display the tagged classes in the class list. Additionally they can select a 'View' option which will pass the tagged classes to a graphing program to show the selected classes and relationships graphically. The graphing program can be a publically available graphing packages such as OntoGraf or the like.

[0485] Alternatively the user can open the starting class in the Class Display screen by clicking on the class in the Class List screen displayed by the browser module 1310. The user can then tag all the data properties which they wish to retain, plus any sub/super classes plus any classes specified in the object properties frame. This process can be performed iteratively by clicking on the link to any related class displayed. At any time the user can return to the Class List screen to Validate or View their progress.

[0486] Once the user has finished tagging the classes require for the pruned ontology, they return to the Class List screen and select the "Generate Ontology" option. This results in the pruned ontology being generated in the location specified in the application metadata. The tags can be saved to allow easy re-editing of the pruning process.

[0487] An example of pruning overlapping ontologies will now be described in more detail.

[0488] In this example, the user only specifies starting and end points from which to run the pruning process. The process proceeds as described in the multiple overlapping ontologies as described above.

[0489] Assuming that the application metadata parameters have been set to pause between shells the process will stop as each shell is completed. At this point the user can validate or view the automatically tagged items and may remove any tags that they recognise as irrelevant. Until a path connecting the starting and end points is established the view function will display two partial ontologies. By selecting a "Resume" option the program will start on the determination of the next shell.

[0490] At any time after one Path has been identified the process can be stopped. However alternatively, a number of different possible paths between the start and end points can be determined.

[0491] Once the specified end of processing conditions have been met the process stops and returns to the user with a status message which would include one of the following:

- [0492] Specified maximum shell depth reached. No paths found. Ontologies may be disparate. (Failure)
- [0493] Specified maximum shell depth reached. 'n' paths found. 'm' paths requested. (Partial success)
- [0494] Specified number of paths found. (Complete success)

[0495] The user may decide to extend the process by changing the completion criteria in the application metadata and selecting the Resume option. If the user is satisfied with the result they would select the "Generate Ontology" option. This results in the pruned ontology being generated in the location specified in the application metadata. The tags can be saved to allow easy re-editing of the pruning process.

[0496] If the user decides that the ontologies are in fact disparate then they would proceed as described below.

[0497] In this example, the user specifies starting and end points and a set of related bridging concepts from which to run the pruning process. They may have saved tags from an earlier attempt to prune and merge the ontologies.

[0498] By selecting a commence pruning option the process will start as described in as per the disparate ontology process described above. Assuming that the application metadata parameters have been set to pause between shells the process will stop as each shell is completed.

[0499] At this point the user can validate or view the automatically tagged items and may remove any tags that they recognise as irrelevant. Until a path connecting the starting and end points to one of the user defined bridging points is established the view function will display many partial ontologies, one for each user defined point and one for the starting and end points.

[0500] By selecting a resume option the process starts on the determination of the next shell. At any time after one path in the source ontology, and one path in the target ontology can be connected via the bridging classes the process can be stopped. However alternatively as many paths as possible between the start and end points can be determined.

[0501] Once the specified end of processing conditions have been met the process stops and returns to the user with a status message which would include one of the following.

- [0502] Specified maximum shell depth reached. No paths found. Ontologies may be disparate. (Failure)
- [0503] Specified maximum shell depth reached. 'n' paths found. 'm' paths requested. (Partial success.)
- [0504] Specified number of paths found. (Complete success)

[0505] The user may decide to extend the process by changing the completion criteria in the application metadata and selecting the Resume option.

[0506] If the user decides that the ontologies are in fact still disparate then they would need to spend some effort in examining their bridging concepts. They may need to perform manual tagging to ensure that the paths meet.

[0507] If the user is satisfied with the result they can select a generate ontology option resulting in a pruned ontology being generated in the location specified in the application metadata. The tags can be saved to allow easy re-editing of the pruning process.

### The Semantic Matcher Module

**[0508]** The semantic matcher module enables a mathematical value to be applied to the degree to which two concepts are similar when considered within a particular context. The name for this process is ‘semantic matching’ and it is of particular importance when trying to align the concepts in two ontologies. For example the words ‘company’ and ‘organisation’ in a business context do not have exactly the same meaning. All companies are organisations but not all organisations are companies. In fact the class companies are a subset of the class organisation. For example “This organisation is a listed company but that organisation is a golf club”.

**[0509]** In a social context company is not related to organisation but may be related to a set of associates. For example “John Doe keeps bad company”. A club and a company are both organisations so there is some similarity. A listed company and an unlisted company are also similar and share a common parent. Are they as conceptually close as a club and a company? What about a public unlisted company (>50 shareholders) and a private unlisted company (<51 shareholders)? Are they closer than a listed company and an unlisted company?

**[0510]** To give a mathematical basis to measure how similar two concepts may be we introduce the concept of ‘sameness’. There are a number of formulaic metrics. For example, the Levenshtein distance (Levenshtein, 1966) counts the insertions and deletions needed to match two strings, the Needleman-Wunsch (Needleman, 1970) distance assigns a different cost on the edit operations, the Smith-Waterman (Smith, 1981) additionally uses an alphabet mapping to costs and the Monge-Elkan (Monge, 1996) uses variable costs depending on the substring gaps between the words. Moreover we used the Jaro-Winkler similarity, which counts the common characters between two strings even if they are misplaced by a “short” distance, the Q-Gram (Sutinen, 1995), which counts the number of tri-grams shared between the two strings and the sub-string distance which searches for the largest common substring. However, none of these have proved to be particularly effective.

**[0511]** Another common technique is to arrange the concepts in a single hierarchical tree with the concept of ‘thing’ as the root. Most Sameness formulae are functions of the number of concepts between those being measured and their common parent, and the distance to the root of the hierarchy.

**[0512]** However given the fact that the distance to the root of the hierarchy can differ significantly, depending upon the ontologist who built the ontology and whether the ontology has been pruned by the person using the ontology, the distance to the root is generally irrelevant.

**[0513]** In general, sameness is measured by the number of edges between concepts. Other possibilities exist based upon the number of data properties. For example, a club and a company may have “5” data properties each, the balance being held in the definition of an organisation, whereas a public listed company and a public unlisted company may only have one attribute each, the balance being held in the company definition. Thus a public unlisted company is more similar to a public listed company than a company is to a club (“2” attributes instead of “10”, or in other words there is less difference and difference is equivalent to distance).

**[0514]** The concept of ‘distance’ is considered important. How far apart are two concepts? There are formulae based upon the number of concepts between those being measured

and their common parent. If the distance is “1” then obviously one concept is a superclass of the other. However if the distance is “2” then they are either siblings or grandchildren. This is not a particularly useful fact.

**[0515]** There are some relationships between distance and sameness. Obviously if the distance is “0” then the sameness is “1.0”, in other words, the concepts are identical, so in effect there is only one concept in this instance.

**[0516]** A good semantic matcher module should be able to calculate the sameness and distance of a match using any appropriate formula.

**[0517]** Given that there are many thousands of public and private ontologies describing every aspect of the scientific, engineering and business worlds. In order to align two ontologies it is necessary to determine whether there is a semantic match between the concepts in the two ontologies.

**[0518]** Currently the manipulation of ontologies defining linked concepts is confined to academics and professional ontologists. Definitions and names of concepts vary enormously depending upon context. In order to compare terms in and across ontologies we need to have some mechanism for examining the terms semantically. Are two concepts actually synonyms for the same thing or are they related in some other way. For example, organisations and companies have some attributes in common so there is some degree of sameness. All companies are organisations but not at all organisations are companies (Subsumption).

**[0519]** In another example the existence of fingers implies the existence of hands. Although they are not the same there is a relationship between them and the existence of one implies the existence of the other because one is a part of the other (Meronym).

**[0520]** Given any two concepts we would like to know how similar they are; i.e. Sameness 0->1 where 1.0 implies they are identical, whether one is a subclass or superclass of the other (-1,0,1), and whether one is a part of another (-1,0,1).

**[0521]** The semantic matcher module includes a database of concepts, their meaning and relationships between them. It has tools for loading the concepts from ontologies, for manually editing the relationships between concepts and their definitions and for analysing concepts in a mathematically defined manner. These mathematically defined properties of concepts and their relationships can then be used in a variety of situations, such as aligning ontologies, as a dictionary and as a semantic concept matcher module.

**[0522]** The semantic matcher module concept finds synonyms, subsumptions (class hierarchy) and meronyms (part of) in a particular context (e.g. Medical, Business). It is initially loaded by parsing an ontology and obtaining the classes, their annotations, class structure and any ‘part-of’ Object properties. The class name is then used in something such as WordNet or Watson to determine the meaning and possible synonyms. The meaning is parsed into triples, as are any notations. The matcher module then looks for mathematical correspondences in the triples determine synonymity.

**[0523]** The semantic matcher module is a stand-alone process which either evaluates two lists of concepts, typically from two ontologies or else evaluates a single concept, matching this against reference terms to determine a meaning for the concept.

[0524] In the first instance the matcher module will pair each item in the first list with each item in the second list. Each pair  $i,j$  is then analysed to determine the following items:

- [0525] The semantic similarity  $S_{ij}$ .
- [0526] If The terms are synonyms then the similarity is  $S_{ij}=1.0$ .
- [0527] If Antonyms then  $S_{ij}=-1$ .
- [0528] If there is no relationship then  $S_{ij}=0$ .
- [0529] The subsumption relationship  $\text{Sub}_{ij}$ .
- [0530] If  $C_i$  is a subclass of then  $\text{Sub}_{ij}=-1$ .
- [0531] If  $C_i$  is a superclass of then  $\text{Sub}_{ij}=1$
- [0532] else  $\text{Sub}_{ij}=0$ .
- [0533] The meronym relationship  $\text{Mer}_{ij}$ .
- [0534] If  $C_i$  is a part off  $C_j$  then  $\text{Mer}_{ij}=1$ .
- [0535] If  $C_j$  is a part of  $C_i$  then  $\text{Mer}_{ij}=1$
- [0536] else  $\text{Mer}_{ij}=0$ .

[0537] In the second instance the matcher module takes a single concept and a context definition and produces a list of synonyms, sub and superclasses and meronyms for that concept in that context. If the context is not supplied the evaluation is performed across all contexts.

[0538] Some examples follow based upon the presumption that a medical ontology and a Human Resources Ontology have been defined to SemMatch:

- [0539]  $\text{SemMat}(\text{Party}, \text{Client}, \text{Business})=(1.0, 0, 0)$
- [0540]  $\text{SemMat}(\text{Party}, \text{Individual}, \text{Business})=(0.25, 1, 0)$
- [0541]  $\text{SemMat}(\text{Individual}, \text{Client}, \text{Business})=(0.25, -1, 0)$
- [0542]  $\text{SemMat}(\text{Car}, \text{Engine}, \text{Automotive})=(0.1, 0, 1)$
- [0543]  $\text{SemMat}(\text{Car}, \text{Wheels}, \text{Automotive})=(0.1, 0, 1)$
- [0544]  $\text{SemMat}(\text{Patient}, \text{Person}, \text{Medical})=(0.25, -1, 0)$
- [0545]  $\text{SemMat}(\text{Patient}, \text{Person}, \text{HR})=(0, 0, 0)$
- [0546]  $\text{SemMat}(\text{Patient}, \text{Person},)=(0.25, -1, 0)$
- [0547]  $\text{SemMat}(\text{Person}, , \text{Medical})=\text{Definition: A single human being}$
- [0548] Synonyms: Individual, Body
- [0549] SuperClass: Entity, Role

- [0550] SubClass: Patient, Practitioner, Performer
- [0551] Meronyms: -1, None
- [0552] +1, Organs, Limbs

- [0553]  $\text{SemMat}(\text{Person}, ,)=\text{Context: Medical}$
- [0554] Definition: A single human being
- [0555] Synonyms: Individual, Body
- [0556] SuperClass: Entity, Role
- [0557] SubClass: Patient, Practitioner, Performer
- [0558] Meronyms: -1, None
- [0559] +1, Organs, Limbs
- [0560]  $\text{SemMat}(\text{Person}, ,)=\text{Context: HR}$
- [0561] Definition: A single human being
- [0562] Synonyms: Individual
- [0563] SuperClass: Entity, Party, Involved Party
- [0564] SubClass: Employee
- [0565] Meronyms: -1, Family
- [0566] +1, None

[0567] The two different usage methods will now be described in more detail with reference to FIGS. 18A and 18B.

[0568] The Semantic Matcher module 1350 uses a Concept Matching Database 1604 to perform its evaluations. In the example of FIG. 18A, two lists of concepts 1801, 1802, such as ontology terms A, B and X, Y are received and then compared by the semantic matcher module 1350 to generate sameness scores 1803 for each possible pairing of ontology terms.

[0569] In the example of FIG. 18B, a single concept, such as a single ontology term 1804 is received, and the semantic matcher module 1350 compares this to the concept matching database 1604 and returns a list of synonyms 1805.

[0570] The concept matching database (CMD) 1604 is constructed using the indexer module 1320. Before it can be used the database must first be loaded, which is typically it would be loaded by parsing an ontology based upon the context of interest. The database can be updated by the user at any time to add new contexts.

[0571] The CMD 1604 contains a number of tables as defined in Table 8, with the relationships between the tables being shown in 18C.

TABLE 8

Table	Column	Description
Cept_W_Ctext	Word	The name of a concept from a particular source.
	Word_ID	An automatically generated unique computer key.
	Meaning	A paragraph defining the meaning of this version of the word.
	Meaning_RDF	The meaning above transformed to RDF triples.
	Source_ID	The Ontology from which the word was sourced.
	Word_ID	An automatically generated unique computer key
Concept	Concept ID	An automatically generated computer key which is updated to ensure that synonyms all have the same key.
	Lemma	Boolean switch showing whether the word is the main root word for synonyms.
	Context_ID	Foreign key identifying context. The context in which the concepts have these meanings and synonyms.
	Concept	A concept name. It may be more than one word. For example 'Involved Party'.
Context	Concept ID	An automatically generated computer key which is updated to ensure that synonyms all have the same key.
	Context	Name of a context. Typically the name of an ontology e.g. SNOMED CT, HL7 RIM.
	Context_ID	An automatically generated unique computer key.

TABLE 8-continued

Table	Column	Description
ContextSource	Context_ID	Foreign key identifying context.
	Source_ID	Foreign key identifying source.
Source	Source_ID	An automatically generated unique computer key.
	Address	Typically the URL/URI of the site from which the ontology use to load the database was obtained.
Relation Type	Relation_Type_ID	An automatically generated unique computer key.
	Name	The name of the Object Property used in the relationship between the two CWCs e.g. ‘subclassof’ or ‘ispartof’.
Meronyms	Description	A description of the Object Property e.g. subsumption, meronym.
	CCW_ID_P	The input concept key
	CCW_ID_C	The concept key of the concept of which the input concept key is a part i.e. Concept_ID is part of Part_of_ID.
	Relation_Type_ID	The concept key of the concept of which the Includes_ID concept key is a part i.e. Includes_ID is part of Concept_ID.
Word to Word	Word_ID_P	The parent word key.
	Word_ID_C	The child word key-a synonym from a different source, typically Wordnet.

[0572] The load mechanism will now be described in detail with reference to FIG. 18D.

[0573] Initially, an overall context of the ontologies 1801 to be loaded is determined and entered into the Context table with an ID of 1. For example, if medical ontologies are loaded, the context would be identified as “medical”.

[0574] An example of the ontologies in this category and the context name for each as shown below:

- [0575] Adverse Event Reporting Ontology AERO
- [0576] African Traditional Medicine Ontology ATMO
- [0577] Allen Brain Atlas (ABA) Adult Mouse Brain Ontology ABA-AMB
- [0578] Alzheimer’s disease ontology ADO
- [0579] Amino Acid Ontology AMINO-ACID
- [0580] Amphibian Gross Anatomy Ontology AAO
- [0581] Amphibian Taxonomy Ontology ATO
- [0582] Anatomic Pathology Lexicon PATHLEX
- [0583] Anatomical Entity Ontology AEO

[0584] Each of these ontologies has a source which will be loaded into the Source table thus allowing the Source 2 Context table to also be loaded.

[0585] Next, the following information is extracted and parsed from each of the ontologies:

- [0586] Classes
- [0587] Object Properties
- [0588] Annotations
- [0589] Labels

[0590] As all words are coming from one ontology the Context\_ID is known. Each Class becomes a Word in the Word table. The Annotations are loaded as the Meaning in the Word table. Temporary tables are created relating Word\_ID 2 Context\_ID with lemma (root meaning) and Concept, both set to null, and Class2Object-Property2Class with Word\_IDs for each class and Concept\_ID set to null.

[0591] Following this, the extracted classes and their annotations are then loaded into Word table. Each Class becomes a Word. Each Word is assigned with a unique Word\_ID and a class annotation becomes the Meaning in the Word table. As all words are coming from one ontology the Context\_ID is known as previously described.

[0592] Temporary tables are created relating Word\_ID2Context\_ID with lemma and Concept, both set to null,

and Class2Object-Property2Class with Word\_IDs for each class and Concept\_ID set to null.

[0593] For each context, the first step is to match each word to a meaning and synonym obtained from a standard dictionary, such as the WordNet 1802. Any unmatched words are then matched against words from other contexts to identify synonyms. These steps are now described in more detail.

[0594] Each word in the Word table is passed to WordNet 1802 to obtain a meaning and potentially the root word or lemma for the group of synonyms or lexeme, based upon that Word. The WordNet meaning is lexically compared with the meaning derived from the annotation.

[0595] This is done by converting the meaning to RDF triples and evaluating the triples. This process is described in more detail below.

[0596] If the meanings match then the Wordnet Word and Meaning are loaded into the Word table with a new Word\_ID. The new Word\_ID is assigned to Word\_ID\_C and the original Word\_ID is assigned to Word\_ID\_P both are then loaded into the Word2Word.

[0597] The Word\_ID2Context\_ID table is loaded with the Word\_ID assigned to the Wordnet Lemma as the Word\_ID and the same Context\_ID as the related Word\_ID, which was loaded as the Word\_ID\_P. The Word\_ID2Context\_ID table has only two columns lemma and concept. So the lemma is assigned with new Word\_ID\_C and concept is assigned from Word\_ID\_P.

[0598] Finally the Class2Object-Property2Class is loaded with the Word\_ID information from Wordnet 1802.

[0599] All words for which a Lemma was defined are then loaded into the Concept table. The Word\_ID2Context\_ID can now be updated with known Concept\_ID and Lemma and used to load the Concept\_Word\_Context table resulting in the CWC\_ID being assigned to each Concept and Word used in the named Context. The CWC\_ID can be used to identify the words in the Class2Object-Property2Class and together to populate the CWC2CWC table and the Relation\_Type table.

[0600] A second pass of the Word table examines the meanings of every word for which there is no related lemma, by syntactically comparing the meaning with the meanings

of words in the other contexts. The Word\_ID of the first meaning to match is chosen as the lemma. The process then continues as for Wordnet identified lemmas.

[0601] A third pass simply identifies each word which is not related to a lemma as being a lemma. At the completion of these three passes every word will have been identified in every possible context in the concept table 1809.

[0602] Following this a sameness value is calculated. If the full ontology were known then the calculation of Sameness could be performed by matching the attributes (Data Properties) of the concepts being compared. The attribute list would of necessity include the attributes of the super-classes of the concepts.

[0603] In the current example sameness is calculated by analysing the meaning of two words. The meaning in English is converted to rdf triples of the form Subject Predicate Object (spo). This is done using a Natural Processing Language (NLP) to RDF converter. (Arndt & Auer, 2014) (Augenstein, et al., 2013).

[0604] For example—a club has meaning “A type of organisation which has members, not shareholders and exists to meet some vocational need of its members” could be converted as shown in table 9 below:

TABLE 9

Subject	Predicate	Object
Club	Is a	Organisation
Club	Has	Members
Members	Have	need
Needs	Are	vocational
Club	Meets	needs

[0605] An organisation is a concept which is defined as follows; “An organisation is a collection of individuals with an agreed reason for being their collection”, which could be converted as shown in Table 10.

TABLE 10

Subject	Predicate	Object
Organisation	Is a	Collection of individuals
Organisation	Has	individuals
Individuals	Have	Agreed reason for being a collection

[0606] Inserting the Organisation definition into the Club definition we obtain the definition shown in Table 11.

TABLE 11

Subject	Predicate	Object
Club	Is a	Organisation
organisation	Is a	Collection of individuals
organisation	Has	individuals
Club	Has	Members
Members	Have	needs
Needs	Are	vocational
Club	Meets	needs
Individuals	Have	Agreed reason for being a collection

[0607] However we cannot infer that a member is an individual. Analysis of this can be used to determine that:

[0608] A Member of a Club is an Individual. This could have been inferred if the Membership concept had the Object Properties more correctly defined as Member is An Individual instead of Individual Holds Membership.

[0609] The agreed reason for being a collection is to meet vocational needs.

[0610] Applying the same process to a Qango in the example ontology described above we would obtain from the Meaning that a Qango is “an organisation created by a government to meet a specified government need”, leading to the triples shown in Table 12.

TABLE 12

Subject	Predicate	Object
Qango	Is a	Organisation
Organisation	Is a	Collection of individuals
Organisation	Has	Individuals
Qango	Created by	Government
Government	Has	need
Qango	Meets	need
Individuals	Have	Agreed reason for being a collection

[0611] This can be used to construct a comparison table based upon common predicates and objects as shown in Table 13.

TABLE 13

Predicate Object	Club	Qango
Is A organisation (+2 other organisational matches)	Y	Y
Created by Government	N	Y
<b>Created by Members</b>	Y	N
Meets Needs	Y	Y
Specified by Government	N	Y
Specified by Members	Y	N

[0612] This allows a formula for sameness to be used based upon the following factors.

[0613] Number of triples for concepts of Club and Qango are denoted by N1 and N2 respectively where N1=9 and N2=7

[0614] Number of shared predicates (SP) between the two concepts Club and Qango is 5, i.e. SP=5

[0615] Number of shared predicate object (SPO) pairs between the two concepts Club and Qango is 4, i.e. SPO=4

[0616] For example:

[0617] Sameness=SPO/SP=4/5=0.8 OR

[0618] Sameness=(SP+SPO)/(N1+N2)=9/16=0.5625

[0619] The actual formula used is irrelevant. The important fact is that we can derive a formula which gives a measure of Sameness.

[0620] It will be appreciated that throughout this process the user can interact with the semantic matcher module using screens 1808, typically displayed by the browser module.

#### Aligner Module

[0621] The need for ontology alignment arises out of the need to integrate heterogeneous databases, ones developed independently and thus each having their own data vocabulary. In the Semantic Web context involving many actors providing their own ontologies, ontology matching has

taken a critical place for helping heterogeneous resources to interoperate. Ontology alignment tools find classes of data that are “semantically equivalent”, for example, “Truck” and “Lorry”. The classes are not necessarily logically identical.

[0622] The result of an ontology alignment is a set of statements representing correspondences between the entities of different ontologies. This may be expressed in the purpose built language ‘Expressive and Declarative Ontology Alignment Language’ (EDOAL) (David, et al., 2013) or other languages (ZIMMERMANN, et al., 2006).

[0623] The first requirement is to determine if there is a semantic match between the concepts in the ontologies being aligned, which can be determined using the semantic matcher module described above. For example the words ‘company’ and ‘organisation’ in a business context do not have exactly the same meaning. All companies are organisations but not all organisations are companies. In fact the class companies is a subset of the class organisation. For example “This organisation is a listed company but that organisation is a golf club”. In a social context company is not related to organisation but may be related to a set of associates. For example “John Doe keeps bad company”.

[0624] A club and a company are both organisations so there is some similarity. A listed company and an unlisted company are also similar and share a common parent viz. company. Are they as conceptually close as a club and a company? What about a public unlisted company (>50 shareholders) and a private unlisted company (<51 shareholders)? Are they closer than a listed company and an unlisted company?

[0625] To give a mathematical basis to measure how similar two concepts may be we introduce the concept of ‘sameness’. There are a number of formulaic metrics for sameness. The most common technique is to arrange the concepts in a single hierarchical tree with the concept of ‘thing’ as the root. Most formulae are functions of the number of concepts between those being measured and their common parent, and the distance to the root of the hierarchy.

[0626] However given the fact that the distance to the root of the hierarchy is can differ significantly, depending on the ontologist who built the ontology and whether the ontology has been pruned by the person using the ontology, the distance to the root is probably irrelevant.

[0627] In general, sameness is measured by the number of edges between concepts. Other possibilities exist based upon the number of data properties. For example, a club and a company may have 5 data properties each, the balance being held in the definition of an organisation, whereas a public listed company and a private listed company may only have one attribute each, the balance being held in the company definition. Thus a private listed company is more similar to a public listed company than a company is to a club. (2 attributes instead of 10, or in other words there is less difference and difference is equivalent to distance.)

[0628] A Putative Ontology (PO) is an ontology created from a structured source, typically a relational database, an xml file or a spreadsheet. Such an alignment may have some very complex mappings in which data instances in the putative ontology map to classes in the full ontology. This is a special case of alignment.

[0629] A simple example will now be described with reference to FIG. 19A, which shows a “Thing Database”, which is an example of a totally de-normalised data structure

as it can contain the metadata (and hence structure) as well as the data within four tables.

[0630] For example, if the Thing Type table contains a Thing Type of ‘Class’, then every related row in the Thing table would contain the name of a class. The relationship between classes would be defined in the ‘Thing to Thing’ table where the ‘Thing Type to Thing Type’ specifies the type of relationship.

[0631] In ontological terms, any Type table can give rise to a set of classes. Consider a table containing details of a set of vehicles. A vehicle type table could have been used to ensure that only valid types of vehicles are included. For example Cars, trucks, tractors but not prams, bicycles, ships. Ontologically, we could then have a separate class for each type of vehicle specified in the Vehicle Type table. This concept can be generalised but is not always appropriate. It could result in every personnel table being split into male and female classes! Consequently the program should identify every situation in which hidden classes contained in the data can be exposed and present them to the user for validation.

[0632] In some cases the Type table may contain many types of types. For example Concepts, Data Properties and Properties of Data Properties, such as Vehicles, trucks, Cars, engine type, weight, kilograms. This could be shown as:

[0633] Car has engine type diesel

[0634] Car has weight 2000

[0635] Weight has Unit of Measure kilograms

[0636] Car is subclass of Vehicle

[0637] An example of the thing database will now be described assuming the database is populated as shown in Tables 14 to 17.

TABLE 14

Table	Thing		
	ID	Name	Thing Type ID
1	Fingers		A
2	Hand		A
3	Person		A
4	Living Organism		A
5	Organisation		B
6	Individual		B
7	Client		B

TABLE 15

Table	Thing to Thing		
	Thing ID_P	Thing_ID_C	Thing Type to Thing Type ID
1	2		Aa
2	3		Aa
3	4		Aa
3	6		Cc
6	7		Bb
5	7		Bb

TABLE 16

Table	Thing Type Name
A	Organic Structure
B	Business Component

TABLE 17

Thing Type to Thing Type			
ID	Thing_Type_ID_P	Thing_Type_ID_C	Name
Aa	A	A	Is Part of
Bb	B	B	Is a
Cc	A	B	Is the same as

[0638] A putative Ontology based on the Relational Schema would only show four classes with names related to the table names. However, an ontology based upon the data would show eight classes based upon the names in the ‘Thing’ and ‘Thing Type’ tables, plus all the Object Properties identified in the other two tables, as shown in FIG. 19B. In this example, the “business component” and “organic structure” terms are obtained from the thing type table (Table 16), whereas the remaining terms are obtained from the thing table (Table 14).

[0639] This is an example of the problem where the classes in one ontology match to data instances in another ontology. For clarity this is identified as a ‘Putative Mapping Problem’ (PMP). It can manifest during alignment when the putative ontology has data properties with names matching ‘Primary Key’ or ‘Foreign Key’, or a class with multiple instances of the same foreign key, as in ‘Parent’ and ‘Child’ (BOM) or a class with an associated type class. These examples potentially disguise a Class hierarchy hidden in Data Instances!

[0640] A common alignment technique is to arrange the concepts from each ontology into two hierarchical trees, each with the concept of ‘thing’ as the root. The mathematical concept of ‘Distance’ is then introduced to give some mathematical mechanism for determining alignment. Most Distance formulae are functions of the number of concepts between those being measured and their common parent, and the distance to the root of the hierarchy.

[0641] However given the fact that the distance to the root of the hierarchy can differ significantly, depending upon the ontologist who built the ontology, whether the ontology has been pruned by the person using the ontology, and whether there is a ‘top’ ontology acting as a conceptual umbrella, the distance to the root is probably irrelevant.

[0642] The ontology aligner module looks for common concepts in multiple ontologies and maps the concepts from one ontology to the other thus allowing the two ontologies to be treated as one ontology. Using the alignment it is also possible to merge the two ontologies although this is a risky process and is not generally recommended due to the potential for semantic mis-match propagation.

[0643] In general no ontology is perfect. For example there are many modelling errors in the sample ontologies used here. It is obvious that ‘Shares’ should be ‘Owned’ by ‘Clients’ rather than ‘Individuals’ and that ‘Work History’ should be ‘Employed’ by ‘Client’ rather than ‘Company’. Both these instances show that the relationship is moving from a more restrictive relationship to a less restrictive relationship. Although that would be possible in these cases it would probably be invalid to move membership of a club from an ‘Individual’ to a ‘Client’.

[0644] The class ‘Membership’ is also badly named as the relationship between Membership and Individual is ‘Holds’. If the Class had been named ‘Member’ then the relationship would have been ‘is A’. This would have allowed the

Member to inherit the Properties of an Individual. Unless the Object Property ‘Has’ is fully defined then it use in inferencing is restricted.

[0645] These errors were introduced to the samples to illustrate some of the complexities of alignment.

[0646] Operation of the aligner module will now be described in more detail with reference to FIG. 19C.

[0647] In this regard, in use, ontologies 1901, 1902 defined in OWL and RDFS files are opened using the aligner module 1340, with the user then interacting with the ontology using a set of screens as defined below, ultimately resulting in ontologies 1903, 1904 connected by a series of alignments 1905 and potentially a merged aligned ontology 1906.

[0648] The process consists of a number of sub processes, including:

[0649] Initialisation

[0650] Low level Class matching—identifies minimal mappings

[0651] Putative Mapping Problem Identification

[0652] Object Property Analysis

[0653] Data Property Analysis

[0654] Multi Class mappings

[0655] PMP resolution

[0656] Sibling Analysis

[0657] Minimal mapping resolution

[0658] Because an alignment can be identified in many steps there is the potential to recalculate the alignment for a particular pair of concepts. This problem is overcome by maintaining an Alignment Map. This map is updated every time an alignment is identified and is consulted by the program before a new alignment pair is considered for evaluation to prevent duplication of effort. The Alignment Map can be displayed to the user enabling them to follow the alignment process, query and override any potential alignment and instruct the program to re-perform any process.

[0659] These steps will now be described in more detail. Each step i can be assigned a weighting factor  $W_i$ , with the results being combined to provide an overall alignment score. These weighting factors are applied at certain steps. A possible Weight Accumulation formula is given, but there are many possible weighting schemes that could be used. This is an area where machine learning or statistical analysis and inferencing can be used to determine suitable weighting formulas.

[0660] During the initialisation process, an index 1603 is obtained from the Indexer module. Following this the ontologies 1901, 1902 are loaded into the semantic matcher module 1340. When the alignment table has not been pre-loaded then  $W_0=0.0$ .

[0661] In the following examples  $W_i=i$  for illustration of the technique. Otherwise the weights  $W_i$  are assigned by the user or a heuristic mechanism determined by machine learning or experience. In general for any step i the accumulatively determine Match Value  $MV_i^A$  is determined by:

$$MV_i^A = MV_{i-1}^A/W_i + (W_i-1)*MV_i/W_i$$

[0662] where  $MV_i$  is the raw Match Value calculated in step i.

[0663] Another, more traditional weighting scheme would be:

$$MV = \sum MV_i * W_i / \sum W_i$$

[0664] where  $MV$  is the weighted match value and  $MV_i$  is the match value at step i.

[0665] This process could be performed at each step or only at the end of the procedure, depending on the preferred implementation.

[0666] Next, class matching is performed on the basis of the semantic meaning of terms in the ontologies. This process examines each potential alignment pair using the semantic matcher module to find a potential match based on the class name. If it finds an alignment it then traverses the inheritance chains (Object Property='SubClassof') from that alignment, checking the class names for another alignment using the semantic matcher module.

[0667] This may only require a small number of matches although it is possible to find all matching classes. A complete 1-1 match is possible if the ontologies being matched are using the same basic ontology. For example:

[0668] Adverse Event Reporting Ontology AERO

[0669] African Traditional Medicine Ontology ATMO

[0670] Both are based on the standard Galen ontology so one would expect a 1-1 match.

[0671] MV for each pair is based upon the score provided by the semantic matcher module and Set  $W_1=1.0$  for purposes of this example.

[0672] Starting at the root of the first Ontology, examine each class starting at the root class of the second ontology. A match occurs when the sameness found using the semantic matcher module for the concept pair exceeds the threshold Match Value for alignment ( $MV_{AT}$ ). If an acceptable match is found it is called a potential alignment and the details are recorded in the alignment map.

[0673] The alignment map records the two concepts, assigns an alignment Id, a minimal map Id, any tags associated with the alignment, any PMP Id assigned, any enrichment Id and the last processing step Id. A separate table, related on the Alignment Id stores the Match Value for each step. These values can be manually overridden if desired.

[0674] The alignment map may be pre-loaded with any known alignments. These are tagged with a user tag 'User Initiated' and the Match value must be set, generally to 1.00 although lower values are possible. The combination of 'User Initiated' and  $MV=1.00$  will prevent further processing of this alignment.

[0675] The process continues to the next class related to the current class in the first Ontology by an Object Property. Superclasses of the current class are processed first. The program processes Inheritance Object Properties before other Object Properties. Superclasses of the current class are processed before any subclasses are examined. The process stops as soon as an alignment with  $MV < MV_{AT}$  is found.

[0676] Each time a potential alignment is identified it is assigned to a minimal mapping set and given a minimal map Id mm\_ID. If a hierarchically related class is identified it is added to the same mm\_ID. At the end of this step we will defined have a number of minimal maps which potentially meet the criteria of minimal mapping. This cumulative match value is refined at each succeeding step.

[0677] The recognition of a potential PMP is always performed. PMP resolution is only performed if requested in a configuration file. If not requested the recognition of the potential PMP is recorded in the activity log created as the alignment is performed as an Information Message and is added to the cumulative statistics report.

[0678] In some instances it may not be desirable to resolve the PMP as both ontologies may be putative ontologies and it may be desirable to retain the BOM structure.

[0679] If PMP resolution was requested then PMP tagging is performed. The Data Property names are examined for the existence of key words such as:

[0680] Object Property names contain:

[0681] Type

[0682] Relation

[0683] Class

[0684] Concept

[0685] . . .

[0686] Data Property names contain:

[0687] Identifier

[0688] ID

[0689] Key

[0690] Parent

[0691] Child

[0692] Primary Key

[0693] Foreign Key

[0694] . . .

[0695] The existence of data properties containing these key words does not necessarily imply a PMP. Further algorithms need to be applied to be certain. Any structure which maps to a standard.

[0696] 'Type' table in ERA diagrams must be identified.

The user must select each row in the type table which is to be.

[0697] 'Bill of Materials' structure must be identified and potentially expanded into the appropriate class structure.

[0698] At this stage the classes involved in each PMP are tagged as 'PMP' and given a PMP-set-identifier PMP01, PMP02, . . . for each set of equivalent BOM tables. They are resolved later on, as will be described in more detail below. As each PMP class is identified the details may be presented to the user who may decide that that instance is not a PMP.

[0699] No MV is calculated for this step so  $MV_2^A = MV_1^A = 0.5$ .

[0700] Following this, object properties and their related classes associated with the each alignment pair from the previous steps are analysed. This step is sometimes called 'Structural Analysis'. This will identify:

[0701] If the names of all the related classes and the Object Properties match then tag the pair as an "Anchor Point".  $MV=1.0$ . Add the related classes to the minimal map if they are not already there and repeat step 2 Data Property Analysis for the related classes in that Minimal Map

[0702] If the name and related super class match but not any of the subclasses then tag the pair as "Possible Siblings".  $MV=0.3$ . Add the Super class to the Minimal Map. Go to multi class mappings below.

[0703] If the name and related super class match but only some of the subclasses match then tag the pair as "Related subset".

[0704] MV is calculated as follows:

[0705] Assign weights of 2.0 to each matching subclass and 1.0 to each other matching related class.

[0706] Sum these weights as the Number matching  $N_M$ .

[0707] Assign weights of 1.0 to each subclass and 0.5 to each other related class.

[0708] Sum these weights across both superclasses as the Total Number  $N_A$ .

[0709] The Match Value  $MV_3 = N_M/N_A$

[0710] If no related classes match then the  $MV_3 = 0.001$ .

- [0711] Add the Super class to the Minimal Map. Go to multi class mappings below.
- [0712] For each pair calculate the cumulative weighted Match Value as below:
- $$MV_3^A = MV_2^A/W_3 + (W_3-1)*MV_3/W_3$$
- [0713] Assume  $W_3=3$
- [0714]  $MV_2^A=0.5$  from previous example
- [0715]  $MV_3=1.0$  from Object Property match
- [0716] Then  $MV_3^A=0.5/3+2/3*1.0=0.83333$
- [0717] Following this, data property analysis is performed to analyse whether the data properties (attributes) of matching classes are similar. The analysis, for each pair of classes:
- [0718] Compare the Data Properties for each class using SemMat where there is no exact name match.
- [0719] Assign a "Match Value" (MV) based on the Data Properties.
- [0720] Tag the alignment pair with a match type. Select the next pair in the minimal map and repeat the processes above. If there are no more alignments within the minimal map, move to the next minimal map.
- [0721] In more detail, if  $A=\{a_1, a_2, a_3, \dots, a_j\}$  is a set of Data Properties of the first concept and  $B=\{b_1, b_2, b_3, \dots, b_j\}$  is a set of Data Properties of the second concept then the following possibilities exist:
- [0722] All Data Properties in the classes match. Tag as "Exact Match" i.e.  $\forall a \in A = \forall b \in B$
  - [0723] Match Value=1.00.
  - [0724] A subset of Data Properties from one ontology match all the Data Properties in the other ontology. Tag as "Subset".
  - [0725] i.e.  $A \subset B$  or  $\forall a \in A = \exists b \in B$
  - [0726]  $MV_i = (N(A \cap B)/N(B))^{0.5}$  where  $N(A)$  is the number of Data properties in A, assuming  $N(A) < N(B)$
  - [0727] A subset of Data Properties from one ontology match a subset of Data Properties in the other ontology. Tag as "PartMatch"
  - [0728] i.e.  $\exists a \in A = \exists b \in B$
  - [0729]  $MV_i = N(A \cap B)/N(B)$  where  $N(A)$  is the number of Data properties in A, assuming  $N(A) < N(B)$
  - [0730] No Data Properties match. MV=0.1, Tag as "NameOnly"
  - [0732] i.e.  $\forall a \in A \neq \forall b \in B$
- [0733] If MV is less than a predetermined threshold, (default value=0.1) then discard that match pair from the Minimal Map and proceed with the next match pair. This process is repeated until all minimal maps have been analysed, at which point a matching value is calculated:
- $$MV_4^A = MV_3^A/W_4 + (W_4-1)*MV_4/W_4$$
- [0734] Assume  $W_4=4$
- [0735]  $MV_3^A=0.833333$
- [0736]  $MV_4=1.0$  from Data Property match
- [0737] Then  $MV_4^A=0.8333/4+3/4*1.0=0.9583$
- [0738] Multi class mappings occur when the class in one ontology has been split into a number of subclasses in another ontology. In such cases we would expect the pair to be have already been tagged as either "Possible Siblings" or "Multi Class Mappings" and "Subset".
- [0739] The multiclass mapping is usually detected by analysing the number of Data Properties for the potentially related classes in the class and sub classes in each ontology. If the ontology class which does not have a subclass has the number of Data Properties approximately equal to the class

in the other Ontology plus the Data Properties of the sub-class with the most Data Properties then it is probable that the sub classes of the class in the second ontology have been denormalised into the class in the first ontology.

[0740] There are the following possible scenarios.

[0741] The Data Properties in single class in one ontology maps to the Data Properties in a class and subclass or subclasses in the other ontology.

[0742] The Data Properties in a class and subclasses match to the Data Properties in a class and some subclasses in the other ontology.

[0743] In the first case the Data Property count is performed by considering matching the Data Properties of the first ontology class with the Data Properties of each pair composed of Class+Subclass in the second ontology.

[0744] For example, Company in Ontology 1 has no children and 2 children in Ontology 2. If analysing the Data Properties of company (1) with Company+Listed Company (2) shows that the number of Data Properties match but not all the meanings.

[0745] Analysing the Data Properties of company (A) with Company+Unlisted Company (B) shows that both the number and meaning of Data Properties match. This can be tagged as 'Different Normalisation' and assigned a matching value MV=1.0.

[0746] Because the Listed Company and Unlisted Company are siblings it is possible to infer that the Listed Company is an enrichment in ontology 2 so it can be tagged as 'Enrichment' and the matching value calculated by dividing twice the number of matching Data Properties by the total number of Data Properties.

$$MV_i = 2 * N(A \cap B) / (N(A) + N(B))$$

[0747] where  $N(\cdot)$  is the function to produce the number of Data properties in concepts A, B and  $A \cap B$ .

[0748] This method can be generalised to the situation where the two classes have a different number of children. This situation can be tagged as 'Enrichment Possible' and each class involved is given a single enrichment ID.

[0749] Another case of multiclass mappings is when classes have been normalised differently. For example, a Vehicle class could be subclassed as (SUV, Sedan, Coupe, Convertible) or it could be subclassed by manufacturer (Citroen, Peugeot, Fiat, Rover). Thus two vehicle ontologies could parse the data properties differently. However, the attributes of vehicles would be identical in the two ontologies.

[0750] In the general case, if a set of Data Properties are assigned to a set of sub Classes from two ontologies, and the sub Classes are different in each ontology but the set of Data properties defining these classes are either identical or very similar, then there is a many to many mapping between the subclasses defined. This is also tagged as 'Enrichment Possible' and each class involved is given a single enrichment ID.

[0751] No MV is calculated for this step so  $MV_5^A = MV_4^A = 0.9583$ .

[0752] PMP resolution involves identifying additional classes in putative ontologies by identifying denormalised classes stored in tables, and results in a major enrichment of the ontology from which it was derived.

[0753] Each PMP-set-identifier is analysed to determine its mapping to a Type structure or a BOM structure as described above. These generally map to some arrangement

of the ERA diagram shown in FIG. 19A, as is determined by mapping only the Object properties with the matching structural relationships in that diagram. An example of extracted classes from Data Property instances is shown in Tables 14 to 17.

[0754] Once the mapping is determined it is a relatively simple matter to generate the denormalised ontology captured in the BOM structure. This generated ontology component can then be aligned by returning to the step of low level class matching based on the semantic meaning of the classes as previous described. In this step the Classes generated from the BOM analysis will be added to the appropriate minimal maps.

[0755] No MV is calculated for this step as it results in the return to the step of low level class matching and the recalculation of the MV values for the newly identified classes.

[0756] Following this enrichment analysis is performed, with each enrichment\_ID identified in the multi class mappings process being analysed to determine whether the subclass sets from the two ontologies match or contain siblings. For example ontology 1 class organisation may have sub classes Club and Company. Ontology 2 contains Qango, Club and Company. Qango is a sibling in Ontology 2 but does not appear in ontology 1. Rather than say the Qango does not align with anything it would be better to identify it as an enrichment to Ontology 1.

[0757] Before the enrichment can be applied it would necessary to determine whether the Qango has been denormalised into one of the other subclasses by analysing the Data Properties of Club and Company.

[0758] Assuming that the class meets the criteria to be added as a sibling it should be possible to ensure that the minimal maps containing the class and subclass are identical at this stage.

[0759] No new MV is calculated for this step. Each sibling retains its current MV. This MV could be raised by a small factor by assigning a current MV of 1.0 to components identified as siblings.

[0760] Once all classes are resolved and enrichment completed, any major restructuring should have already occurred and accordingly, minimal maps can be resolved. Further restructuring occurred if enrichments were added in the previous section. Both these facts would result in improve minimal mappings.

[0761] Alignments with  $MV_7 < MV_{AT}$  the threshold would be rejected.  $MV_{AT}$  is the Match Value threshold for alignment.

[0762] The next step is to apply redundancy recognition patterns, so that, within each Minimal Map, redundancy, disjointedness and subsumption is determined. This will have largely been performed already by the preceding steps.

[0763] Once the Minimal Map has been fully processed it is recorded along with its classes as a set of RDF triples.

[0764] Finally the Minimal Maps must be assembled into a single map by querying the RDF triple generated above. This will be a map of all the classes for which an alignment with acceptable threshold value was found. There may be unaligned items.

[0765] Using the cumulative matching formula the final match value  $MV_g = 0.9375$ .

[0766] Using the linear matching formula  $MV = (1 * 5 + 2 * 1 + 3 * 1) / (1 + 2 + 3) = 5.5 / 6 = 0.9167$ .

[0767] An example alignment index is shown in Table 18, which shows an alignment map for the example ontologies described above. The results have been sequenced by alignment pair and step number to highlight the effects of the various algorithms. In reality they would be performed in the # sequence (Column 1).

TABLE 18

#	Ontology 1	Ontology 2	Align Id	Mini Map Tags	Step	MV	Cum MV
4	Club	Club	4	1 Exact Match	1 SemMat	1.0000	1.0000
11	Club	Club	4	1 Anchor Point	2 Obj Prop	1.0000	1.0000
18	Club	Club	4	1 Exact Match	3 Data Prop	1.0000	1.0000
27	Club	Qango	11	1 Possible Match	1 SemMat	0.5000	0.5000
28	Club	Qango	11	1 Related Subset	2 Obj Prop	0.8000	0.6500
29	Club	Qango	11	1 Subset	3 Data Prop	0.5000	0.5500
30	Club	Qango	11	1 Sibling	4 Multi Class	0.5000	0.5125
5	Company	Company	5	1 Exact Match	1 SemMat	1.0000	1.0000
12	Company	Company	5	1 Subclass Mismatch	2 Obj Prop	0.0001	0.5000
19	Company	Company	5	1 Subset	3 Data Prop	0.7072	0.6381
25	Company	Company + Listed Company	9	1 Different Normalisation	4 Multi Class	0.5000	0.5000
26	Company	Company + Unlisted Company	10	1 Different Normalisation	4 Multi Class	1.0000	1.0000
22	employment	Work History	8	2 Possible Match	1 SemMat	0.0010	0.0010
23	employment	Work History	8	2 Related Subset	2 Obj Prop	0.8000	0.4510

TABLE 18-continued

#	Ontology 1	Ontology 2	Align Id	Mini Map	Tags	Step	MV	Cum MV
24	employment	Work History	8	2	Subset	3 Data Prop	0.8660	0.7273
2	Individual	Person	2	1	Exact Match	1 SemMat	1.0000	1.0000
9	Individual	Person	2	1	Related Subset	2 Obj Prop	0.8000	0.9000
16	Individual	Person	2	1	Exact Match	3 Data Prop	1.0000	0.9666
6	Member	Membership	6	1	Near Match	1 SemMat	0.7000	1.0000
13	Member	Membership	6	1	Anchor Point	2 Obj Prop	1.0000	0.8500
20	Member	Membership	6	1	Exact Match	3 Data Prop	1.0000	0.9500
3	Organisation	Organisation	3	1	Exact Match	1 SemMat	1.0000	1.0000
10	Organisation	Organisation	3	1	Related Subset	2 Obj Prop	0.3333	0.6667
17	Organisation	Organisation	3	1	Exact Match	3 Data Prop	1.0000	0.8888
1	Party	Client	1	1	Exact Match	1 SemMat	1.0000	1.0000
8	Party	Client	1	1	Related Subset	2 Obj Prop	0.8889	0.9259
15	Party	Client	1	1	Exact Match	3 Data Prop	1.0000	0.9815
7	Shares	Shares	7	1	Exact Match	1 SemMat	1.0000	1.0000
14	Shares	Shares	7	1	Related Subset	2 Obj Prop	0.8571	0.9285
21	Shares	Shares	7	1	Exact Match	3 Data Prop	1.0000	0.9762

[0768] A merge process can then be performed to produce a merged ontology **1906**, although this is optional and will depend on the preferred implementation. If the user decides to merge the ontologies then a number of decisions need to be made, including:

[0769] Determine whether the merged ontology should be Ontology 1 into Ontology 2, or vice versa, or whether the merged ontology should be given a new URI. These cases are shown diagrammatically in FIGS. **19D** and **19E**.

[0770] Select  $MV_{MT}$  as the Match Value threshold for merging. Generally the  $MV_{MT}$  would be lower than the  $MV_{AT}$  as we may include related classes which do not actually align.

[0771] If classes are not to be merged then a decision is required as to whether both, neither or only one of the classes should be included in the merged ontology. This can be specified as a rule, or as ‘Ask’, in which case the merge process would pause to allow the user to decide the action.

[0772] Should classes for which no alignment was found be added to the merged ontology? For example, if Ontology 1 consists of classes A,B and Ontology 2 of classes B,C where B is the set of classes which are aligned, then should the merged ontology be A,B,C, or A,B or B,C or just B?

[0773] Once the merge parameters have been determined then it is a simple matter to merge the Classes, Data Properties and Object Properties of the two ontologies.

[0774] Any Data Property instances would retain their original URI unless specified otherwise. Thus if an aligned class has instance data in each ontology then the single merged class would contain the instances from both ontologies.

[0775] In general user interaction with the aligner module will be for the purpose of controlling the alignment process.

[0776] The first step is to load the configuration file specifying parameters to be used in the alignment and merge. There are a number of metadata parameters which can be set. These include:

[0777] URI of the ontologies to be aligned.

[0778] Location to store the alignment map.

[0779] Location to store the merged ontology.

[0780] The Match Value threshold for aligning  $MV_{AT}$ .

[0781] The Match Value threshold for merging  $MV_{MT}$ .

[0782] Match quality for accepting sameness during low level class matching.

[0783] Optionally preload the Alignment Table with known alignments.

[0784] Weights to be applied at each analysis step. These could be determined by a machine learning algorithm.

[0785] Whether to pause the process during merge to allow user input on merging

[0786] Maximum run time.

[0787] Verbosity of error and log messages.

[0788] Etcetera.

[0789] The user then runs or schedules the process. If a pause for user input was specified the user provides input as requested, and as provided via screens typically displayed by the browser module.

[0790] Upon completion of the process the user examines:

[0791] a report produced giving statistics of:

[0792] number of input classes in each ontology;

[0793] number of classes aligned;

[0794] number of PMPs identified;

[0795] number of PMPs expanded;

[0796] number of classes expanded from PMPs;

[0797] number of Data Property instances expanded from PMPs;

[0798] maximum and minimum Match values;

[0799] number of Classes merged;

- [0800] number of classes in Merged Ontology;
- [0801] number of data instances in merged ontology;
- [0802] etcetera;
- [0803] the runtime logs to evaluate error, warning and information messages.
- [0804] Based upon this information the user decides to accept the alignment or merge or to vary some of the configuration parameters and reschedule the process.
- [0805] Accordingly, the above described processes allow for users to interact with ontologies to perform a variety of tasks including browsing, pruning and aligning ontologies. These processes can use a variety of modules and allow operations to be performed such as determining mappings between ontologies, including putative and formalised ontologies, which can in turn be used in mapping source and target data structures for the purpose of facilitating transfer of content between source and target data stores.
- [0806] Throughout this specification and claims which follow, unless the context requires otherwise, the word "comprise", and variations such as "comprises" or "comprising", will be understood to imply the inclusion of a stated integer or group of integers or steps but not the exclusion of any other integer or group of integers.
- [0807] Persons skilled in the art will appreciate that numerous variations and modifications will become apparent. All such variations and modifications which become apparent to persons skilled in the art, should be considered to fall within the spirit and scope that the invention broadly appearing before described.
1. Apparatus for use in generating a mapping using ontologies, the apparatus including at least one electronic processing device that:
    - determines ontologies having respective ontology terms;
    - determines a group of ontology terms from at least one of the ontologies;
    - determines alignment between ontology terms in the ontologies for at least some of the group of ontology terms, the alignment being determined at least partially in accordance with an ontology term meaning of the ontology terms; and
    - generates a mapping in accordance with the alignment; wherein the ontology term meaning is determined by comparing an ontology term with a number of potential ontology terms and generating a matching score for the results of each comparison, the ontology term meaning determining the matching score.
  2. Apparatus according to claim 1, wherein the ontology term meaning is converted to a set of RDF triples and that set of RDF triples is compared with a set of other RDF triples for potential ontology terms and generating the matching score for the comparison, the matching score determining the nature of a relationship between the ontology term meanings, thereby enabling semantic matching between the terms.
  3. Apparatus according to claim 1, wherein the mapping is at least one of:
    - a merged ontology; and
    - an alignment index.
  4. Apparatus according to claim 1, wherein the ontologies include:
    - a putative ontology; and
    - a formalized ontology.
5. Apparatus according to claim 4, wherein the electronic processing device generates a mapping to map a putative ontology to a formalized ontology.
6. Apparatus according to claim 1, wherein the mapping is for use in transferring content from a source data store having a source data structure including a number of source data fields to a target data store having a target data structure including a number of target data fields, wherein the ontologies are associated with the source data structure and the target data structure and wherein the electronic processing device transfers content between the source data fields and target data fields using the mapping.
7. Apparatus according to claim 6, wherein the ontologies include a source ontology associated with the source data structure and a target ontology associated with the target data structure.
8. Apparatus according to claim 7, wherein the source and target ontologies are putative ontologies and the source and target ontologies are mapped to one or more formalized ontologies.
9. Apparatus according to claim 1, wherein the electronic processing device determines an ontology by:
  - generating a putative ontology; and
  - selecting one of a number of existing ontologies.
10. Apparatus according to claim 9, wherein the electronic processing device selects an ontology using at least one of:
  - metadata associated with a data structure; and
  - data fields of a data structure.
11. Apparatus according to claim 10, wherein the electronic processing device selects one of a number of existing ontologies by:
  - comparing data fields to ontology terms of a number of existing ontologies; and
  - selecting one of the number of existing ontologies in accordance with the results of the comparison.
12. Apparatus according to claim 9, wherein the electronic processing device generates a putative ontology from a database schema by:
  - identifying tables in the schema;
  - creating an ontology term corresponding to each table;
  - identifying at least one bill of material table; and
  - creating an ontology term corresponding to each entry in the bill of materials table.
13. Apparatus according to claim 10, wherein the electronic processing device generates a putative ontology from a database schema by:
  - displaying an indication of the ontology term corresponding to each entry in the bill of materials table; and
  - adding the ontology term to the putative ontology in response to user input commands.
14. Apparatus according to claim 10, wherein the electronic processing device generates relationships between ontology terms using a table structure defined by the database schema.
15. Apparatus according to claim 9, wherein the putative ontology includes:
  - classes corresponding to ontology terms;
  - data properties for at least some of the classes; and
  - object properties defining relationships between classes.
16. Apparatus according to claim 1, wherein the electronic processing device:
  - determines an index for at least one ontology, the index including an indication of the ontology terms of the at least one ontology; and

uses the index to determine:  
the group of ontology terms; and  
alignment between ontology terms.

**17.** Apparatus according to claim **16**, wherein, for each ontology term, the index includes an indication of:  
an ontology term meaning; and  
an ontology term type.

**18.** Apparatus according to claim **16**, wherein the electronic processing device generates an index by, for each ontology term:  
identifying an ontology term name;  
identifying an ontology term type;  
identifying an ontology meaning of each ontology term using a semantic matching process; and  
creating an index entry including an indication of the ontology term name, the ontology term type and the ontology term meaning.

**19.** Apparatus according to claim **16**, wherein the electronic processing device generates a single index for multiple ontologies and wherein the index entry includes an indication of the ontology associated with the ontology term.

**20.** Apparatus according to claim **1**, wherein the electronic processing device:  
displays an indication of one or more ontology terms in an ontology;  
identifies at least one ontology term in response to user input commands; and  
displays details of at least one of:  
at least one identified ontology term; and  
data properties associated with the at least one identified ontology term.

**21.** Apparatus according to claim **20**, wherein the electronic processing device determines the indication of the one or more ontology terms using an index.

**22.** Apparatus according to claim **20**, wherein the electronic processing device:  
determines user selected ontology terms and corresponding user selected data properties in response to user input commands; and  
generates executable code in accordance with the user selected ontology terms and corresponding user selected data properties, the executable code, when executed on a computer system, causes the computer system to display a user interface for allowing a user to interact with content stored in a data store having a data structure, the content being stored in data fields corresponding to the user selected ontology terms.

**23.** Apparatus according to claim **22**, wherein the executable code causes the computer system to generate queries for interacting with data stored in a source or target data structure in accordance with a source or target ontology.

**24.** Apparatus according to claim **22**, wherein the executable code causes the computer system to generate queries in accordance with at least one of data properties and relationships between the user selected ontology terms of an ontology.

**25.** Apparatus according to claim **22**, wherein the executable code causes the computer system to:  
display an indication of one or more ontology terms;  
determine selection of at least one ontology term in response to user input commands; and  
query data stored in a data field associated with the selected ontology term.

**26.** Apparatus according to claim **1**, wherein the electronic processing device:  
compares an ontology term to a number of potential ontology term meanings; and  
selects one of the potential ontology term meanings as the ontology term meaning in accordance with the results of the comparison.

**27.** Apparatus according to claim **26**, wherein the electronic processing device, determines if a potential ontology term meaning is at least one of:  
a synonym;  
an antonym;  
a meronym;  
a superclass; and  
a subclass.

**28.** Apparatus according to claim **26**, wherein the electronic processing device stores an indication of the ontology term meaning in an index, the indication including at least one of:  
an identifier indicative of a defined meaning; and  
a list of equivalent meanings.

**29.** Apparatus according to claim **1**, wherein the electronic processing device, determines an alignment between ontology terms by:  
comparing ontology term meanings of a number of ontology terms;  
generating a matching score for the results of each comparison; and  
determining an alignment in accordance with matching scores.

**30.** Apparatus according to claim **9**, wherein the electronic processing device, further determines the alignment based on at least one of:  
relationships between ontology terms; and  
data properties of the ontology terms.

**31.** Apparatus according to claim **29**, wherein the electronic processing device:  
determines relationships between source ontology terms in a source ontology;  
determines relationships between the target ontology terms in a target ontology;  
compares the relationships; and  
determines the alignment in accordance with the results of the comparison.

**32.** Apparatus according to claim **29**, wherein the electronic processing device, determines alignment between one or more source ontology terms from a source ontology and one or more target ontology terms from a target ontology.

**33.** Apparatus according to claim **1**, wherein the electronic processing device, determines the group of ontology terms by:  
determining selected ontology terms; and  
determining the group of ontology terms at least partially

in accordance with the selected ontology terms and relationships between the selected ontology terms.

**34.** Apparatus according to claim **33**, wherein the electronic processing device, determines related ontology terms by:  
for each selected ontology term, identifying ontology terms within a defined relationship path length for each of a number of different types of relationship; and  
adding the identified ontology terms to the group of ontology terms.

**35.** Apparatus according to claim **34**, wherein the electronic processing device repeatedly adds identified ontology terms until all the selected ontology terms are connected by relationships.

**36.** Apparatus according to claim **33** wherein the electronic processing device uses a different relationship path length for different types of relationship.

**37.** Apparatus according to claim **36**, wherein the electronic processing device determines a relationship path length in accordance with user input commands.

**38.** Apparatus according to claim **1**, wherein the apparatus includes:

- an indexer module that generates an index indicative of ontology terms in an ontology;
- a browser module that enables browsing of ontology terms in an ontology and generates code embodying at least part of the ontology thereby allowing a user to interact with data stored in a data structure in accordance with the ontology;
- an aligner module that determines alignment between ontology terms different ontologies;
- a pruner module that determines a group of ontology terms within at least one ontology at least in part using relationships between the ontology terms; and

a semantic matcher module that identifies ontology term meanings.

**39.** A method for use in generating a mapping using ontologies, the method including, in at least one electronic processing device:

- determining ontologies having respective ontology terms;
- determining a group of ontology terms from at least one of the ontologies;
- determining alignment between ontology terms in the ontologies for at least some of the group of ontology terms, the alignment being determined at least partially in accordance with an ontology term meaning of the ontology terms; and
- generating a mapping in accordance with the alignment; wherein the ontology term meaning is determined by comparing an ontology term with a number of potential ontology terms and generating a matching score for the results of each comparison, the ontology term meaning determining the matching score.

**40.** A mapping using ontologies produced using the apparatus of any one of claims **1** to **38**.

**41.** A Putative ontology produced by the apparatus of claim **11**.

\* \* \* \* \*